

A1_Part3_solution

September 10, 2023

0.1 Part 3: Convolutional Neural Networks and Image Classification

[Total marks for this part: 40 points]

This part of the assignment is designed to assess your knowledge and coding skill with Tensorflow as well as hands-on experience with training Convolutional Neural Network (CNN).

The dataset we use for this part is the [STL10 dataset](#) which consists of 5,000 training images of airplane, bird, car, cat, deer, dog, horse, monkey, ship, truck; each of which has 500 images. You can download the dataset at [download here](#) and then decompress to the folder `datasets\Animals` in your assignment folder.**

Your task is to build a CNN model using *TF 2.x* to classify these animals. You're provided with the module `models.py`, which you can find in the assignment folder, with some of the following classes:

1. **DatasetManager**: Support with loading and splitting the dataset into the train-val-test sets. It also supports generating next batches for training.
2. **BaseImageClassifier**: A base class image classification, which is basically a CNN model.

Note: You may need to install the package `imutils` if you have not installed yet

Firstly, we need to run the following cells to load required packages.

```
[ ]: %load_ext autoreload
      %autoreload 2
```

```
[ ]: import os
      import tensorflow as tf
      import matplotlib.pyplot as plt
      plt.style.use('ggplot')
      %matplotlib inline
      from A1_S2_2023 import DatasetManager, BaseImageClassifier
```

Note that the class `DatasetManager` has attributes related to *the training, validation, and testing sets*. You can use them in training your developed models in the sequel.

```
[ ]: dataset_name = 'stl10'
      # Choose path to store dataset
      data_dir = '{}'/tensorflow_datasets'.format(os.path.expanduser('~'))

      data_manager = DatasetManager(dataset_name, data_dir)
      data_manager.load_dataset()
```

```
data_manager.preprocess_dataset()
data_manager.show_examples()
```

Downloading and preparing dataset 2.46 GiB (download: 2.46 GiB, generated: 1.86 GiB, total: 4.32 GiB) to /root/tensorflow_datasets/stl10/1.0.0...

Dl Completed...: 0 url [00:00, ? url/s]

Dl Size...: 0 MiB [00:00, ? MiB/s]

Extraction completed...: 0 file [00:00, ? file/s]

Generating splits...: 0%| | 0/3 [00:00<?, ? splits/s]

Generating train examples...: 0%| | 0/5000 [00:00<?, ? examples/s]

Shuffling /root/tensorflow_datasets/stl10/1.0.0.incompleteHZ8LUJ/stl10-train.

↵tfrecord*...: 0%| | 0/...

Generating test examples...: 0%| | 0/8000 [00:00<?, ? examples/s]

Shuffling /root/tensorflow_datasets/stl10/1.0.0.incompleteHZ8LUJ/stl10-test.

↵tfrecord*...: 0%| | 0/8...

Generating unlabelled examples...: 0%| | 0/100000 [00:00<?, ? examples/s]


Shuffling /root/tensorflow_datasets/stl10/1.0.0.incompleteHZ8LUJ/

↵stl10-unlabelled.tfrecord*...: 0%| ...

Dataset stl10 downloaded and prepared to /root/tensorflow_datasets/stl10/1.0.0. Subsequent calls will reuse this data.



bird (1)



airplane (0)



ship (8)



cat (3)



truck (9)



ship (8)



airplane (0)



car (2)



deer (4)

```
[ ]: # Choose a random example
import random
num_examples = tf.data.experimental.cardinality(data_manager.ds_train).numpy()
random_index = random.randint(0, num_examples - 1)
example = next(iter(data_manager.ds_train.skip(random_index).take(1)))[0]

# Print the shape and value of the image
print("Image shape:", example.shape)
print("Image value range:", example.numpy().min(), "to", example.numpy().max())
```

Image shape: (32, 32, 3)

Image value range: 0.09803922 to 0.9647059

```
[ ]: # Check the number of examples in each dataset
print(tf.data.experimental.cardinality(data_manager.ds_train))
print(tf.data.experimental.cardinality(data_manager.ds_val))
print(tf.data.experimental.cardinality(data_manager.ds_test))
```

```
tf.Tensor(4500, shape=(), dtype=int64)
tf.Tensor(500, shape=(), dtype=int64)
tf.Tensor(8000, shape=(), dtype=int64)
```

We now use **BaseImageClassifier** built in the **A1_S2_2023.py** file which serves as a basic baseline to start the investigation. Follow the following steps to realize how to run a model and know the built-in methods associated with.

```
[ ]: network1 = BaseImageClassifier(name='network1',
                                   num_classes=10,
                                   optimizer='sgd',
                                   batch_size=128,
                                   num_epochs=20,
                                   learning_rate=0.001)
```

We first initialize a default model from the DefaultModel class. Basically, we can define the relevant parameters of training a model including `num_classes`, `optimizer`, `learning_rate`, `batch_size`, and `num_epochs`.

The method `build_cnn()` assists us in building your convolutional neural network. You can view the code (in the **A1_S2_2023.py** file) of the model behind a default model to realize how simple it is. Additionally, the method `summary()` shows the architecture of a model.

```
[ ]: network1.build_cnn()
network1.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 32)	896
conv2d_1 (Conv2D)	(None, 32, 32, 32)	9248
average_pooling2d (AveragePooling2D)	(None, 16, 16, 32)	0
conv2d_2 (Conv2D)	(None, 16, 16, 64)	18496
conv2d_3 (Conv2D)	(None, 16, 16, 64)	36928
average_pooling2d_1 (AveragePooling2D)	(None, 8, 8, 64)	0

flatten (Flatten)	(None, 4096)	0
dense (Dense)	(None, 10)	40970

```
=====
Total params: 106,538
Trainable params: 106,538
Non-trainable params: 0
-----
None
```

```
[ ]: x_train_batch = network1.optimize_data_pipeline(data_manager.ds_train,
    ↪batch_size=32)
x_val_batch = network1.optimize_data_pipeline(data_manager.ds_val,
    ↪batch_size=32)
network1.fit(x_train_batch, x_val_batch, num_epochs=20)
```

```
Epoch 1/20
141/141 [=====] - 16s 31ms/step - loss: 2.2981 -
accuracy: 0.1222 - val_loss: 2.2889 - val_accuracy: 0.1500
Epoch 2/20
141/141 [=====] - 1s 5ms/step - loss: 2.2765 -
accuracy: 0.1427 - val_loss: 2.2474 - val_accuracy: 0.1940
Epoch 3/20
141/141 [=====] - 1s 5ms/step - loss: 2.1806 -
accuracy: 0.1969 - val_loss: 2.1057 - val_accuracy: 0.2500
Epoch 4/20
141/141 [=====] - 1s 5ms/step - loss: 2.0548 -
accuracy: 0.2629 - val_loss: 1.9850 - val_accuracy: 0.2960
Epoch 5/20
141/141 [=====] - 1s 4ms/step - loss: 1.9539 -
accuracy: 0.3029 - val_loss: 1.9003 - val_accuracy: 0.3360
Epoch 6/20
141/141 [=====] - 1s 4ms/step - loss: 1.8941 -
accuracy: 0.3187 - val_loss: 1.8511 - val_accuracy: 0.3460
Epoch 7/20
141/141 [=====] - 1s 5ms/step - loss: 1.8360 -
accuracy: 0.3376 - val_loss: 1.8105 - val_accuracy: 0.3560
Epoch 8/20
141/141 [=====] - 1s 5ms/step - loss: 1.7824 -
accuracy: 0.3649 - val_loss: 1.7775 - val_accuracy: 0.3620
Epoch 9/20
141/141 [=====] - 1s 5ms/step - loss: 1.7403 -
accuracy: 0.3756 - val_loss: 1.7436 - val_accuracy: 0.3720
Epoch 10/20
141/141 [=====] - 1s 4ms/step - loss: 1.7052 -
accuracy: 0.3900 - val_loss: 1.7134 - val_accuracy: 0.3800
Epoch 11/20
```

```

141/141 [=====] - 1s 6ms/step - loss: 1.6717 -
accuracy: 0.4027 - val_loss: 1.6897 - val_accuracy: 0.3780
Epoch 12/20
141/141 [=====] - 1s 6ms/step - loss: 1.6402 -
accuracy: 0.4124 - val_loss: 1.6700 - val_accuracy: 0.3880
Epoch 13/20
141/141 [=====] - 1s 5ms/step - loss: 1.6102 -
accuracy: 0.4251 - val_loss: 1.6527 - val_accuracy: 0.3940
Epoch 14/20
141/141 [=====] - 1s 5ms/step - loss: 1.5811 -
accuracy: 0.4329 - val_loss: 1.6379 - val_accuracy: 0.4040
Epoch 15/20
141/141 [=====] - 1s 4ms/step - loss: 1.5540 -
accuracy: 0.4376 - val_loss: 1.6254 - val_accuracy: 0.4260
Epoch 16/20
141/141 [=====] - 1s 5ms/step - loss: 1.5261 -
accuracy: 0.4478 - val_loss: 1.6139 - val_accuracy: 0.4340
Epoch 17/20
141/141 [=====] - 1s 4ms/step - loss: 1.4978 -
accuracy: 0.4564 - val_loss: 1.6024 - val_accuracy: 0.4440
Epoch 18/20
141/141 [=====] - 1s 5ms/step - loss: 1.4693 -
accuracy: 0.4700 - val_loss: 1.5910 - val_accuracy: 0.4540
Epoch 19/20
141/141 [=====] - 1s 4ms/step - loss: 1.4404 -
accuracy: 0.4800 - val_loss: 1.5783 - val_accuracy: 0.4600
Epoch 20/20
141/141 [=====] - 1s 5ms/step - loss: 1.4111 -
accuracy: 0.4931 - val_loss: 1.5668 - val_accuracy: 0.4560

```

To train a model regarding to the datasets stored in `data_manager`, you can invoke the method `fit()` for which you can specify the batch size and number of epochs for your training.

```

[ ]: x_test_batch = network1.optimize_data_pipeline(data_manager.ds_test,
↳ batch_size=32)
network1.compute_accuracy(x_test_batch)

```

```

250/250 [=====] - 4s 17ms/step - loss: 1.5981 -
accuracy: 0.4289
loss: 1.5980982780456543
accuracy: 0.42887499928474426

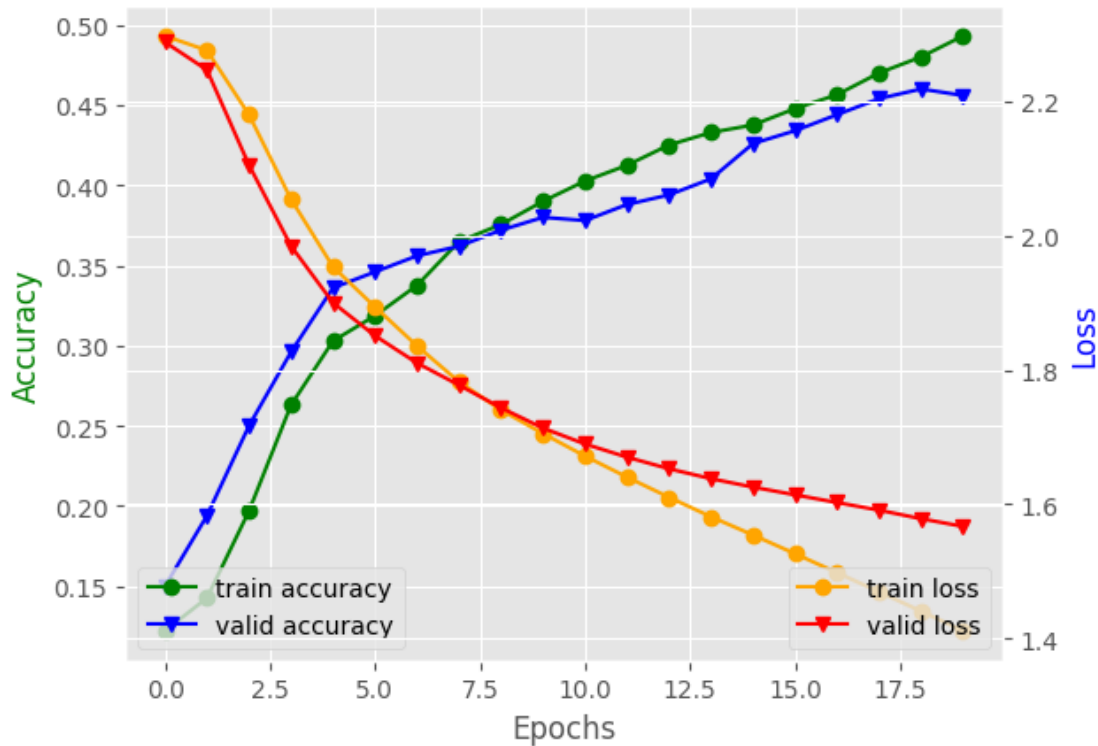
```

Below shows how you can inspect the training progress.

```

[ ]: network1.plot_progress()

```



You can use the method `predict()` to predict labels for data examples in a test set.

```
[ ]: num_samples = 25
      sample_dataset = data_manager.ds_test.take(num_samples)
      network1.predict(sample_dataset.batch(num_samples), data_manager.ds_info)
```

1/1 [=====] - 0s 298ms/step

Sample 1: Predicted label - ship
 Sample 2: Predicted label - monkey
 Sample 3: Predicted label - cat
 Sample 4: Predicted label - bird
 Sample 5: Predicted label - ship
 Sample 6: Predicted label - dog
 Sample 7: Predicted label - deer
 Sample 8: Predicted label - cat
 Sample 9: Predicted label - airplane
 Sample 10: Predicted label - deer
 Sample 11: Predicted label - deer
 Sample 12: Predicted label - ship
 Sample 13: Predicted label - horse
 Sample 14: Predicted label - ship
 Sample 15: Predicted label - deer
 Sample 16: Predicted label - cat

Sample 17: Predicted label - deer
 Sample 18: Predicted label - airplane
 Sample 19: Predicted label - truck
 Sample 20: Predicted label - ship
 Sample 21: Predicted label - dog
 Sample 22: Predicted label - horse
 Sample 23: Predicted label - ship
 Sample 24: Predicted label - truck
 Sample 25: Predicted label - airplane

Finally, the method `plot_prediction()` visualizes the predictions for a test set in which several images are chosen to show the predictions.

```
[ ]: num_samples = 20
sample_dataset = data_manager.ds_test.take(num_samples)
network1.plot_predictions(sample_dataset, data_manager.ds_info,
    ↪num_samples=num_samples, grid_shape=(4, 5))
```

1/1 [=====] - 0s 79ms/step



0.1.1 Question 3.1: Observe the learning curve

After running the above cells to train the default model and observe the learning curve. Report your observation (i.e. did the model learn well? if not, what is the problem? What would you do to improve it?). Write your answer below.

[4 points]

The model behaves not too good. The accuracy for test dataset is only 0.422.

What could be done to improve: increase number of epoches, add skip connection, add dropout

For questions 3.2 to 3.9, you'll need to write your own model in a way that makes it easy for you to experiment with different architectures and parameters. The goal is to be able to pass the parameters to initialize a new instance of `YourModel` to build different network architectures with different parameters. Below are descriptions of some parameters which you can find in function `__init__()` for the class `BaseImageClassifier`:

1. **num_blocks**: an integer specifying the number of blocks in our network. Each block has the pattern `[conv, batch norm, activation, conv, batch norm, activation, mean pool, dropout]`. All convolutional layers have filter size (3,3), strides (1,1) and 'SAME' padding, and all mean pool layers have strides (2,2) and 'SAME' padding. The network will consists of a few blocks before applying a linear layer to output the logits for the softmax layer.
2. **feature_maps**: the number of feature maps in the first block of the network. The number of feature_maps will double in each of the following block. To make it convenient for you, we already calculated the number of feature maps for each block for you in line 106
3. **drop_rate**: the keep probability for dropout. Setting **drop_rate** to 0.0 means not using dropout.
4. **batch_norm**: the batch normalization function is used or not. Setting **batch_norm** to `None` means not using batch normalization.
5. The **skip connection** is added to the output of the second **batch norm**. Additionally, your class has a boolean property (i.e., instance variable) named **use_skip**. If **use_skip**=`True`, the skip connectnion is enable. Otherwise, if **use_skip**=`False`, the skip connectnion is disable.

Below is the architecture of one block:

Below is the architecture of the entire deep net with **two blocks**:

Here we assume that the first block has `feature_maps = feature_maps[0] = 32`. Note that the initial number of feature maps of the first block is declared in the instance variable **feature_maps** and is multiplied by 2 in each follpwng block.

```
[ ]: import tensorflow as tf
      from tensorflow import keras
      from tensorflow.keras import layers, models

      tf.random.set_seed(3181)
```

0.1.2 Question 3.2: Define your CNN

Write the code of the YourModel class here. Note that this class will be inherited from the BaseImageClassifier class. You'll only need to re-write the code for the build_cnn method in the YourModel class from the cell below.

[4 points]

```
[ ]: import tensorflow as tf
from tensorflow.keras.layers import Input, Conv2D, BatchNormalization,
    ↪Activation, AveragePooling2D, Flatten, Dropout
from tensorflow.keras.models import Model

class YourModel(BaseImageClassifier):
    def __init__(self,
                  name='network1',
                  width=32, height=32, depth=3,
                  num_blocks=2,
                  feature_maps=32,
                  num_classes=4,
                  drop_rate=0.2,
                  batch_norm=None,
                  is_augmentation=False,
                  activation_func='relu',
                  optimizer='adam',
                  use_skip=True,
                  batch_size=10,
                  num_epochs=20,
                  learning_rate=0.0001,
                  verbose=True):
        super(YourModel, self).__init__(name, width, height, depth, num_blocks,
    ↪feature_maps, num_classes, drop_rate, batch_norm, is_augmentation,
                                   activation_func, use_skip, optimizer,
    ↪batch_size, num_epochs, learning_rate, verbose)

    def custom_block(self, x, filters):
        # skip connection
        shortcut = x

        # First Conv layer
        x = layers.Conv2D(filters=filters, kernel_size=(3, 3), strides=(1, 1),
    ↪padding="same")(x)
        if self.batch_norm:
            x = layers.BatchNormalization()(x)
        x = layers.Activation(self.activation_func)(x)

        # Second Conv layer
```

```

        x = layers.Conv2D(filters=filters, kernel_size=(3, 3), strides=(1, 1),
↪padding="same")(x)
        if self.batch_norm:
            x = layers.BatchNormalization()(x)

        # Add the skip connection to the output
        if self.use_skip:
            x = tf.keras.layers.add([shortcut, x])
        x = layers.Activation(self.activation_func)(x)

        # Mean Pooling layer
        x = layers.AveragePooling2D(pool_size=(2, 2), strides=(2, 2),
↪padding="same")(x)

        # Dropout
        x = layers.Dropout(rate=self.drop_rate)(x)

        return x

    def build_cnn(self):
        input_tensor = layers.Input(shape=(self.height, self.width, self.depth))
        x = input_tensor

        for current_feature_maps in self.feature_maps:
            # Initial convolution
            x = layers.Conv2D(current_feature_maps, (3, 3), padding='same',
↪activation=None)(x)
            if self.batch_norm:
                x = layers.BatchNormalization()(x)
            x = layers.Activation(self.activation_func)(x)

            # Second convolution with skip connection
            shortcut = x
            x = layers.Conv2D(current_feature_maps, (3, 3), padding='same',
↪activation=None)(x)
            if self.batch_norm:
                x = layers.BatchNormalization()(x)
            if self.use_skip:
                x = layers.Add()([shortcut, x])
            x = layers.Activation(self.activation_func)(x)

            # Mean Pooling and Dropout
            x = layers.AveragePooling2D(pool_size=(2, 2), padding='same')(x)
            if self.drop_rate > 0:
                x = layers.Dropout(self.drop_rate)(x)

        x = layers.Flatten()(x)

```

```

x = layers.Dense(self.num_classes, activation='softmax')(x)

self.model = models.Model(inputs=input_tensor, outputs=x)
self.model.compile(optimizer=self.optimizer,
    ↪loss='sparse_categorical_crossentropy', metrics=['accuracy'])

```

0.1.3 Question 3.3: Experiment with skip connection

Once writing your own model, you need to compare two cases: (i) *using the skip connection* and (ii) *not using the skip connection*. You should set the instance variable `use_skip` to either `True` or `False`. For your runs, report which case is better and if you confront overfitting in training.

[6 points]

WRITE YOUR ANSWER AND OBSERVATION HERE

....

```

[ ]: import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers, models

my_network_skip = YourModel(name='network1',
                             feature_maps=32,
                             num_classes=data_manager.n_classes,
                             num_blocks=3,
                             drop_rate=0.0,
                             batch_norm=True,
                             use_skip=True,
                             optimizer='sgd',
                             learning_rate=0.001)
my_network_skip.build_cnn()
my_network_skip.summary()

```

Model: "model"

```

-----
-----
Layer (type)                Output Shape          Param #   Connected to
=====
input_1 (InputLayer)        [(None, 32, 32, 3)]  0         []

conv2d_4 (Conv2D)           (None, 32, 32, 32)   896       ['input_1[0][0]']

batch_normalization (BatchNorm (None, 32, 32, 32)  128       ['conv2d_4[0][0]']
alization)

```

activation (Activation) ['batch_normalization[0][0]']	(None, 32, 32, 32)	0	
conv2d_5 (Conv2D) ['activation[0][0]']	(None, 32, 32, 32)	9248	
batch_normalization_1 (BatchNormal- ization) ['conv2d_5[0][0]']	(None, 32, 32, 32)	128	
add (Add) ['activation[0][0]', 'batch_normalization_1[0][0]']	(None, 32, 32, 32)	0	
activation_1 (Activation)	(None, 32, 32, 32)	0	['add[0][0]']
average_pooling2d_2 (AveragePool- ing2D) ['activation_1[0][0]']	(None, 16, 16, 32)	0	
conv2d_6 (Conv2D) ['average_pooling2d_2[0][0]']	(None, 16, 16, 64)	18496	
batch_normalization_2 (BatchNormal- ization) ['conv2d_6[0][0]']	(None, 16, 16, 64)	256	
activation_2 (Activation) ['batch_normalization_2[0][0]']	(None, 16, 16, 64)	0	
conv2d_7 (Conv2D) ['activation_2[0][0]']	(None, 16, 16, 64)	36928	
batch_normalization_3 (BatchNormal- ization) ['conv2d_7[0][0]']	(None, 16, 16, 64)	256	
add_1 (Add) ['activation_2[0][0]', 'batch_normalization_3[0][0]']	(None, 16, 16, 64)	0	
activation_3 (Activation)	(None, 16, 16, 64)	0	['add_1[0][0]']
average_pooling2d_3 (AveragePool- ing2D) ['activation_3[0][0]']	(None, 8, 8, 64)	0	
conv2d_8 (Conv2D)	(None, 8, 8, 128)	73856	

```

['average_pooling2d_3[0][0]']

batch_normalization_4 (BatchNo (None, 8, 8, 128) 512
['conv2d_8[0][0]']
rmalization)

activation_4 (Activation) (None, 8, 8, 128) 0
['batch_normalization_4[0][0]']

conv2d_9 (Conv2D) (None, 8, 8, 128) 147584
['activation_4[0][0]']

batch_normalization_5 (BatchNo (None, 8, 8, 128) 512
['conv2d_9[0][0]']
rmalization)

add_2 (Add) (None, 8, 8, 128) 0
['activation_4[0][0]',
'batch_normalization_5[0][0]']

activation_5 (Activation) (None, 8, 8, 128) 0 ['add_2[0][0]']

average_pooling2d_4 (AveragePo (None, 4, 4, 128) 0
['activation_5[0][0]']
oling2D)

flatten_1 (Flatten) (None, 2048) 0
['average_pooling2d_4[0][0]']

dense_1 (Dense) (None, 10) 20490
['flatten_1[0][0]']

```

```

=====
=====
Total params: 309,290
Trainable params: 308,394
Non-trainable params: 896
-----
-----
None

```

```
[ ]: my_network_skip.fit(x_train_batch, x_val_batch, num_epochs=20)
```

```

Epoch 1/20
141/141 [=====] - 3s 9ms/step - loss: 1.8172 -
accuracy: 0.3380 - val_loss: 2.8307 - val_accuracy: 0.1060
Epoch 2/20
141/141 [=====] - 1s 9ms/step - loss: 1.3395 -

```

accuracy: 0.5089 - val_loss: 3.2411 - val_accuracy: 0.1680
Epoch 3/20
141/141 [=====] - 1s 7ms/step - loss: 1.1227 -
accuracy: 0.5958 - val_loss: 2.2395 - val_accuracy: 0.2900
Epoch 4/20
141/141 [=====] - 1s 7ms/step - loss: 0.9691 -
accuracy: 0.6582 - val_loss: 1.5108 - val_accuracy: 0.4720
Epoch 5/20
141/141 [=====] - 1s 7ms/step - loss: 0.8469 -
accuracy: 0.7169 - val_loss: 1.3289 - val_accuracy: 0.5520
Epoch 6/20
141/141 [=====] - 1s 7ms/step - loss: 0.7444 -
accuracy: 0.7638 - val_loss: 1.3557 - val_accuracy: 0.5360
Epoch 7/20
141/141 [=====] - 1s 7ms/step - loss: 0.6531 -
accuracy: 0.8058 - val_loss: 1.3614 - val_accuracy: 0.5520
Epoch 8/20
141/141 [=====] - 1s 7ms/step - loss: 0.5696 -
accuracy: 0.8438 - val_loss: 1.4007 - val_accuracy: 0.5420
Epoch 9/20
141/141 [=====] - 1s 7ms/step - loss: 0.4970 -
accuracy: 0.8738 - val_loss: 1.4979 - val_accuracy: 0.5180
Epoch 10/20
141/141 [=====] - 1s 6ms/step - loss: 0.4394 -
accuracy: 0.8938 - val_loss: 1.6191 - val_accuracy: 0.5200
Epoch 11/20
141/141 [=====] - 1s 8ms/step - loss: 0.3891 -
accuracy: 0.9113 - val_loss: 2.0143 - val_accuracy: 0.4480
Epoch 12/20
141/141 [=====] - 1s 10ms/step - loss: 0.3448 -
accuracy: 0.9344 - val_loss: 2.4209 - val_accuracy: 0.4020
Epoch 13/20
141/141 [=====] - 1s 7ms/step - loss: 0.3029 -
accuracy: 0.9460 - val_loss: 2.4210 - val_accuracy: 0.4380
Epoch 14/20
141/141 [=====] - 1s 7ms/step - loss: 0.2607 -
accuracy: 0.9649 - val_loss: 2.1427 - val_accuracy: 0.4680
Epoch 15/20
141/141 [=====] - 1s 7ms/step - loss: 0.2190 -
accuracy: 0.9764 - val_loss: 1.7753 - val_accuracy: 0.4940
Epoch 16/20
141/141 [=====] - 1s 7ms/step - loss: 0.1865 -
accuracy: 0.9871 - val_loss: 1.5216 - val_accuracy: 0.5520
Epoch 17/20
141/141 [=====] - 1s 7ms/step - loss: 0.1598 -
accuracy: 0.9933 - val_loss: 1.4127 - val_accuracy: 0.5760
Epoch 18/20
141/141 [=====] - 1s 7ms/step - loss: 0.1373 -

```

accuracy: 0.9964 - val_loss: 1.4383 - val_accuracy: 0.5740
Epoch 19/20
141/141 [=====] - 1s 7ms/step - loss: 0.1188 -
accuracy: 0.9980 - val_loss: 1.4569 - val_accuracy: 0.5700
Epoch 20/20
141/141 [=====] - 1s 7ms/step - loss: 0.1036 -
accuracy: 0.9989 - val_loss: 1.4557 - val_accuracy: 0.5740

```

```

[ ]: my_network_no_skip = YourModel(name='network1',
    feature_maps=32,
    num_classes=data_manager.n_classes,
    num_blocks=3,
    drop_rate=0.0,
    batch_norm=True,
    use_skip=False,
    optimizer='sgd',
    learning_rate=0.001)
my_network_no_skip.build_cnn()
my_network_no_skip.summary()

```

Model: "model_1"

Layer (type)	Output Shape	Param #
input_2 (InputLayer)	[(None, 32, 32, 3)]	0
conv2d_10 (Conv2D)	(None, 32, 32, 32)	896
batch_normalization_6 (Batch Normalization)	(None, 32, 32, 32)	128
activation_6 (Activation)	(None, 32, 32, 32)	0
conv2d_11 (Conv2D)	(None, 32, 32, 32)	9248
batch_normalization_7 (Batch Normalization)	(None, 32, 32, 32)	128
activation_7 (Activation)	(None, 32, 32, 32)	0
average_pooling2d_5 (Average Pooling2D)	(None, 16, 16, 32)	0
conv2d_12 (Conv2D)	(None, 16, 16, 64)	18496
batch_normalization_8 (Batch Normalization)	(None, 16, 16, 64)	256

activation_8 (Activation)	(None, 16, 16, 64)	0
conv2d_13 (Conv2D)	(None, 16, 16, 64)	36928
batch_normalization_9 (Batch Normalization)	(None, 16, 16, 64)	256
activation_9 (Activation)	(None, 16, 16, 64)	0
average_pooling2d_6 (Average Pooling2D)	(None, 8, 8, 64)	0
conv2d_14 (Conv2D)	(None, 8, 8, 128)	73856
batch_normalization_10 (Batch Normalization)	(None, 8, 8, 128)	512
activation_10 (Activation)	(None, 8, 8, 128)	0
conv2d_15 (Conv2D)	(None, 8, 8, 128)	147584
batch_normalization_11 (Batch Normalization)	(None, 8, 8, 128)	512
activation_11 (Activation)	(None, 8, 8, 128)	0
average_pooling2d_7 (Average Pooling2D)	(None, 4, 4, 128)	0
flatten_2 (Flatten)	(None, 2048)	0
dense_2 (Dense)	(None, 10)	20490

```

=====
Total params: 309,290
Trainable params: 308,394
Non-trainable params: 896
-----
None

```

```
[ ]: my_network_no_skip.fit(x_train_batch, x_val_batch, num_epochs=20)
```

```

Epoch 1/20
141/141 [=====] - 3s 9ms/step - loss: 1.7145 -
accuracy: 0.3591 - val_loss: 2.5450 - val_accuracy: 0.1120
Epoch 2/20
141/141 [=====] - 1s 6ms/step - loss: 1.3194 -
accuracy: 0.5147 - val_loss: 2.6301 - val_accuracy: 0.1600

```

Epoch 3/20
141/141 [=====] - 1s 7ms/step - loss: 1.1367 - accuracy: 0.5996 - val_loss: 1.9719 - val_accuracy: 0.3020

Epoch 4/20
141/141 [=====] - 1s 6ms/step - loss: 0.9927 - accuracy: 0.6609 - val_loss: 1.5810 - val_accuracy: 0.4420

Epoch 5/20
141/141 [=====] - 1s 6ms/step - loss: 0.8759 - accuracy: 0.7138 - val_loss: 1.5433 - val_accuracy: 0.4580

Epoch 6/20
141/141 [=====] - 1s 7ms/step - loss: 0.7781 - accuracy: 0.7636 - val_loss: 1.9402 - val_accuracy: 0.4120

Epoch 7/20
141/141 [=====] - 1s 8ms/step - loss: 0.6990 - accuracy: 0.7980 - val_loss: 2.0749 - val_accuracy: 0.4020

Epoch 8/20
141/141 [=====] - 1s 9ms/step - loss: 0.6278 - accuracy: 0.8309 - val_loss: 2.0775 - val_accuracy: 0.4140

Epoch 9/20
141/141 [=====] - 1s 7ms/step - loss: 0.5581 - accuracy: 0.8598 - val_loss: 2.0063 - val_accuracy: 0.4400

Epoch 10/20
141/141 [=====] - 1s 7ms/step - loss: 0.4959 - accuracy: 0.8862 - val_loss: 2.0783 - val_accuracy: 0.4420

Epoch 11/20
141/141 [=====] - 1s 6ms/step - loss: 0.4391 - accuracy: 0.9091 - val_loss: 2.1168 - val_accuracy: 0.4320

Epoch 12/20
141/141 [=====] - 1s 7ms/step - loss: 0.3874 - accuracy: 0.9307 - val_loss: 1.9595 - val_accuracy: 0.4360

Epoch 13/20
141/141 [=====] - 1s 7ms/step - loss: 0.3393 - accuracy: 0.9467 - val_loss: 1.7357 - val_accuracy: 0.4820

Epoch 14/20
141/141 [=====] - 1s 7ms/step - loss: 0.2954 - accuracy: 0.9613 - val_loss: 1.6131 - val_accuracy: 0.5120

Epoch 15/20
141/141 [=====] - 1s 7ms/step - loss: 0.2531 - accuracy: 0.9744 - val_loss: 1.6467 - val_accuracy: 0.5040

Epoch 16/20
141/141 [=====] - 1s 7ms/step - loss: 0.2167 - accuracy: 0.9844 - val_loss: 1.5854 - val_accuracy: 0.5160

Epoch 17/20
141/141 [=====] - 1s 7ms/step - loss: 0.1854 - accuracy: 0.9918 - val_loss: 1.4717 - val_accuracy: 0.5200

Epoch 18/20
141/141 [=====] - 1s 9ms/step - loss: 0.1580 - accuracy: 0.9951 - val_loss: 1.3932 - val_accuracy: 0.5460

Epoch 19/20

141/141 [=====] - 1s 8ms/step - loss: 0.1354 - accuracy: 0.9976 - val_loss: 1.3732 - val_accuracy: 0.5460

Epoch 20/20

141/141 [=====] - 1s 7ms/step - loss: 0.1174 - accuracy: 0.9984 - val_loss: 1.3709 - val_accuracy: 0.5680

The new model works apparently much much better. The original accuracy was only around 0.422 with skip, the accuracy becomes 0.612 impressive without skip, the accuracy becomes 0.564

It improves, but seems not much. And apparently, there is overfitting issue as accuracy is much higher than the val_accuracy

0.1.4 Question 3.4: Tune hyperparameters with grid search

Now, let us tune the $num_blocks \in \{2, 3, 4\}$, $use_skip \in \{True, False\}$, and $learning_rate \in \{0.001, 0.0001\}$. Write your code for this tuning and report the result of the best model on the testing set. Note that you need to show your code for tuning and evaluating on the test set to earn the full marks. During tuning, you can set the instance variable `verbose` of your model to `False` for not showing the training details of each epoch.

[4 points]

REPORT THE BEST PARAMETERS AND THE TESTING ACCURACY HERE

.....

```
[ ]: # YOU ARE REQUIRED TO INSERT YOUR CODES IN THIS CELL
# You can add more cells if necessary
best_accuracy = 0
best_model = None
best_params = {}

#set up different parameters for number_blocks, skip, and learning rate
for blocks in [2, 3, 4]:
    for skip in [True, False]:
        for lr in [0.001, 0.0001]:
            print(f"Training model with num_blocks={blocks}, use_skip={skip},
↪and learning_rate={lr}")

            model_3_4=YourModel(name='q3_4',
                                feature_maps=32,
                                num_classes=data_manager.n_classes,
                                num_blocks=blocks,
                                drop_rate=0.0,
                                batch_norm=True,
                                use_skip=skip,
                                optimizer='sgd',
```

```

learning_rate=lr)

model_3_4.build_cnn()
model_3_4.fit(x_train_batch, x_val_batch,num_epochs=20)

val_accuracy = max(model_3_4.history.history['val_accuracy'])
if val_accuracy > best_accuracy:
    best_accuracy = val_accuracy
    best_model = model_3_4
    best_params = {
        "num_blocks": blocks,
        "use_skip": skip,
        "learning_rate": lr
    }

print(f"Best validation accuracy is {best_accuracy} with params: {best_params}")

```

Training model with num_blocks=2, use_skip=True, and learning_rate=0.001

Epoch 1/20

141/141 [=====] - 3s 8ms/step - loss: 1.9165 - accuracy: 0.3236 - val_loss: 2.7488 - val_accuracy: 0.0920

Epoch 2/20

141/141 [=====] - 1s 6ms/step - loss: 1.3802 - accuracy: 0.4924 - val_loss: 2.7620 - val_accuracy: 0.1980

Epoch 3/20

141/141 [=====] - 1s 5ms/step - loss: 1.1641 - accuracy: 0.5789 - val_loss: 2.1526 - val_accuracy: 0.3280

Epoch 4/20

141/141 [=====] - 1s 5ms/step - loss: 1.0179 - accuracy: 0.6398 - val_loss: 1.7824 - val_accuracy: 0.4180

Epoch 5/20

141/141 [=====] - 1s 5ms/step - loss: 0.9008 - accuracy: 0.6876 - val_loss: 1.6467 - val_accuracy: 0.4600

Epoch 6/20

141/141 [=====] - 1s 5ms/step - loss: 0.8078 - accuracy: 0.7273 - val_loss: 1.7132 - val_accuracy: 0.4660

Epoch 7/20

141/141 [=====] - 1s 5ms/step - loss: 0.7297 - accuracy: 0.7631 - val_loss: 1.9207 - val_accuracy: 0.4160

Epoch 8/20

141/141 [=====] - 1s 5ms/step - loss: 0.6576 - accuracy: 0.7907 - val_loss: 1.9553 - val_accuracy: 0.4280

Epoch 9/20

141/141 [=====] - 1s 5ms/step - loss: 0.5889 - accuracy: 0.8233 - val_loss: 2.1519 - val_accuracy: 0.4160

Epoch 10/20
141/141 [=====] - 1s 6ms/step - loss: 0.5281 - accuracy: 0.8460 - val_loss: 1.8633 - val_accuracy: 0.4640

Epoch 11/20
141/141 [=====] - 1s 7ms/step - loss: 0.4728 - accuracy: 0.8691 - val_loss: 1.7808 - val_accuracy: 0.4920

Epoch 12/20
141/141 [=====] - 1s 7ms/step - loss: 0.4238 - accuracy: 0.8920 - val_loss: 1.9491 - val_accuracy: 0.4660

Epoch 13/20
141/141 [=====] - 1s 5ms/step - loss: 0.3834 - accuracy: 0.9131 - val_loss: 2.0269 - val_accuracy: 0.4840

Epoch 14/20
141/141 [=====] - 1s 5ms/step - loss: 0.3549 - accuracy: 0.9233 - val_loss: 1.8934 - val_accuracy: 0.4880

Epoch 15/20
141/141 [=====] - 1s 5ms/step - loss: 0.3300 - accuracy: 0.9322 - val_loss: 1.7375 - val_accuracy: 0.5140

Epoch 16/20
141/141 [=====] - 1s 6ms/step - loss: 0.3065 - accuracy: 0.9407 - val_loss: 1.6873 - val_accuracy: 0.5260

Epoch 17/20
141/141 [=====] - 1s 5ms/step - loss: 0.2790 - accuracy: 0.9480 - val_loss: 1.7562 - val_accuracy: 0.5060

Epoch 18/20
141/141 [=====] - 1s 5ms/step - loss: 0.2550 - accuracy: 0.9569 - val_loss: 1.7542 - val_accuracy: 0.5100

Epoch 19/20
141/141 [=====] - 1s 5ms/step - loss: 0.2329 - accuracy: 0.9649 - val_loss: 1.6554 - val_accuracy: 0.5280

Epoch 20/20
141/141 [=====] - 1s 6ms/step - loss: 0.2136 - accuracy: 0.9729 - val_loss: 1.6399 - val_accuracy: 0.5220

Training model with num_blocks=2, use_skip=True, and learning_rate=0.0001

Epoch 1/20
141/141 [=====] - 2s 8ms/step - loss: 2.1670 - accuracy: 0.2229 - val_loss: 2.3032 - val_accuracy: 0.1140

Epoch 2/20
141/141 [=====] - 1s 6ms/step - loss: 1.8024 - accuracy: 0.3469 - val_loss: 2.3262 - val_accuracy: 0.1840

Epoch 3/20
141/141 [=====] - 1s 6ms/step - loss: 1.6595 - accuracy: 0.3978 - val_loss: 1.9562 - val_accuracy: 0.3120

Epoch 4/20
141/141 [=====] - 1s 5ms/step - loss: 1.5599 - accuracy: 0.4360 - val_loss: 1.6757 - val_accuracy: 0.3920

Epoch 5/20
141/141 [=====] - 1s 5ms/step - loss: 1.4812 -

accuracy: 0.4640 - val_loss: 1.5819 - val_accuracy: 0.4080
 Epoch 6/20
 141/141 [=====] - 1s 5ms/step - loss: 1.4160 -
 accuracy: 0.4893 - val_loss: 1.5388 - val_accuracy: 0.4320
 Epoch 7/20
 141/141 [=====] - 1s 5ms/step - loss: 1.3604 -
 accuracy: 0.5122 - val_loss: 1.5037 - val_accuracy: 0.4620
 Epoch 8/20
 141/141 [=====] - 1s 5ms/step - loss: 1.3117 -
 accuracy: 0.5329 - val_loss: 1.4731 - val_accuracy: 0.4700
 Epoch 9/20
 141/141 [=====] - 1s 5ms/step - loss: 1.2681 -
 accuracy: 0.5471 - val_loss: 1.4462 - val_accuracy: 0.4820
 Epoch 10/20
 141/141 [=====] - 1s 5ms/step - loss: 1.2287 -
 accuracy: 0.5616 - val_loss: 1.4231 - val_accuracy: 0.4880
 Epoch 11/20
 141/141 [=====] - 1s 6ms/step - loss: 1.1928 -
 accuracy: 0.5796 - val_loss: 1.4003 - val_accuracy: 0.4920
 Epoch 12/20
 141/141 [=====] - 1s 8ms/step - loss: 1.1597 -
 accuracy: 0.5949 - val_loss: 1.3824 - val_accuracy: 0.5120
 Epoch 13/20
 141/141 [=====] - 1s 6ms/step - loss: 1.1290 -
 accuracy: 0.6102 - val_loss: 1.3665 - val_accuracy: 0.5240
 Epoch 14/20
 141/141 [=====] - 1s 6ms/step - loss: 1.1004 -
 accuracy: 0.6220 - val_loss: 1.3530 - val_accuracy: 0.5240
 Epoch 15/20
 141/141 [=====] - 1s 5ms/step - loss: 1.0735 -
 accuracy: 0.6313 - val_loss: 1.3416 - val_accuracy: 0.5260
 Epoch 16/20
 141/141 [=====] - 1s 5ms/step - loss: 1.0480 -
 accuracy: 0.6391 - val_loss: 1.3307 - val_accuracy: 0.5260
 Epoch 17/20
 141/141 [=====] - 1s 6ms/step - loss: 1.0237 -
 accuracy: 0.6489 - val_loss: 1.3217 - val_accuracy: 0.5340
 Epoch 18/20
 141/141 [=====] - 1s 6ms/step - loss: 1.0007 -
 accuracy: 0.6618 - val_loss: 1.3138 - val_accuracy: 0.5400
 Epoch 19/20
 141/141 [=====] - 1s 5ms/step - loss: 0.9787 -
 accuracy: 0.6702 - val_loss: 1.3069 - val_accuracy: 0.5400
 Epoch 20/20
 141/141 [=====] - 1s 5ms/step - loss: 0.9578 -
 accuracy: 0.6791 - val_loss: 1.2993 - val_accuracy: 0.5440
 Training model with num_blocks=2, use_skip=False, and learning_rate=0.001
 Epoch 1/20

141/141 [=====] - 2s 7ms/step - loss: 1.7959 -
accuracy: 0.3511 - val_loss: 2.4202 - val_accuracy: 0.1260
Epoch 2/20
141/141 [=====] - 1s 8ms/step - loss: 1.3637 -
accuracy: 0.5091 - val_loss: 2.2538 - val_accuracy: 0.2160
Epoch 3/20
141/141 [=====] - 1s 6ms/step - loss: 1.1944 -
accuracy: 0.5742 - val_loss: 1.6436 - val_accuracy: 0.3920
Epoch 4/20
141/141 [=====] - 1s 5ms/step - loss: 1.0677 -
accuracy: 0.6256 - val_loss: 1.4147 - val_accuracy: 0.4900
Epoch 5/20
141/141 [=====] - 1s 6ms/step - loss: 0.9632 -
accuracy: 0.6787 - val_loss: 1.3391 - val_accuracy: 0.5260
Epoch 6/20
141/141 [=====] - 1s 5ms/step - loss: 0.8723 -
accuracy: 0.7136 - val_loss: 1.3621 - val_accuracy: 0.5080
Epoch 7/20
141/141 [=====] - 1s 5ms/step - loss: 0.7920 -
accuracy: 0.7524 - val_loss: 1.4365 - val_accuracy: 0.4960
Epoch 8/20
141/141 [=====] - 1s 5ms/step - loss: 0.7195 -
accuracy: 0.7882 - val_loss: 1.4109 - val_accuracy: 0.5100
Epoch 9/20
141/141 [=====] - 1s 5ms/step - loss: 0.6544 -
accuracy: 0.8140 - val_loss: 1.4271 - val_accuracy: 0.5020
Epoch 10/20
141/141 [=====] - 1s 6ms/step - loss: 0.5962 -
accuracy: 0.8369 - val_loss: 1.4128 - val_accuracy: 0.5040
Epoch 11/20
141/141 [=====] - 1s 5ms/step - loss: 0.5429 -
accuracy: 0.8607 - val_loss: 1.3887 - val_accuracy: 0.5160
Epoch 12/20
141/141 [=====] - 1s 6ms/step - loss: 0.4942 -
accuracy: 0.8851 - val_loss: 1.3566 - val_accuracy: 0.5360
Epoch 13/20
141/141 [=====] - 1s 7ms/step - loss: 0.4492 -
accuracy: 0.9027 - val_loss: 1.3395 - val_accuracy: 0.5480
Epoch 14/20
141/141 [=====] - 1s 6ms/step - loss: 0.4090 -
accuracy: 0.9191 - val_loss: 1.3364 - val_accuracy: 0.5620
Epoch 15/20
141/141 [=====] - 1s 5ms/step - loss: 0.3740 -
accuracy: 0.9309 - val_loss: 1.3676 - val_accuracy: 0.5600
Epoch 16/20
141/141 [=====] - 1s 5ms/step - loss: 0.3403 -
accuracy: 0.9449 - val_loss: 1.3567 - val_accuracy: 0.5540
Epoch 17/20

141/141 [=====] - 1s 5ms/step - loss: 0.3129 -
accuracy: 0.9540 - val_loss: 1.3934 - val_accuracy: 0.5400
Epoch 18/20
141/141 [=====] - 1s 6ms/step - loss: 0.2879 -
accuracy: 0.9613 - val_loss: 1.4009 - val_accuracy: 0.5380
Epoch 19/20
141/141 [=====] - 1s 5ms/step - loss: 0.2654 -
accuracy: 0.9709 - val_loss: 1.4073 - val_accuracy: 0.5500
Epoch 20/20
141/141 [=====] - 1s 5ms/step - loss: 0.2425 -
accuracy: 0.9762 - val_loss: 1.3865 - val_accuracy: 0.5580
Training model with num_blocks=2, use_skip=False, and learning_rate=0.0001
Epoch 1/20
141/141 [=====] - 2s 7ms/step - loss: 2.2027 -
accuracy: 0.2073 - val_loss: 2.3339 - val_accuracy: 0.0960
Epoch 2/20
141/141 [=====] - 1s 7ms/step - loss: 1.8676 -
accuracy: 0.3216 - val_loss: 2.2317 - val_accuracy: 0.1620
Epoch 3/20
141/141 [=====] - 1s 6ms/step - loss: 1.7387 -
accuracy: 0.3762 - val_loss: 1.9520 - val_accuracy: 0.2960
Epoch 4/20
141/141 [=====] - 1s 5ms/step - loss: 1.6540 -
accuracy: 0.4116 - val_loss: 1.7425 - val_accuracy: 0.3540
Epoch 5/20
141/141 [=====] - 1s 6ms/step - loss: 1.5878 -
accuracy: 0.4349 - val_loss: 1.6662 - val_accuracy: 0.3840
Epoch 6/20
141/141 [=====] - 1s 6ms/step - loss: 1.5322 -
accuracy: 0.4582 - val_loss: 1.6237 - val_accuracy: 0.4100
Epoch 7/20
141/141 [=====] - 1s 5ms/step - loss: 1.4841 -
accuracy: 0.4769 - val_loss: 1.5901 - val_accuracy: 0.4240
Epoch 8/20
141/141 [=====] - 1s 5ms/step - loss: 1.4418 -
accuracy: 0.4936 - val_loss: 1.5582 - val_accuracy: 0.4240
Epoch 9/20
141/141 [=====] - 1s 5ms/step - loss: 1.4034 -
accuracy: 0.5113 - val_loss: 1.5300 - val_accuracy: 0.4400
Epoch 10/20
141/141 [=====] - 1s 5ms/step - loss: 1.3680 -
accuracy: 0.5251 - val_loss: 1.5036 - val_accuracy: 0.4500
Epoch 11/20
141/141 [=====] - 1s 6ms/step - loss: 1.3355 -
accuracy: 0.5402 - val_loss: 1.4833 - val_accuracy: 0.4540
Epoch 12/20
141/141 [=====] - 1s 5ms/step - loss: 1.3054 -
accuracy: 0.5549 - val_loss: 1.4634 - val_accuracy: 0.4640

Epoch 13/20
141/141 [=====] - 1s 7ms/step - loss: 1.2772 - accuracy: 0.5622 - val_loss: 1.4443 - val_accuracy: 0.4700

Epoch 14/20
141/141 [=====] - 1s 7ms/step - loss: 1.2508 - accuracy: 0.5718 - val_loss: 1.4279 - val_accuracy: 0.4780

Epoch 15/20
141/141 [=====] - 1s 5ms/step - loss: 1.2258 - accuracy: 0.5809 - val_loss: 1.4139 - val_accuracy: 0.4860

Epoch 16/20
141/141 [=====] - 1s 6ms/step - loss: 1.2022 - accuracy: 0.5924 - val_loss: 1.4026 - val_accuracy: 0.4920

Epoch 17/20
141/141 [=====] - 1s 5ms/step - loss: 1.1799 - accuracy: 0.6020 - val_loss: 1.3903 - val_accuracy: 0.5040

Epoch 18/20
141/141 [=====] - 1s 5ms/step - loss: 1.1586 - accuracy: 0.6102 - val_loss: 1.3763 - val_accuracy: 0.5040

Epoch 19/20
141/141 [=====] - 1s 5ms/step - loss: 1.1381 - accuracy: 0.6187 - val_loss: 1.3666 - val_accuracy: 0.5140

Epoch 20/20
141/141 [=====] - 1s 5ms/step - loss: 1.1185 - accuracy: 0.6267 - val_loss: 1.3551 - val_accuracy: 0.5120

Training model with num_blocks=3, use_skip=True, and learning_rate=0.001

Epoch 1/20
141/141 [=====] - 3s 8ms/step - loss: 1.7995 - accuracy: 0.3520 - val_loss: 2.7965 - val_accuracy: 0.1380

Epoch 2/20
141/141 [=====] - 1s 7ms/step - loss: 1.3253 - accuracy: 0.5176 - val_loss: 3.0005 - val_accuracy: 0.2200

Epoch 3/20
141/141 [=====] - 1s 8ms/step - loss: 1.1307 - accuracy: 0.5900 - val_loss: 2.1731 - val_accuracy: 0.3180

Epoch 4/20
141/141 [=====] - 1s 9ms/step - loss: 0.9842 - accuracy: 0.6544 - val_loss: 1.5053 - val_accuracy: 0.4600

Epoch 5/20
141/141 [=====] - 1s 7ms/step - loss: 0.8627 - accuracy: 0.7124 - val_loss: 1.3924 - val_accuracy: 0.4860

Epoch 6/20
141/141 [=====] - 1s 7ms/step - loss: 0.7649 - accuracy: 0.7527 - val_loss: 1.5152 - val_accuracy: 0.4720

Epoch 7/20
141/141 [=====] - 1s 7ms/step - loss: 0.6734 - accuracy: 0.7929 - val_loss: 1.6826 - val_accuracy: 0.4540

Epoch 8/20
141/141 [=====] - 1s 7ms/step - loss: 0.5861 -

```

accuracy: 0.8293 - val_loss: 1.6894 - val_accuracy: 0.4420
Epoch 9/20
141/141 [=====] - 1s 7ms/step - loss: 0.5076 -
accuracy: 0.8687 - val_loss: 1.5566 - val_accuracy: 0.4820
Epoch 10/20
141/141 [=====] - 1s 7ms/step - loss: 0.4415 -
accuracy: 0.8998 - val_loss: 1.5980 - val_accuracy: 0.4820
Epoch 11/20
141/141 [=====] - 1s 7ms/step - loss: 0.3833 -
accuracy: 0.9176 - val_loss: 1.5905 - val_accuracy: 0.4900
Epoch 12/20
141/141 [=====] - 1s 7ms/step - loss: 0.3310 -
accuracy: 0.9424 - val_loss: 1.4653 - val_accuracy: 0.5420
Epoch 13/20
141/141 [=====] - 1s 9ms/step - loss: 0.2898 -
accuracy: 0.9573 - val_loss: 1.4417 - val_accuracy: 0.5500
Epoch 14/20
141/141 [=====] - 1s 7ms/step - loss: 0.2537 -
accuracy: 0.9678 - val_loss: 1.4241 - val_accuracy: 0.5500
Epoch 15/20
141/141 [=====] - 1s 7ms/step - loss: 0.2234 -
accuracy: 0.9760 - val_loss: 1.5265 - val_accuracy: 0.5420
Epoch 16/20
141/141 [=====] - 1s 7ms/step - loss: 0.1947 -
accuracy: 0.9822 - val_loss: 1.6417 - val_accuracy: 0.5260
Epoch 17/20
141/141 [=====] - 1s 7ms/step - loss: 0.1692 -
accuracy: 0.9878 - val_loss: 1.6998 - val_accuracy: 0.5100
Epoch 18/20
141/141 [=====] - 1s 7ms/step - loss: 0.1464 -
accuracy: 0.9936 - val_loss: 1.4843 - val_accuracy: 0.5500
Epoch 19/20
141/141 [=====] - 1s 7ms/step - loss: 0.1236 -
accuracy: 0.9969 - val_loss: 1.3803 - val_accuracy: 0.5860
Epoch 20/20
141/141 [=====] - 1s 7ms/step - loss: 0.1045 -
accuracy: 0.9980 - val_loss: 1.3364 - val_accuracy: 0.6160
Training model with num_blocks=3, use_skip=True, and learning_rate=0.0001
Epoch 1/20
141/141 [=====] - 4s 13ms/step - loss: 2.2147 -
accuracy: 0.2167 - val_loss: 2.3670 - val_accuracy: 0.1000
Epoch 2/20
141/141 [=====] - 1s 7ms/step - loss: 1.7773 -
accuracy: 0.3569 - val_loss: 2.3118 - val_accuracy: 0.1840
Epoch 3/20
141/141 [=====] - 1s 7ms/step - loss: 1.6247 -
accuracy: 0.4164 - val_loss: 1.9523 - val_accuracy: 0.2760
Epoch 4/20

```

141/141 [=====] - 1s 7ms/step - loss: 1.5251 -
accuracy: 0.4567 - val_loss: 1.6289 - val_accuracy: 0.3940
Epoch 5/20
141/141 [=====] - 1s 7ms/step - loss: 1.4491 -
accuracy: 0.4842 - val_loss: 1.5233 - val_accuracy: 0.4380
Epoch 6/20
141/141 [=====] - 1s 7ms/step - loss: 1.3862 -
accuracy: 0.5062 - val_loss: 1.4841 - val_accuracy: 0.4440
Epoch 7/20
141/141 [=====] - 1s 7ms/step - loss: 1.3328 -
accuracy: 0.5289 - val_loss: 1.4552 - val_accuracy: 0.4420
Epoch 8/20
141/141 [=====] - 1s 7ms/step - loss: 1.2861 -
accuracy: 0.5473 - val_loss: 1.4340 - val_accuracy: 0.4440
Epoch 9/20
141/141 [=====] - 1s 7ms/step - loss: 1.2443 -
accuracy: 0.5618 - val_loss: 1.4127 - val_accuracy: 0.4740
Epoch 10/20
141/141 [=====] - 1s 7ms/step - loss: 1.2066 -
accuracy: 0.5747 - val_loss: 1.3934 - val_accuracy: 0.4780
Epoch 11/20
141/141 [=====] - 1s 10ms/step - loss: 1.1719 -
accuracy: 0.5896 - val_loss: 1.3751 - val_accuracy: 0.4820
Epoch 12/20
141/141 [=====] - 1s 7ms/step - loss: 1.1399 -
accuracy: 0.6062 - val_loss: 1.3582 - val_accuracy: 0.4880
Epoch 13/20
141/141 [=====] - 1s 7ms/step - loss: 1.1103 -
accuracy: 0.6189 - val_loss: 1.3473 - val_accuracy: 0.4900
Epoch 14/20
141/141 [=====] - 1s 7ms/step - loss: 1.0821 -
accuracy: 0.6353 - val_loss: 1.3325 - val_accuracy: 0.4960
Epoch 15/20
141/141 [=====] - 1s 7ms/step - loss: 1.0557 -
accuracy: 0.6473 - val_loss: 1.3231 - val_accuracy: 0.5040
Epoch 16/20
141/141 [=====] - 1s 7ms/step - loss: 1.0307 -
accuracy: 0.6578 - val_loss: 1.3144 - val_accuracy: 0.5000
Epoch 17/20
141/141 [=====] - 1s 7ms/step - loss: 1.0068 -
accuracy: 0.6702 - val_loss: 1.3044 - val_accuracy: 0.5140
Epoch 18/20
141/141 [=====] - 1s 7ms/step - loss: 0.9839 -
accuracy: 0.6796 - val_loss: 1.2964 - val_accuracy: 0.5140
Epoch 19/20
141/141 [=====] - 1s 7ms/step - loss: 0.9617 -
accuracy: 0.6900 - val_loss: 1.2892 - val_accuracy: 0.5220
Epoch 20/20

141/141 [=====] - 1s 9ms/step - loss: 0.9404 -
accuracy: 0.7013 - val_loss: 1.2806 - val_accuracy: 0.5240
Training model with num_blocks=3, use_skip=False, and learning_rate=0.001
Epoch 1/20
141/141 [=====] - 3s 9ms/step - loss: 1.7919 -
accuracy: 0.3500 - val_loss: 2.6187 - val_accuracy: 0.1320
Epoch 2/20
141/141 [=====] - 1s 7ms/step - loss: 1.3576 -
accuracy: 0.5151 - val_loss: 2.5274 - val_accuracy: 0.1800
Epoch 3/20
141/141 [=====] - 1s 7ms/step - loss: 1.1632 -
accuracy: 0.5942 - val_loss: 1.7175 - val_accuracy: 0.3780
Epoch 4/20
141/141 [=====] - 1s 7ms/step - loss: 1.0162 -
accuracy: 0.6573 - val_loss: 1.4810 - val_accuracy: 0.4860
Epoch 5/20
141/141 [=====] - 1s 7ms/step - loss: 0.8937 -
accuracy: 0.7164 - val_loss: 1.6572 - val_accuracy: 0.4440
Epoch 6/20
141/141 [=====] - 1s 9ms/step - loss: 0.7865 -
accuracy: 0.7638 - val_loss: 1.6164 - val_accuracy: 0.4520
Epoch 7/20
141/141 [=====] - 1s 9ms/step - loss: 0.6920 -
accuracy: 0.8033 - val_loss: 1.5467 - val_accuracy: 0.4760
Epoch 8/20
141/141 [=====] - 1s 9ms/step - loss: 0.6186 -
accuracy: 0.8382 - val_loss: 1.6237 - val_accuracy: 0.4760
Epoch 9/20
141/141 [=====] - 1s 7ms/step - loss: 0.5561 -
accuracy: 0.8609 - val_loss: 1.4653 - val_accuracy: 0.4760
Epoch 10/20
141/141 [=====] - 1s 7ms/step - loss: 0.5024 -
accuracy: 0.8802 - val_loss: 1.5936 - val_accuracy: 0.4720
Epoch 11/20
141/141 [=====] - 1s 7ms/step - loss: 0.4391 -
accuracy: 0.9107 - val_loss: 1.5274 - val_accuracy: 0.4860
Epoch 12/20
141/141 [=====] - 1s 7ms/step - loss: 0.3817 -
accuracy: 0.9327 - val_loss: 1.3972 - val_accuracy: 0.5040
Epoch 13/20
141/141 [=====] - 1s 7ms/step - loss: 0.3332 -
accuracy: 0.9489 - val_loss: 1.3769 - val_accuracy: 0.5280
Epoch 14/20
141/141 [=====] - 1s 7ms/step - loss: 0.2918 -
accuracy: 0.9633 - val_loss: 1.4067 - val_accuracy: 0.5020
Epoch 15/20
141/141 [=====] - 1s 7ms/step - loss: 0.2529 -
accuracy: 0.9769 - val_loss: 1.5056 - val_accuracy: 0.4800

Epoch 16/20
141/141 [=====] - 1s 7ms/step - loss: 0.2209 - accuracy: 0.9842 - val_loss: 1.6467 - val_accuracy: 0.4760
Epoch 17/20
141/141 [=====] - 1s 8ms/step - loss: 0.1952 - accuracy: 0.9900 - val_loss: 1.6863 - val_accuracy: 0.4640
Epoch 18/20
141/141 [=====] - 1s 9ms/step - loss: 0.1732 - accuracy: 0.9927 - val_loss: 1.5936 - val_accuracy: 0.4860
Epoch 19/20
141/141 [=====] - 1s 8ms/step - loss: 0.1512 - accuracy: 0.9964 - val_loss: 1.5589 - val_accuracy: 0.4860
Epoch 20/20
141/141 [=====] - 1s 7ms/step - loss: 0.1317 - accuracy: 0.9982 - val_loss: 1.5157 - val_accuracy: 0.5160
Training model with num_blocks=3, use_skip=False, and learning_rate=0.0001
Epoch 1/20
141/141 [=====] - 3s 9ms/step - loss: 2.1465 - accuracy: 0.2236 - val_loss: 2.3016 - val_accuracy: 0.1040
Epoch 2/20
141/141 [=====] - 1s 7ms/step - loss: 1.8046 - accuracy: 0.3469 - val_loss: 2.2845 - val_accuracy: 0.1660
Epoch 3/20
141/141 [=====] - 1s 7ms/step - loss: 1.6620 - accuracy: 0.3976 - val_loss: 1.9542 - val_accuracy: 0.2960
Epoch 4/20
141/141 [=====] - 1s 7ms/step - loss: 1.5682 - accuracy: 0.4367 - val_loss: 1.6419 - val_accuracy: 0.4020
Epoch 5/20
141/141 [=====] - 1s 9ms/step - loss: 1.4970 - accuracy: 0.4664 - val_loss: 1.5449 - val_accuracy: 0.4400
Epoch 6/20
141/141 [=====] - 1s 10ms/step - loss: 1.4399 - accuracy: 0.4949 - val_loss: 1.4979 - val_accuracy: 0.4600
Epoch 7/20
141/141 [=====] - 1s 7ms/step - loss: 1.3920 - accuracy: 0.5160 - val_loss: 1.4633 - val_accuracy: 0.4740
Epoch 8/20
141/141 [=====] - 1s 7ms/step - loss: 1.3503 - accuracy: 0.5358 - val_loss: 1.4373 - val_accuracy: 0.4800
Epoch 9/20
141/141 [=====] - 1s 7ms/step - loss: 1.3128 - accuracy: 0.5533 - val_loss: 1.4142 - val_accuracy: 0.4880
Epoch 10/20
141/141 [=====] - 1s 7ms/step - loss: 1.2784 - accuracy: 0.5702 - val_loss: 1.3931 - val_accuracy: 0.5000
Epoch 11/20
141/141 [=====] - 1s 7ms/step - loss: 1.2467 -

```

accuracy: 0.5796 - val_loss: 1.3750 - val_accuracy: 0.5020
Epoch 12/20
141/141 [=====] - 1s 7ms/step - loss: 1.2171 -
accuracy: 0.5944 - val_loss: 1.3591 - val_accuracy: 0.5040
Epoch 13/20
141/141 [=====] - 1s 7ms/step - loss: 1.1890 -
accuracy: 0.6016 - val_loss: 1.3464 - val_accuracy: 0.5120
Epoch 14/20
141/141 [=====] - 1s 8ms/step - loss: 1.1625 -
accuracy: 0.6118 - val_loss: 1.3360 - val_accuracy: 0.5160
Epoch 15/20
141/141 [=====] - 1s 9ms/step - loss: 1.1370 -
accuracy: 0.6256 - val_loss: 1.3234 - val_accuracy: 0.5200
Epoch 16/20
141/141 [=====] - 1s 7ms/step - loss: 1.1127 -
accuracy: 0.6371 - val_loss: 1.3119 - val_accuracy: 0.5200
Epoch 17/20
141/141 [=====] - 1s 7ms/step - loss: 1.0893 -
accuracy: 0.6502 - val_loss: 1.3022 - val_accuracy: 0.5200
Epoch 18/20
141/141 [=====] - 1s 7ms/step - loss: 1.0668 -
accuracy: 0.6624 - val_loss: 1.2934 - val_accuracy: 0.5260
Epoch 19/20
141/141 [=====] - 1s 7ms/step - loss: 1.0448 -
accuracy: 0.6740 - val_loss: 1.2837 - val_accuracy: 0.5280
Epoch 20/20
141/141 [=====] - 1s 7ms/step - loss: 1.0234 -
accuracy: 0.6851 - val_loss: 1.2750 - val_accuracy: 0.5320
Training model with num_blocks=4, use_skip=True, and learning_rate=0.001
Epoch 1/20
141/141 [=====] - 5s 15ms/step - loss: 1.8263 -
accuracy: 0.3364 - val_loss: 2.8242 - val_accuracy: 0.1120
Epoch 2/20
141/141 [=====] - 1s 9ms/step - loss: 1.3530 -
accuracy: 0.5000 - val_loss: 3.0550 - val_accuracy: 0.1300
Epoch 3/20
141/141 [=====] - 1s 8ms/step - loss: 1.1058 -
accuracy: 0.6002 - val_loss: 2.0405 - val_accuracy: 0.3000
Epoch 4/20
141/141 [=====] - 1s 9ms/step - loss: 0.9042 -
accuracy: 0.6938 - val_loss: 1.6576 - val_accuracy: 0.4500
Epoch 5/20
141/141 [=====] - 1s 9ms/step - loss: 0.7417 -
accuracy: 0.7649 - val_loss: 1.5670 - val_accuracy: 0.4640
Epoch 6/20
141/141 [=====] - 1s 8ms/step - loss: 0.6125 -
accuracy: 0.8229 - val_loss: 1.5563 - val_accuracy: 0.4700
Epoch 7/20

```

```

141/141 [=====] - 1s 8ms/step - loss: 0.5179 -
accuracy: 0.8671 - val_loss: 2.4837 - val_accuracy: 0.3900
Epoch 8/20
141/141 [=====] - 1s 9ms/step - loss: 0.4439 -
accuracy: 0.8922 - val_loss: 3.8181 - val_accuracy: 0.3420
Epoch 9/20
141/141 [=====] - 1s 8ms/step - loss: 0.3611 -
accuracy: 0.9247 - val_loss: 3.1617 - val_accuracy: 0.3700
Epoch 10/20
141/141 [=====] - 1s 8ms/step - loss: 0.2806 -
accuracy: 0.9524 - val_loss: 2.5320 - val_accuracy: 0.4100
Epoch 11/20
141/141 [=====] - 1s 8ms/step - loss: 0.2236 -
accuracy: 0.9689 - val_loss: 1.9231 - val_accuracy: 0.4640
Epoch 12/20
141/141 [=====] - 1s 8ms/step - loss: 0.1874 -
accuracy: 0.9756 - val_loss: 1.7596 - val_accuracy: 0.5040
Epoch 13/20
141/141 [=====] - 1s 8ms/step - loss: 0.1607 -
accuracy: 0.9793 - val_loss: 1.7542 - val_accuracy: 0.4820
Epoch 14/20
141/141 [=====] - 1s 8ms/step - loss: 0.1397 -
accuracy: 0.9847 - val_loss: 2.0238 - val_accuracy: 0.4840
Epoch 15/20
141/141 [=====] - 1s 8ms/step - loss: 0.1143 -
accuracy: 0.9880 - val_loss: 1.7279 - val_accuracy: 0.5160
Epoch 16/20
141/141 [=====] - 1s 8ms/step - loss: 0.0941 -
accuracy: 0.9909 - val_loss: 1.7014 - val_accuracy: 0.5300
Epoch 17/20
141/141 [=====] - 1s 10ms/step - loss: 0.0737 -
accuracy: 0.9949 - val_loss: 1.5870 - val_accuracy: 0.5700
Epoch 18/20
141/141 [=====] - 1s 9ms/step - loss: 0.0539 -
accuracy: 0.9971 - val_loss: 1.4798 - val_accuracy: 0.5960
Epoch 19/20
141/141 [=====] - 1s 8ms/step - loss: 0.0398 -
accuracy: 0.9991 - val_loss: 1.5578 - val_accuracy: 0.5820
Epoch 20/20
141/141 [=====] - 1s 8ms/step - loss: 0.0286 -
accuracy: 0.9993 - val_loss: 1.6076 - val_accuracy: 0.5780
Training model with num_blocks=4, use_skip=True, and learning_rate=0.0001
Epoch 1/20
141/141 [=====] - 4s 10ms/step - loss: 2.0299 -
accuracy: 0.2631 - val_loss: 2.4121 - val_accuracy: 0.1100
Epoch 2/20
141/141 [=====] - 1s 8ms/step - loss: 1.6729 -
accuracy: 0.3880 - val_loss: 2.5415 - val_accuracy: 0.1120

```

Epoch 3/20
141/141 [=====] - 1s 10ms/step - loss: 1.5266 -
accuracy: 0.4462 - val_loss: 1.9680 - val_accuracy: 0.2860
Epoch 4/20
141/141 [=====] - 1s 10ms/step - loss: 1.4239 -
accuracy: 0.4851 - val_loss: 1.5731 - val_accuracy: 0.4340
Epoch 5/20
141/141 [=====] - 1s 8ms/step - loss: 1.3428 -
accuracy: 0.5171 - val_loss: 1.4764 - val_accuracy: 0.4540
Epoch 6/20
141/141 [=====] - 1s 8ms/step - loss: 1.2749 -
accuracy: 0.5538 - val_loss: 1.4382 - val_accuracy: 0.4760
Epoch 7/20
141/141 [=====] - 1s 8ms/step - loss: 1.2159 -
accuracy: 0.5800 - val_loss: 1.4089 - val_accuracy: 0.4940
Epoch 8/20
141/141 [=====] - 1s 9ms/step - loss: 1.1629 -
accuracy: 0.5991 - val_loss: 1.3883 - val_accuracy: 0.5000
Epoch 9/20
141/141 [=====] - 1s 8ms/step - loss: 1.1132 -
accuracy: 0.6204 - val_loss: 1.3682 - val_accuracy: 0.5100
Epoch 10/20
141/141 [=====] - 1s 8ms/step - loss: 1.0674 -
accuracy: 0.6373 - val_loss: 1.3468 - val_accuracy: 0.5220
Epoch 11/20
141/141 [=====] - 1s 8ms/step - loss: 1.0244 -
accuracy: 0.6587 - val_loss: 1.3303 - val_accuracy: 0.5220
Epoch 12/20
141/141 [=====] - 2s 11ms/step - loss: 0.9834 -
accuracy: 0.6767 - val_loss: 1.3169 - val_accuracy: 0.5160
Epoch 13/20
141/141 [=====] - 1s 8ms/step - loss: 0.9445 -
accuracy: 0.6931 - val_loss: 1.2995 - val_accuracy: 0.5240
Epoch 14/20
141/141 [=====] - 1s 8ms/step - loss: 0.9074 -
accuracy: 0.7107 - val_loss: 1.2881 - val_accuracy: 0.5320
Epoch 15/20
141/141 [=====] - 1s 8ms/step - loss: 0.8717 -
accuracy: 0.7269 - val_loss: 1.2732 - val_accuracy: 0.5420
Epoch 16/20
141/141 [=====] - 1s 8ms/step - loss: 0.8374 -
accuracy: 0.7478 - val_loss: 1.2640 - val_accuracy: 0.5420
Epoch 17/20
141/141 [=====] - 1s 8ms/step - loss: 0.8041 -
accuracy: 0.7622 - val_loss: 1.2550 - val_accuracy: 0.5380
Epoch 18/20
141/141 [=====] - 1s 8ms/step - loss: 0.7721 -
accuracy: 0.7778 - val_loss: 1.2504 - val_accuracy: 0.5480

Epoch 19/20
141/141 [=====] - 1s 8ms/step - loss: 0.7413 - accuracy: 0.7953 - val_loss: 1.2438 - val_accuracy: 0.5460

Epoch 20/20
141/141 [=====] - 1s 9ms/step - loss: 0.7113 - accuracy: 0.8122 - val_loss: 1.2378 - val_accuracy: 0.5520
Training model with num_blocks=4, use_skip=False, and learning_rate=0.001

Epoch 1/20
141/141 [=====] - 4s 10ms/step - loss: 1.7467 - accuracy: 0.3518 - val_loss: 2.4245 - val_accuracy: 0.1060

Epoch 2/20
141/141 [=====] - 1s 9ms/step - loss: 1.3043 - accuracy: 0.5218 - val_loss: 2.8335 - val_accuracy: 0.1520

Epoch 3/20
141/141 [=====] - 1s 8ms/step - loss: 1.0767 - accuracy: 0.6238 - val_loss: 2.1762 - val_accuracy: 0.2960

Epoch 4/20
141/141 [=====] - 1s 8ms/step - loss: 0.9019 - accuracy: 0.6996 - val_loss: 1.5227 - val_accuracy: 0.4480

Epoch 5/20
141/141 [=====] - 1s 8ms/step - loss: 0.7540 - accuracy: 0.7680 - val_loss: 1.4505 - val_accuracy: 0.4560

Epoch 6/20
141/141 [=====] - 1s 8ms/step - loss: 0.6320 - accuracy: 0.8216 - val_loss: 1.4151 - val_accuracy: 0.4960

Epoch 7/20
141/141 [=====] - 1s 9ms/step - loss: 0.5225 - accuracy: 0.8653 - val_loss: 1.5505 - val_accuracy: 0.4680

Epoch 8/20
141/141 [=====] - 1s 10ms/step - loss: 0.4262 - accuracy: 0.9033 - val_loss: 1.5769 - val_accuracy: 0.4700

Epoch 9/20
141/141 [=====] - 1s 10ms/step - loss: 0.3423 - accuracy: 0.9351 - val_loss: 1.5709 - val_accuracy: 0.4920

Epoch 10/20
141/141 [=====] - 1s 8ms/step - loss: 0.2663 - accuracy: 0.9640 - val_loss: 2.3646 - val_accuracy: 0.3600

Epoch 11/20
141/141 [=====] - 1s 8ms/step - loss: 0.2170 - accuracy: 0.9756 - val_loss: 2.5921 - val_accuracy: 0.3280

Epoch 12/20
141/141 [=====] - 1s 8ms/step - loss: 0.1750 - accuracy: 0.9833 - val_loss: 2.0552 - val_accuracy: 0.4300

Epoch 13/20
141/141 [=====] - 1s 8ms/step - loss: 0.1356 - accuracy: 0.9907 - val_loss: 1.7173 - val_accuracy: 0.5140

Epoch 14/20
141/141 [=====] - 1s 8ms/step - loss: 0.0994 -

```

accuracy: 0.9953 - val_loss: 1.6295 - val_accuracy: 0.5240
Epoch 15/20
141/141 [=====] - 1s 8ms/step - loss: 0.0753 -
accuracy: 0.9980 - val_loss: 1.7714 - val_accuracy: 0.5240
Epoch 16/20
141/141 [=====] - 1s 8ms/step - loss: 0.0542 -
accuracy: 0.9996 - val_loss: 1.7023 - val_accuracy: 0.5520
Epoch 17/20
141/141 [=====] - 1s 8ms/step - loss: 0.0400 -
accuracy: 0.9996 - val_loss: 1.8312 - val_accuracy: 0.5340
Epoch 18/20
141/141 [=====] - 2s 11ms/step - loss: 0.0289 -
accuracy: 1.0000 - val_loss: 1.6550 - val_accuracy: 0.5600
Epoch 19/20
141/141 [=====] - 1s 8ms/step - loss: 0.0212 -
accuracy: 1.0000 - val_loss: 1.5509 - val_accuracy: 0.5660
Epoch 20/20
141/141 [=====] - 1s 8ms/step - loss: 0.0167 -
accuracy: 1.0000 - val_loss: 1.5023 - val_accuracy: 0.5720
Training model with num_blocks=4, use_skip=False, and learning_rate=0.0001
Epoch 1/20
141/141 [=====] - 4s 11ms/step - loss: 2.1205 -
accuracy: 0.2304 - val_loss: 2.3534 - val_accuracy: 0.0880
Epoch 2/20
141/141 [=====] - 1s 8ms/step - loss: 1.7792 -
accuracy: 0.3682 - val_loss: 2.4094 - val_accuracy: 0.1020
Epoch 3/20
141/141 [=====] - 1s 8ms/step - loss: 1.6251 -
accuracy: 0.4247 - val_loss: 2.0627 - val_accuracy: 0.2360
Epoch 4/20
141/141 [=====] - 1s 10ms/step - loss: 1.5212 -
accuracy: 0.4638 - val_loss: 1.7005 - val_accuracy: 0.3820
Epoch 5/20
141/141 [=====] - 2s 12ms/step - loss: 1.4425 -
accuracy: 0.4956 - val_loss: 1.5809 - val_accuracy: 0.4340
Epoch 6/20
141/141 [=====] - 1s 8ms/step - loss: 1.3778 -
accuracy: 0.5220 - val_loss: 1.5305 - val_accuracy: 0.4460
Epoch 7/20
141/141 [=====] - 1s 8ms/step - loss: 1.3219 -
accuracy: 0.5456 - val_loss: 1.4987 - val_accuracy: 0.4600
Epoch 8/20
141/141 [=====] - 1s 8ms/step - loss: 1.2712 -
accuracy: 0.5691 - val_loss: 1.4723 - val_accuracy: 0.4620
Epoch 9/20
141/141 [=====] - 1s 9ms/step - loss: 1.2244 -
accuracy: 0.5940 - val_loss: 1.4470 - val_accuracy: 0.4720
Epoch 10/20

```

```

141/141 [=====] - 1s 8ms/step - loss: 1.1808 -
accuracy: 0.6116 - val_loss: 1.4280 - val_accuracy: 0.4760
Epoch 11/20
141/141 [=====] - 1s 8ms/step - loss: 1.1394 -
accuracy: 0.6302 - val_loss: 1.4109 - val_accuracy: 0.4800
Epoch 12/20
141/141 [=====] - 1s 9ms/step - loss: 1.1003 -
accuracy: 0.6480 - val_loss: 1.3933 - val_accuracy: 0.4860
Epoch 13/20
141/141 [=====] - 2s 11ms/step - loss: 1.0629 -
accuracy: 0.6633 - val_loss: 1.3827 - val_accuracy: 0.4920
Epoch 14/20
141/141 [=====] - 2s 11ms/step - loss: 1.0269 -
accuracy: 0.6776 - val_loss: 1.3685 - val_accuracy: 0.4900
Epoch 15/20
141/141 [=====] - 1s 8ms/step - loss: 0.9922 -
accuracy: 0.6960 - val_loss: 1.3606 - val_accuracy: 0.4900
Epoch 16/20
141/141 [=====] - 1s 9ms/step - loss: 0.9585 -
accuracy: 0.7129 - val_loss: 1.3465 - val_accuracy: 0.4940
Epoch 17/20
141/141 [=====] - 1s 9ms/step - loss: 0.9259 -
accuracy: 0.7287 - val_loss: 1.3301 - val_accuracy: 0.5000
Epoch 18/20
141/141 [=====] - 1s 9ms/step - loss: 0.8944 -
accuracy: 0.7433 - val_loss: 1.3156 - val_accuracy: 0.5040
Epoch 19/20
141/141 [=====] - 1s 9ms/step - loss: 0.8640 -
accuracy: 0.7558 - val_loss: 1.3010 - val_accuracy: 0.5220
Epoch 20/20
141/141 [=====] - 1s 9ms/step - loss: 0.8345 -
accuracy: 0.7687 - val_loss: 1.2893 - val_accuracy: 0.5300
Best validation accuracy is 0.615999996621399 with params: {'num_blocks': 3,
'use_skip': True, 'learning_rate': 0.001}

Result:      The best validation accuracy is 0.598 with num_blocks=4,
use_skip=true,learning_rate=0.001

```

comments: the next time I run the same code, the result changes to: num_block=3,use_skip=True,Learning_rate=0.001. validation accuracy becomes 0.616

It is noticed that the overfitting problem has been solved a bit

0.1.5 Question 3.5: Apply data augmentation

We now try to apply data augmentation to improve the performance. Extend the code of the class `YourModel` so that if the attribute `is_augmentation` is set to `True`, we apply the data augmentation. Also you need to incorporate early stopping to your training process. Specifically, you early stop the training if the valid accuracy cannot increase in three consecutive epochs.

[4 points]

```
[ ]: from tensorflow.keras.preprocessing.image import ImageDataGenerator
```

Write your code in the cell below. Hint that you can rewrite the code of the `fit` method to apply the data augmentation. In addition, you can copy the code of `build_cnn` method above to reuse here.

```
[ ]: import numpy as np
class YourModel(BaseImageClassifier):
    def __init__(self,
                  name='network1',
                  width=32, height=32, depth=3,
                  num_blocks=2,
                  feature_maps=32,
                  num_classes=4,
                  drop_rate=0.2,
                  batch_norm = None,
                  is_augmentation = False,
                  activation_func='relu',
                  use_skip = True,
                  optimizer='adam',
                  batch_size=10,
                  num_epochs= 20,
                  learning_rate=0.0001):
        super(YourModel, self).__init__(name, width, height, depth, num_blocks,
        ↪feature_maps, num_classes, drop_rate, batch_norm, is_augmentation,
        ↪activation_func, use_skip, optimizer,
        ↪batch_size, num_epochs, learning_rate)

    def custom_block(self, x, filters):
        # skip connection
        shortcut = x

        # First Conv layer
        x = layers.Conv2D(filters=filters, kernel_size=(3, 3), strides=(1, 1),
        ↪padding="same")(x)
        if self.batch_norm:
            x = layers.BatchNormalization()(x)
        x = layers.Activation(self.activation_func)(x)

        # Second Conv layer
        x = layers.Conv2D(filters=filters, kernel_size=(3, 3), strides=(1, 1),
        ↪padding="same")(x)
        if self.batch_norm:
            x = layers.BatchNormalization()(x)

        # Add the shortcut to the output
```

```

        if self.use_skip:
            x = tf.keras.layers.add([shortcut, x])
        x = layers.Activation(self.activation_func)(x)

        # Mean Pooling layer
        x = layers.AveragePooling2D(pool_size=(2, 2), strides=(2, 2),
↪padding="same")(x)

        # Dropout
        x = layers.Dropout(rate=self.drop_rate)(x)

    return x

def build_cnn(self):
    input_tensor = layers.Input(shape=(self.height, self.width, self.depth))
    x = input_tensor

    for current_feature_maps in self.feature_maps:
        # Initial convolution
        x = layers.Conv2D(current_feature_maps, (3, 3), padding='same',
↪activation=None)(x)
        if self.batch_norm:
            x = layers.BatchNormalization()(x)
        x = layers.Activation(self.activation_func)(x)

        # Second convolution with skip connection
        shortcut = x
        x = layers.Conv2D(current_feature_maps, (3, 3), padding='same',
↪activation=None)(x)
        if self.batch_norm:
            x = layers.BatchNormalization()(x)
        if self.use_skip:
            x = layers.Add()([shortcut, x])
        x = layers.Activation(self.activation_func)(x)

        # Mean Pooling and Dropout
        x = layers.AveragePooling2D(pool_size=(2, 2), padding='same')(x)
        if self.drop_rate > 0:
            x = layers.Dropout(self.drop_rate)(x)

    x = layers.Flatten()(x)
    x = layers.Dense(self.num_classes, activation='softmax')(x)

    self.model = models.Model(inputs=input_tensor, outputs=x)
    self.model.compile(optimizer=self.optimizer,
↪loss='sparse_categorical_crossentropy', metrics=['accuracy'])

```

```

def fit(self, data_manager, batch_size=None, num_epochs=None):
    batch_size = self.batch_size if batch_size is None else batch_size
    num_epochs = self.num_epochs if num_epochs is None else num_epochs

    x_train_batch = self.optimize_data_pipeline(data_manager.ds_train,
    ↪batch_size=batch_size)
    x_val_batch = self.optimize_data_pipeline(data_manager.ds_val,
    ↪batch_size=batch_size)

    if self.is_augmentation: #added code for data augmentation
        train_data_list, train_labels_list = [], []
        for data, labels in x_train_batch:
            train_data_list.append(data.numpy())
            train_labels_list.append(labels.numpy())

        train_data = np.concatenate(train_data_list, axis=0)
        train_labels = np.concatenate(train_labels_list, axis=0)

        datagen = ImageDataGenerator(rotation_range=10,
                                     width_shift_range=0.1,
                                     height_shift_range=0.1,
                                     horizontal_flip=True) #data augmentation

        datagen.fit(train_data)

        self.model.fit(datagen.flow(train_data, train_labels, batch_size=self.
    ↪batch_size),
                        validation_data=x_val_batch,
                        epochs=self.num_epochs,
                        verbose=self.verbose)
    else:
        self.model.fit(x_train_batch, validation_data=x_val_batch, epochs=self.
    ↪num_epochs, batch_size=self.batch_size, verbose=self.verbose)

```

```

[ ]: model_3_5=YourModel(name='q3_5',
                        feature_maps=32,
                        num_classes=data_manager.n_classes,
                        num_blocks=4,
                        drop_rate=0.0,
                        batch_norm=True,
                        is_augmentation=True,
                        use_skip=True,
                        optimizer='sgd',
                        learning_rate=0.001)

```

```
model_3_5.build_cnn()  
model_3_5.fit(data_manager)
```

```
Epoch 1/20  
450/450 [=====] - 9s 12ms/step - loss: 2.0184 -  
accuracy: 0.3062 - val_loss: 1.7373 - val_accuracy: 0.3420  
Epoch 2/20  
450/450 [=====] - 6s 14ms/step - loss: 1.6659 -  
accuracy: 0.4024 - val_loss: 1.5998 - val_accuracy: 0.3980  
Epoch 3/20  
450/450 [=====] - 5s 11ms/step - loss: 1.5359 -  
accuracy: 0.4438 - val_loss: 1.6073 - val_accuracy: 0.3860  
Epoch 4/20  
450/450 [=====] - 6s 12ms/step - loss: 1.4367 -  
accuracy: 0.4691 - val_loss: 1.4128 - val_accuracy: 0.4860  
Epoch 5/20  
450/450 [=====] - 6s 12ms/step - loss: 1.3578 -  
accuracy: 0.5047 - val_loss: 1.3745 - val_accuracy: 0.4980  
Epoch 6/20  
450/450 [=====] - 5s 12ms/step - loss: 1.2977 -  
accuracy: 0.5216 - val_loss: 1.7192 - val_accuracy: 0.4480  
Epoch 7/20  
450/450 [=====] - 6s 14ms/step - loss: 1.2417 -  
accuracy: 0.5478 - val_loss: 1.1789 - val_accuracy: 0.5520  
Epoch 8/20  
450/450 [=====] - 5s 11ms/step - loss: 1.1722 -  
accuracy: 0.5644 - val_loss: 1.5165 - val_accuracy: 0.5140  
Epoch 9/20  
450/450 [=====] - 5s 11ms/step - loss: 1.1578 -  
accuracy: 0.5718 - val_loss: 1.5003 - val_accuracy: 0.5200  
Epoch 10/20  
450/450 [=====] - 6s 13ms/step - loss: 1.0827 -  
accuracy: 0.6007 - val_loss: 1.5834 - val_accuracy: 0.4760  
Epoch 11/20  
450/450 [=====] - 5s 12ms/step - loss: 1.0400 -  
accuracy: 0.6156 - val_loss: 1.1756 - val_accuracy: 0.5620  
Epoch 12/20  
450/450 [=====] - 6s 14ms/step - loss: 1.0200 -  
accuracy: 0.6300 - val_loss: 1.1522 - val_accuracy: 0.6080  
Epoch 13/20  
450/450 [=====] - 5s 11ms/step - loss: 0.9788 -  
accuracy: 0.6467 - val_loss: 1.1999 - val_accuracy: 0.5800  
Epoch 14/20  
450/450 [=====] - 5s 12ms/step - loss: 0.9303 -  
accuracy: 0.6613 - val_loss: 1.2312 - val_accuracy: 0.5840  
Epoch 15/20
```

```

450/450 [=====] - 5s 11ms/step - loss: 0.9233 -
accuracy: 0.6636 - val_loss: 1.8355 - val_accuracy: 0.4180
Epoch 16/20
450/450 [=====] - 6s 13ms/step - loss: 0.8817 -
accuracy: 0.6767 - val_loss: 1.5634 - val_accuracy: 0.5160
Epoch 17/20
450/450 [=====] - 5s 11ms/step - loss: 0.8544 -
accuracy: 0.6818 - val_loss: 1.3581 - val_accuracy: 0.5500
Epoch 18/20
450/450 [=====] - 5s 12ms/step - loss: 0.8375 -
accuracy: 0.6947 - val_loss: 2.0195 - val_accuracy: 0.4120
Epoch 19/20
450/450 [=====] - 6s 12ms/step - loss: 0.8080 -
accuracy: 0.7067 - val_loss: 1.2482 - val_accuracy: 0.5800
Epoch 20/20
450/450 [=====] - 5s 11ms/step - loss: 0.7768 -
accuracy: 0.7142 - val_loss: 1.2149 - val_accuracy: 0.5880

```

0.1.6 Question 3.6: Observe model performance with data augmentation

Leverage your best model with the data augmentation and try to observe the difference in performance between using data augmentation and not using it.

[4 points]

From the observation, in terms of the val_accuracy, it is 0.588, not improved a lot. but I noticed that, without data augmentation, the accuracy is almost 1 while val_accuracy is only around 0.6, which meaning that there was a huge overfitting issue.

With the data augmentation, although validation accuracy does not improve a lot, the accuracy downs to 0.7142, meaning that the overfitting issue has been improved a lot

0.1.7 Question 3.7: Explore data mixup technique

[4 points]

Data mixup is another super-simple technique used to boost the generalization ability of deep learning models. You need to incorporate data mixup technique to the above deep learning model and experiment its performance. There are some papers and documents for data mixup as follows:
- Main paper for data mixup [link for main paper](#) and a good article [article link](#).

You need to extend your model developed above, train a model using data mixup, and write your observations and comments about the result.

```

[ ]: # YOU ARE REQUIRED TO INSERT YOUR CODES IN THIS CELL
# You can add more cells if necessary
import numpy as np
class YourModel(BaseImageClassifier):
    def __init__(self,
                  name='network1',
                  width=32, height=32, depth=3,

```



```

        num_blocks=2,
        feature_maps=32,
        num_classes=4,
        drop_rate=0.2,
        batch_norm = None,
        is_augmentation = False,
        activation_func='relu',
        use_skip = True,
        optimizer='adam',
        batch_size=10,
        num_epochs= 20,
        learning_rate=0.0001):
    super(YourModel, self).__init__(name, width, height, depth, num_blocks,
↪feature_maps, num_classes, drop_rate, batch_norm, is_augmentation,
        activation_func, use_skip, optimizer,
↪batch_size, num_epochs, learning_rate)

    def custom_block(self, x, filters):
        # skip connection
        shortcut = x

        # First Conv layer
        x = layers.Conv2D(filters=filters, kernel_size=(3, 3), strides=(1, 1),
↪padding="same")(x)
        if self.batch_norm:
            x = layers.BatchNormalization()(x)
        x = layers.Activation(self.activation_func)(x)

        # Second Conv layer
        x = layers.Conv2D(filters=filters, kernel_size=(3, 3), strides=(1, 1),
↪padding="same")(x)
        if self.batch_norm:
            x = layers.BatchNormalization()(x)

        # Add the shortcut to the output
        if self.use_skip:
            x = tf.keras.layers.add([shortcut, x])
        x = layers.Activation(self.activation_func)(x)

        # Mean Pooling layer
        x = layers.AveragePooling2D(pool_size=(2, 2), strides=(2, 2),
↪padding="same")(x)

        # Dropout
        x = layers.Dropout(rate=self.drop_rate)(x)

    return x

```

```

def build_cnn(self):
    input_tensor = layers.Input(shape=(self.height, self.width, self.depth))
    x = input_tensor

    for current_feature_maps in self.feature_maps:
        # Initial convolution
        x = layers.Conv2D(current_feature_maps, (3, 3), padding='same',
↪activation=None)(x)
        if self.batch_norm:
            x = layers.BatchNormalization()(x)
        x = layers.Activation(self.activation_func)(x)

        # Second convolution with skip connection
        shortcut = x
        x = layers.Conv2D(current_feature_maps, (3, 3), padding='same',
↪activation=None)(x)
        if self.batch_norm:
            x = layers.BatchNormalization()(x)
        if self.use_skip:
            x = layers.Add()([shortcut, x])
        x = layers.Activation(self.activation_func)(x)

        # Mean Pooling and Dropout
        x = layers.AveragePooling2D(pool_size=(2, 2), padding='same')(x)
        if self.drop_rate > 0:
            x = layers.Dropout(self.drop_rate)(x)

    x = layers.Flatten()(x)
    x = layers.Dense(self.num_classes, activation='softmax')(x)

    self.model = models.Model(inputs=input_tensor, outputs=x)
    self.model.compile(optimizer=self.optimizer,
↪loss='sparse_categorical_crossentropy', metrics=['accuracy'])

def generate_mixup_data(self, x, y, alpha=0.2):
    """
    Apply mixup technique
    """
    if alpha > 0:
        lam = np.random.beta(alpha, alpha)
    else:
        lam = 1

    batch_size = x.shape[0]
    index = np.random.permutation(batch_size)

```

```

mixed_x = lam * x + (1 - lam) * x[index] #mixup
mixed_y = lam * y + (1 - lam) * y[index] #mixup

return mixed_x, mixed_y

def fit(self, data_manager, batch_size=None, num_epochs=None,
↪use_mixup=False, alpha=0.2):
    batch_size = self.batch_size if batch_size is None else batch_size
    num_epochs = self.num_epochs if num_epochs is None else num_epochs

    x_train_batch = self.optimize_data_pipeline(data_manager.ds_train,
↪batch_size=batch_size)
    x_val_batch = self.optimize_data_pipeline(data_manager.ds_val,
↪batch_size=batch_size)

    if self.is_augmentation: #code for data augmentation
        train_data_list, train_labels_list = [], []
        for data, labels in x_train_batch:
            train_data_list.append(data.numpy())
            train_labels_list.append(labels.numpy())

    train_data = np.concatenate(train_data_list, axis=0)
    train_labels = np.concatenate(train_labels_list, axis=0)

    datagen = ImageDataGenerator(rotation_range=10,
                                width_shift_range=0.1,
                                height_shift_range=0.1,
                                horizontal_flip=True)

    datagen.fit(train_data)

    self.model.fit(datagen.flow(train_data, train_labels, batch_size=self.
↪batch_size),

                    validation_data=x_val_batch,
                    epochs=self.num_epochs,
                    verbose=self.verbose)

    if use_mixup:
        for epoch in range(num_epochs):
            for step, (x_batch, y_batch) in enumerate(x_train_batch):
                # Apply mixup
                mixed_x, mixed_y = self.generate_mixup_data(x_batch.numpy(),
↪y_batch.numpy(), alpha=alpha)

                # Training step

```

```

        self.model.train_on_batch(mixed_x, mixed_y)
    else:
        self.model.fit(x_train_batch, validation_data=x_val_batch, epochs=self.
↪num_epochs, batch_size=self.batch_size, verbose=self.verbose)

```

```

[ ]: model_3_7=YourModel(name='q3_7',
                        feature_maps=32,
                        num_classes=data_manager.n_classes,
                        num_blocks=4,
                        drop_rate=0.0,
                        batch_norm=True,
                        is_augmentation=True,
                        use_skip=True,
                        optimizer='sgd',
                        learning_rate=0.001)

```

```

model_3_7.build_cnn()
model_3_7.fit(data_manager,use_mixup=True)

```

```

Epoch 1/20
450/450 [=====] - 9s 13ms/step - loss: 2.0255 -
accuracy: 0.2962 - val_loss: 2.0315 - val_accuracy: 0.3100
Epoch 2/20
450/450 [=====] - 6s 13ms/step - loss: 1.6694 -
accuracy: 0.3902 - val_loss: 1.5157 - val_accuracy: 0.4380
Epoch 3/20
450/450 [=====] - 6s 13ms/step - loss: 1.5102 -
accuracy: 0.4531 - val_loss: 2.1353 - val_accuracy: 0.3340
Epoch 4/20
450/450 [=====] - 5s 12ms/step - loss: 1.4502 -
accuracy: 0.4640 - val_loss: 1.8327 - val_accuracy: 0.4120
Epoch 5/20
450/450 [=====] - 6s 13ms/step - loss: 1.3412 -
accuracy: 0.5102 - val_loss: 1.9333 - val_accuracy: 0.3320
Epoch 6/20
450/450 [=====] - 5s 11ms/step - loss: 1.2960 -
accuracy: 0.5253 - val_loss: 1.2722 - val_accuracy: 0.5240
Epoch 7/20
450/450 [=====] - 6s 14ms/step - loss: 1.2365 -
accuracy: 0.5409 - val_loss: 1.7553 - val_accuracy: 0.4080
Epoch 8/20
450/450 [=====] - 5s 12ms/step - loss: 1.1713 -
accuracy: 0.5664 - val_loss: 1.1427 - val_accuracy: 0.5820
Epoch 9/20
450/450 [=====] - 7s 15ms/step - loss: 1.1278 -
accuracy: 0.5904 - val_loss: 1.2963 - val_accuracy: 0.5280

```

```

Epoch 10/20
450/450 [=====] - 6s 13ms/step - loss: 1.1182 -
accuracy: 0.5798 - val_loss: 1.2589 - val_accuracy: 0.5900
Epoch 11/20
450/450 [=====] - 5s 12ms/step - loss: 1.0554 -
accuracy: 0.6100 - val_loss: 1.2019 - val_accuracy: 0.5740
Epoch 12/20
450/450 [=====] - 6s 13ms/step - loss: 1.0309 -
accuracy: 0.6213 - val_loss: 1.8021 - val_accuracy: 0.4760
Epoch 13/20
450/450 [=====] - 5s 12ms/step - loss: 0.9816 -
accuracy: 0.6318 - val_loss: 1.5859 - val_accuracy: 0.4980
Epoch 14/20
450/450 [=====] - 6s 14ms/step - loss: 0.9381 -
accuracy: 0.6631 - val_loss: 1.3570 - val_accuracy: 0.5500
Epoch 15/20
450/450 [=====] - 5s 12ms/step - loss: 0.9351 -
accuracy: 0.6580 - val_loss: 1.3212 - val_accuracy: 0.5620
Epoch 16/20
450/450 [=====] - 6s 14ms/step - loss: 0.8765 -
accuracy: 0.6804 - val_loss: 1.5741 - val_accuracy: 0.4840
Epoch 17/20
450/450 [=====] - 5s 12ms/step - loss: 0.8558 -
accuracy: 0.6867 - val_loss: 1.2008 - val_accuracy: 0.5880
Epoch 18/20
450/450 [=====] - 5s 12ms/step - loss: 0.8264 -
accuracy: 0.7022 - val_loss: 1.1152 - val_accuracy: 0.6000
Epoch 19/20
450/450 [=====] - 6s 13ms/step - loss: 0.7997 -
accuracy: 0.7091 - val_loss: 1.2903 - val_accuracy: 0.5880
Epoch 20/20
450/450 [=====] - 5s 12ms/step - loss: 0.7928 -
accuracy: 0.7131 - val_loss: 1.0881 - val_accuracy: 0.6220

```

Result: “loss: 0.7928 - accuracy: 0.7131 - val_loss: 1.0881 - val_accuracy: 0.6220”

It is noticed that val_accuracy improves.

```

[ ]: #####The following code is just for further fine tuning the_
      ↪ parameters#####

# try to further fine tune the parameters to find the best combination
drop_rates = [0.0, 0.2, 0.5]
optimizers = ['sgd', 'adam', 'rmsprop']
alphas = [0.1, 0.2, 0.5]

# Placeholder
best_params = None

```

```

best_val_accuracy = 0

# Iterate over all combinations
for drop_rate in drop_rates:
    for optimizer in optimizers:
        for alpha in alphas:
            # Create a model with the current parameter combination
            model_3_7_2 = YourModel(name='q3_7_2',
                                     feature_maps=32,
                                     num_classes=data_manager.n_classes,
                                     num_blocks=4,
                                     drop_rate=drop_rate,
                                     batch_norm=True,
                                     is_augmentation=True,
                                     use_skip=True,
                                     optimizer=optimizer,
                                     learning_rate=0.001)

            # Build and train the model
            model_3_7_2.build_cnn()
            model_3_7_2.fit(data_manager, use_mixup=True, alpha=alpha)

            # Evaluate the model on the validation set
            val_accuracy = max(model_3_7_2.model.history.
                               ↪history['val_accuracy'])

            # Check if this performance is the best so far
            if val_accuracy > best_val_accuracy:
                best_val_performance = val_accuracy
                best_params = {'drop_rate': drop_rate, 'optimizer': optimizer, ↪
                               ↪'alpha': alpha}

print("Best Parameters:", best_params)
print("Best Validation Performance:", best_val_accuracy)

```

Epoch 1/20

450/450 [=====] - 9s 13ms/step - loss: 2.0354 - accuracy: 0.3029 - val_loss: 1.7770 - val_accuracy: 0.3340

Epoch 2/20

450/450 [=====] - 6s 13ms/step - loss: 1.7012 - accuracy: 0.3849 - val_loss: 1.5626 - val_accuracy: 0.4180

Epoch 3/20

450/450 [=====] - 5s 12ms/step - loss: 1.5254 - accuracy: 0.4424 - val_loss: 1.7653 - val_accuracy: 0.3640

Epoch 4/20

450/450 [=====] - 6s 14ms/step - loss: 1.4283 -

accuracy: 0.4678 - val_loss: 1.5592 - val_accuracy: 0.4520
 Epoch 5/20
 450/450 [=====] - 5s 12ms/step - loss: 1.3566 -
 accuracy: 0.5080 - val_loss: 1.5697 - val_accuracy: 0.4660
 Epoch 6/20
 450/450 [=====] - 6s 13ms/step - loss: 1.2920 -
 accuracy: 0.5360 - val_loss: 1.3516 - val_accuracy: 0.5020
 Epoch 7/20
 450/450 [=====] - 5s 12ms/step - loss: 1.2373 -
 accuracy: 0.5447 - val_loss: 1.2674 - val_accuracy: 0.5360
 Epoch 8/20
 450/450 [=====] - 6s 14ms/step - loss: 1.1924 -
 accuracy: 0.5624 - val_loss: 1.5309 - val_accuracy: 0.4540
 Epoch 9/20
 450/450 [=====] - 5s 12ms/step - loss: 1.1282 -
 accuracy: 0.5951 - val_loss: 1.3653 - val_accuracy: 0.5280
 Epoch 10/20
 450/450 [=====] - 6s 13ms/step - loss: 1.1126 -
 accuracy: 0.5949 - val_loss: 1.7026 - val_accuracy: 0.4520
 Epoch 11/20
 450/450 [=====] - 5s 12ms/step - loss: 1.0546 -
 accuracy: 0.6198 - val_loss: 1.8052 - val_accuracy: 0.4560
 Epoch 12/20
 450/450 [=====] - 5s 12ms/step - loss: 1.0238 -
 accuracy: 0.6189 - val_loss: 1.4236 - val_accuracy: 0.5300
 Epoch 13/20
 450/450 [=====] - 6s 14ms/step - loss: 0.9884 -
 accuracy: 0.6373 - val_loss: 1.1623 - val_accuracy: 0.5960
 Epoch 14/20
 450/450 [=====] - 5s 12ms/step - loss: 0.9452 -
 accuracy: 0.6544 - val_loss: 1.4605 - val_accuracy: 0.5200
 Epoch 15/20
 450/450 [=====] - 6s 14ms/step - loss: 0.9289 -
 accuracy: 0.6598 - val_loss: 1.7715 - val_accuracy: 0.4540
 Epoch 16/20
 450/450 [=====] - 6s 12ms/step - loss: 0.8675 -
 accuracy: 0.6827 - val_loss: 1.1270 - val_accuracy: 0.6220
 Epoch 17/20
 450/450 [=====] - 6s 13ms/step - loss: 0.8843 -
 accuracy: 0.6780 - val_loss: 1.2103 - val_accuracy: 0.5900
 Epoch 18/20
 450/450 [=====] - 5s 12ms/step - loss: 0.8475 -
 accuracy: 0.6967 - val_loss: 2.5453 - val_accuracy: 0.3960
 Epoch 19/20
 450/450 [=====] - 6s 14ms/step - loss: 0.7988 -
 accuracy: 0.7047 - val_loss: 1.0827 - val_accuracy: 0.6320
 Epoch 20/20
 450/450 [=====] - 5s 12ms/step - loss: 0.8029 -

accuracy: 0.7038 - val_loss: 1.4081 - val_accuracy: 0.5740
 Epoch 1/20
 450/450 [=====] - 10s 17ms/step - loss: 2.0276 -
 accuracy: 0.3073 - val_loss: 1.7634 - val_accuracy: 0.3740
 Epoch 2/20
 450/450 [=====] - 5s 12ms/step - loss: 1.6939 -
 accuracy: 0.3880 - val_loss: 1.6321 - val_accuracy: 0.4100
 Epoch 3/20
 450/450 [=====] - 6s 14ms/step - loss: 1.5266 -
 accuracy: 0.4431 - val_loss: 2.4063 - val_accuracy: 0.2760
 Epoch 4/20
 450/450 [=====] - 5s 12ms/step - loss: 1.4411 -
 accuracy: 0.4747 - val_loss: 1.5643 - val_accuracy: 0.4740
 Epoch 5/20
 450/450 [=====] - 6s 14ms/step - loss: 1.3557 -
 accuracy: 0.4949 - val_loss: 1.4572 - val_accuracy: 0.5040
 Epoch 6/20
 450/450 [=====] - 5s 12ms/step - loss: 1.3041 -
 accuracy: 0.5211 - val_loss: 1.5791 - val_accuracy: 0.4600
 Epoch 7/20
 450/450 [=====] - 6s 13ms/step - loss: 1.2475 -
 accuracy: 0.5384 - val_loss: 1.4807 - val_accuracy: 0.4800
 Epoch 8/20
 450/450 [=====] - 5s 12ms/step - loss: 1.2113 -
 accuracy: 0.5629 - val_loss: 1.4128 - val_accuracy: 0.5000
 Epoch 9/20
 450/450 [=====] - 6s 13ms/step - loss: 1.1169 -
 accuracy: 0.5940 - val_loss: 1.1977 - val_accuracy: 0.5540
 Epoch 10/20
 450/450 [=====] - 6s 12ms/step - loss: 1.1066 -
 accuracy: 0.5933 - val_loss: 1.2769 - val_accuracy: 0.5420
 Epoch 11/20
 450/450 [=====] - 5s 12ms/step - loss: 1.0626 -
 accuracy: 0.6171 - val_loss: 1.8713 - val_accuracy: 0.4620
 Epoch 12/20
 450/450 [=====] - 6s 13ms/step - loss: 1.0526 -
 accuracy: 0.6167 - val_loss: 1.7249 - val_accuracy: 0.4420
 Epoch 13/20
 450/450 [=====] - 5s 12ms/step - loss: 1.0185 -
 accuracy: 0.6304 - val_loss: 1.3825 - val_accuracy: 0.5100
 Epoch 14/20
 450/450 [=====] - 6s 14ms/step - loss: 0.9636 -
 accuracy: 0.6531 - val_loss: 1.5294 - val_accuracy: 0.5140
 Epoch 15/20
 450/450 [=====] - 5s 12ms/step - loss: 0.9518 -
 accuracy: 0.6504 - val_loss: 1.2967 - val_accuracy: 0.5620
 Epoch 16/20
 450/450 [=====] - 6s 14ms/step - loss: 0.9205 -

accuracy: 0.6691 - val_loss: 1.0866 - val_accuracy: 0.5940
 Epoch 17/20
 450/450 [=====] - 5s 12ms/step - loss: 0.8674 -
 accuracy: 0.6829 - val_loss: 1.2322 - val_accuracy: 0.5700
 Epoch 18/20
 450/450 [=====] - 6s 14ms/step - loss: 0.8370 -
 accuracy: 0.6951 - val_loss: 1.8385 - val_accuracy: 0.4560
 Epoch 19/20
 450/450 [=====] - 5s 12ms/step - loss: 0.8487 -
 accuracy: 0.6987 - val_loss: 1.3269 - val_accuracy: 0.5420
 Epoch 20/20
 450/450 [=====] - 6s 13ms/step - loss: 0.7964 -
 accuracy: 0.7082 - val_loss: 1.2248 - val_accuracy: 0.5980
 Epoch 1/20
 450/450 [=====] - 8s 13ms/step - loss: 2.0375 -
 accuracy: 0.3084 - val_loss: 2.2273 - val_accuracy: 0.3380
 Epoch 2/20
 450/450 [=====] - 6s 14ms/step - loss: 1.6589 -
 accuracy: 0.3953 - val_loss: 2.6954 - val_accuracy: 0.2200
 Epoch 3/20
 450/450 [=====] - 5s 12ms/step - loss: 1.5299 -
 accuracy: 0.4418 - val_loss: 1.4348 - val_accuracy: 0.4540
 Epoch 4/20
 450/450 [=====] - 6s 14ms/step - loss: 1.4245 -
 accuracy: 0.4831 - val_loss: 2.1491 - val_accuracy: 0.3300
 Epoch 5/20
 450/450 [=====] - 5s 12ms/step - loss: 1.3310 -
 accuracy: 0.5113 - val_loss: 1.8443 - val_accuracy: 0.4200
 Epoch 6/20
 450/450 [=====] - 6s 14ms/step - loss: 1.2622 -
 accuracy: 0.5358 - val_loss: 1.3673 - val_accuracy: 0.5080
 Epoch 7/20
 450/450 [=====] - 5s 12ms/step - loss: 1.2205 -
 accuracy: 0.5533 - val_loss: 1.6660 - val_accuracy: 0.4160
 Epoch 8/20
 450/450 [=====] - 6s 14ms/step - loss: 1.1630 -
 accuracy: 0.5733 - val_loss: 1.2215 - val_accuracy: 0.5580
 Epoch 9/20
 450/450 [=====] - 5s 12ms/step - loss: 1.1237 -
 accuracy: 0.5844 - val_loss: 1.1851 - val_accuracy: 0.5780
 Epoch 10/20
 450/450 [=====] - 6s 14ms/step - loss: 1.1110 -
 accuracy: 0.5929 - val_loss: 1.1893 - val_accuracy: 0.5720
 Epoch 11/20
 450/450 [=====] - 5s 12ms/step - loss: 1.0519 -
 accuracy: 0.6153 - val_loss: 1.5565 - val_accuracy: 0.4820
 Epoch 12/20
 450/450 [=====] - 6s 13ms/step - loss: 1.0166 -

accuracy: 0.6329 - val_loss: 1.0919 - val_accuracy: 0.5980
 Epoch 13/20
 450/450 [=====] - 6s 12ms/step - loss: 0.9756 -
 accuracy: 0.6451 - val_loss: 1.6290 - val_accuracy: 0.4820
 Epoch 14/20
 450/450 [=====] - 6s 14ms/step - loss: 0.9385 -
 accuracy: 0.6647 - val_loss: 1.1805 - val_accuracy: 0.5720
 Epoch 15/20
 450/450 [=====] - 5s 12ms/step - loss: 0.9199 -
 accuracy: 0.6689 - val_loss: 1.1943 - val_accuracy: 0.5740
 Epoch 16/20
 450/450 [=====] - 6s 14ms/step - loss: 0.8931 -
 accuracy: 0.6767 - val_loss: 1.1112 - val_accuracy: 0.6000
 Epoch 17/20
 450/450 [=====] - 5s 12ms/step - loss: 0.8654 -
 accuracy: 0.6871 - val_loss: 1.1290 - val_accuracy: 0.6020
 Epoch 18/20
 450/450 [=====] - 5s 12ms/step - loss: 0.8332 -
 accuracy: 0.6969 - val_loss: 1.0585 - val_accuracy: 0.6140
 Epoch 19/20
 450/450 [=====] - 6s 14ms/step - loss: 0.8083 -
 accuracy: 0.7047 - val_loss: 1.5422 - val_accuracy: 0.5420
 Epoch 20/20
 450/450 [=====] - 6s 13ms/step - loss: 0.7943 -
 accuracy: 0.7098 - val_loss: 1.3984 - val_accuracy: 0.5540
 Epoch 1/20
 450/450 [=====] - 13s 14ms/step - loss: 2.1137 -
 accuracy: 0.2707 - val_loss: 1.7632 - val_accuracy: 0.3660
 Epoch 2/20
 450/450 [=====] - 7s 15ms/step - loss: 1.7832 -
 accuracy: 0.3436 - val_loss: 1.8600 - val_accuracy: 0.3700
 Epoch 3/20
 450/450 [=====] - 6s 12ms/step - loss: 1.6610 -
 accuracy: 0.3898 - val_loss: 1.5037 - val_accuracy: 0.4460
 Epoch 4/20
 450/450 [=====] - 7s 14ms/step - loss: 1.5607 -
 accuracy: 0.4240 - val_loss: 1.4796 - val_accuracy: 0.4440
 Epoch 5/20
 450/450 [=====] - 6s 12ms/step - loss: 1.4926 -
 accuracy: 0.4482 - val_loss: 1.4034 - val_accuracy: 0.4700
 Epoch 6/20
 450/450 [=====] - 6s 14ms/step - loss: 1.4231 -
 accuracy: 0.4744 - val_loss: 1.5006 - val_accuracy: 0.4240
 Epoch 7/20
 450/450 [=====] - 6s 12ms/step - loss: 1.3639 -
 accuracy: 0.4900 - val_loss: 1.5254 - val_accuracy: 0.4620
 Epoch 8/20
 450/450 [=====] - 6s 14ms/step - loss: 1.2941 -

accuracy: 0.5158 - val_loss: 1.3231 - val_accuracy: 0.5100
 Epoch 9/20
 450/450 [=====] - 6s 12ms/step - loss: 1.2272 -
 accuracy: 0.5476 - val_loss: 1.6451 - val_accuracy: 0.4380
 Epoch 10/20
 450/450 [=====] - 6s 14ms/step - loss: 1.1920 -
 accuracy: 0.5569 - val_loss: 1.6471 - val_accuracy: 0.4220
 Epoch 11/20
 450/450 [=====] - 5s 12ms/step - loss: 1.1457 -
 accuracy: 0.5851 - val_loss: 1.8055 - val_accuracy: 0.3800
 Epoch 12/20
 450/450 [=====] - 6s 14ms/step - loss: 1.1070 -
 accuracy: 0.5951 - val_loss: 1.6724 - val_accuracy: 0.4480
 Epoch 13/20
 450/450 [=====] - 6s 13ms/step - loss: 1.0424 -
 accuracy: 0.6116 - val_loss: 1.0807 - val_accuracy: 0.6140
 Epoch 14/20
 450/450 [=====] - 6s 13ms/step - loss: 1.0056 -
 accuracy: 0.6293 - val_loss: 1.3152 - val_accuracy: 0.5120
 Epoch 15/20
 450/450 [=====] - 5s 12ms/step - loss: 0.9842 -
 accuracy: 0.6407 - val_loss: 1.1997 - val_accuracy: 0.5900
 Epoch 16/20
 450/450 [=====] - 7s 14ms/step - loss: 0.9537 -
 accuracy: 0.6460 - val_loss: 1.1192 - val_accuracy: 0.6100
 Epoch 17/20
 450/450 [=====] - 6s 12ms/step - loss: 0.9025 -
 accuracy: 0.6698 - val_loss: 1.5659 - val_accuracy: 0.4740
 Epoch 18/20
 450/450 [=====] - 6s 14ms/step - loss: 0.8839 -
 accuracy: 0.6818 - val_loss: 1.0431 - val_accuracy: 0.6540
 Epoch 19/20
 450/450 [=====] - 6s 13ms/step - loss: 0.8552 -
 accuracy: 0.6909 - val_loss: 1.5771 - val_accuracy: 0.4940
 Epoch 20/20
 450/450 [=====] - 6s 14ms/step - loss: 0.8052 -
 accuracy: 0.6998 - val_loss: 1.3468 - val_accuracy: 0.5360
 Epoch 1/20
 450/450 [=====] - 13s 15ms/step - loss: 2.1242 -
 accuracy: 0.2662 - val_loss: 2.4718 - val_accuracy: 0.2660
 Epoch 2/20
 450/450 [=====] - 6s 14ms/step - loss: 1.7727 -
 accuracy: 0.3456 - val_loss: 1.9867 - val_accuracy: 0.2840
 Epoch 3/20
 450/450 [=====] - 6s 13ms/step - loss: 1.6398 -
 accuracy: 0.3967 - val_loss: 1.7196 - val_accuracy: 0.4220
 Epoch 4/20
 450/450 [=====] - 7s 14ms/step - loss: 1.5126 -

accuracy: 0.4433 - val_loss: 1.5613 - val_accuracy: 0.4380
Epoch 5/20
450/450 [=====] - 6s 13ms/step - loss: 1.4180 -
accuracy: 0.4807 - val_loss: 1.3429 - val_accuracy: 0.5180
Epoch 6/20
450/450 [=====] - 7s 15ms/step - loss: 1.3524 -
accuracy: 0.5022 - val_loss: 1.3083 - val_accuracy: 0.4980
Epoch 7/20
450/450 [=====] - 6s 13ms/step - loss: 1.2924 -
accuracy: 0.5147 - val_loss: 1.4205 - val_accuracy: 0.4740
Epoch 8/20
450/450 [=====] - 6s 12ms/step - loss: 1.2380 -
accuracy: 0.5331 - val_loss: 2.0379 - val_accuracy: 0.4020
Epoch 9/20
450/450 [=====] - 7s 15ms/step - loss: 1.1939 -
accuracy: 0.5611 - val_loss: 1.1263 - val_accuracy: 0.6000
Epoch 10/20
450/450 [=====] - 6s 13ms/step - loss: 1.1551 -
accuracy: 0.5733 - val_loss: 1.4301 - val_accuracy: 0.4740
Epoch 11/20
450/450 [=====] - 7s 15ms/step - loss: 1.1111 -
accuracy: 0.5889 - val_loss: 1.2455 - val_accuracy: 0.5660
Epoch 12/20
450/450 [=====] - 6s 13ms/step - loss: 1.0774 -
accuracy: 0.6004 - val_loss: 1.3632 - val_accuracy: 0.5100
Epoch 13/20
450/450 [=====] - 7s 15ms/step - loss: 1.0208 -
accuracy: 0.6198 - val_loss: 1.2225 - val_accuracy: 0.5660
Epoch 14/20
450/450 [=====] - 6s 12ms/step - loss: 0.9744 -
accuracy: 0.6424 - val_loss: 1.6167 - val_accuracy: 0.5060
Epoch 15/20
450/450 [=====] - 7s 14ms/step - loss: 0.9561 -
accuracy: 0.6511 - val_loss: 1.0791 - val_accuracy: 0.6220
Epoch 16/20
450/450 [=====] - 6s 13ms/step - loss: 0.9172 -
accuracy: 0.6627 - val_loss: 1.0447 - val_accuracy: 0.6240
Epoch 17/20
450/450 [=====] - 6s 14ms/step - loss: 0.8756 -
accuracy: 0.6791 - val_loss: 1.1586 - val_accuracy: 0.6160
Epoch 18/20
450/450 [=====] - 6s 12ms/step - loss: 0.8612 -
accuracy: 0.6871 - val_loss: 1.1174 - val_accuracy: 0.6120
Epoch 19/20
450/450 [=====] - 7s 15ms/step - loss: 0.8359 -
accuracy: 0.6918 - val_loss: 1.8062 - val_accuracy: 0.4760
Epoch 20/20
450/450 [=====] - 6s 12ms/step - loss: 0.7982 -

accuracy: 0.7031 - val_loss: 1.2887 - val_accuracy: 0.5780
 Epoch 1/20
 450/450 [=====] - 13s 14ms/step - loss: 2.1532 -
 accuracy: 0.2756 - val_loss: 1.9820 - val_accuracy: 0.2940
 Epoch 2/20
 450/450 [=====] - 7s 15ms/step - loss: 1.7762 -
 accuracy: 0.3469 - val_loss: 2.4789 - val_accuracy: 0.2040
 Epoch 3/20
 450/450 [=====] - 5s 12ms/step - loss: 1.6531 -
 accuracy: 0.3898 - val_loss: 2.7400 - val_accuracy: 0.2960
 Epoch 4/20
 450/450 [=====] - 6s 14ms/step - loss: 1.5396 -
 accuracy: 0.4251 - val_loss: 1.9080 - val_accuracy: 0.3280
 Epoch 5/20
 450/450 [=====] - 6s 13ms/step - loss: 1.4304 -
 accuracy: 0.4662 - val_loss: 1.6940 - val_accuracy: 0.3920
 Epoch 6/20
 450/450 [=====] - 6s 14ms/step - loss: 1.3760 -
 accuracy: 0.4813 - val_loss: 3.2806 - val_accuracy: 0.2300
 Epoch 7/20
 450/450 [=====] - 6s 12ms/step - loss: 1.3063 -
 accuracy: 0.5227 - val_loss: 1.4182 - val_accuracy: 0.4960
 Epoch 8/20
 450/450 [=====] - 6s 14ms/step - loss: 1.2339 -
 accuracy: 0.5373 - val_loss: 1.6909 - val_accuracy: 0.4320
 Epoch 9/20
 450/450 [=====] - 6s 12ms/step - loss: 1.1874 -
 accuracy: 0.5573 - val_loss: 1.4353 - val_accuracy: 0.5320
 Epoch 10/20
 450/450 [=====] - 7s 14ms/step - loss: 1.1393 -
 accuracy: 0.5753 - val_loss: 1.3127 - val_accuracy: 0.5220
 Epoch 11/20
 450/450 [=====] - 6s 12ms/step - loss: 1.1001 -
 accuracy: 0.5980 - val_loss: 1.1823 - val_accuracy: 0.5380
 Epoch 12/20
 450/450 [=====] - 7s 14ms/step - loss: 1.0573 -
 accuracy: 0.6102 - val_loss: 1.2395 - val_accuracy: 0.5720
 Epoch 13/20
 450/450 [=====] - 5s 12ms/step - loss: 1.0240 -
 accuracy: 0.6204 - val_loss: 1.7087 - val_accuracy: 0.4740
 Epoch 14/20
 450/450 [=====] - 6s 14ms/step - loss: 0.9672 -
 accuracy: 0.6433 - val_loss: 1.1337 - val_accuracy: 0.6340
 Epoch 15/20
 450/450 [=====] - 6s 12ms/step - loss: 0.9508 -
 accuracy: 0.6522 - val_loss: 1.7439 - val_accuracy: 0.4740
 Epoch 16/20
 450/450 [=====] - 7s 15ms/step - loss: 0.9058 -

accuracy: 0.6656 - val_loss: 1.1081 - val_accuracy: 0.5900
 Epoch 17/20
 450/450 [=====] - 6s 12ms/step - loss: 0.8605 -
 accuracy: 0.6789 - val_loss: 1.2743 - val_accuracy: 0.5580
 Epoch 18/20
 450/450 [=====] - 7s 15ms/step - loss: 0.8406 -
 accuracy: 0.6891 - val_loss: 1.1196 - val_accuracy: 0.6180
 Epoch 19/20
 450/450 [=====] - 6s 12ms/step - loss: 0.8105 -
 accuracy: 0.7016 - val_loss: 1.2957 - val_accuracy: 0.5540
 Epoch 20/20
 450/450 [=====] - 6s 14ms/step - loss: 0.7908 -
 accuracy: 0.7016 - val_loss: 1.3962 - val_accuracy: 0.5580
 Epoch 1/20
 450/450 [=====] - 11s 14ms/step - loss: 2.2347 -
 accuracy: 0.2516 - val_loss: 2.2594 - val_accuracy: 0.2440
 Epoch 2/20
 450/450 [=====] - 5s 12ms/step - loss: 1.8123 -
 accuracy: 0.3407 - val_loss: 1.6445 - val_accuracy: 0.3840
 Epoch 3/20
 450/450 [=====] - 6s 13ms/step - loss: 1.6761 -
 accuracy: 0.3796 - val_loss: 1.9524 - val_accuracy: 0.3660
 Epoch 4/20
 450/450 [=====] - 5s 12ms/step - loss: 1.5618 -
 accuracy: 0.4220 - val_loss: 1.5073 - val_accuracy: 0.4420
 Epoch 5/20
 450/450 [=====] - 5s 12ms/step - loss: 1.4692 -
 accuracy: 0.4542 - val_loss: 2.3699 - val_accuracy: 0.3000
 Epoch 6/20
 450/450 [=====] - 6s 14ms/step - loss: 1.3974 -
 accuracy: 0.4800 - val_loss: 2.2974 - val_accuracy: 0.3520
 Epoch 7/20
 450/450 [=====] - 5s 12ms/step - loss: 1.3239 -
 accuracy: 0.5131 - val_loss: 2.1349 - val_accuracy: 0.3820
 Epoch 8/20
 450/450 [=====] - 6s 14ms/step - loss: 1.2688 -
 accuracy: 0.5322 - val_loss: 1.4154 - val_accuracy: 0.5040
 Epoch 9/20
 450/450 [=====] - 5s 12ms/step - loss: 1.2362 -
 accuracy: 0.5371 - val_loss: 1.6561 - val_accuracy: 0.4480
 Epoch 10/20
 450/450 [=====] - 6s 14ms/step - loss: 1.1827 -
 accuracy: 0.5582 - val_loss: 1.7021 - val_accuracy: 0.4480
 Epoch 11/20
 450/450 [=====] - 5s 12ms/step - loss: 1.1562 -
 accuracy: 0.5731 - val_loss: 1.6370 - val_accuracy: 0.4780
 Epoch 12/20
 450/450 [=====] - 5s 12ms/step - loss: 1.1095 -

accuracy: 0.5922 - val_loss: 2.7953 - val_accuracy: 0.3160
 Epoch 13/20
 450/450 [=====] - 6s 14ms/step - loss: 1.0570 -
 accuracy: 0.6129 - val_loss: 1.7767 - val_accuracy: 0.4960
 Epoch 14/20
 450/450 [=====] - 5s 12ms/step - loss: 1.0465 -
 accuracy: 0.6189 - val_loss: 1.9187 - val_accuracy: 0.4540
 Epoch 15/20
 450/450 [=====] - 6s 13ms/step - loss: 1.0199 -
 accuracy: 0.6251 - val_loss: 1.0610 - val_accuracy: 0.6340
 Epoch 16/20
 450/450 [=====] - 6s 12ms/step - loss: 0.9839 -
 accuracy: 0.6367 - val_loss: 1.4212 - val_accuracy: 0.5380
 Epoch 17/20
 450/450 [=====] - 6s 13ms/step - loss: 0.9423 -
 accuracy: 0.6560 - val_loss: 1.1197 - val_accuracy: 0.6260
 Epoch 18/20
 450/450 [=====] - 6s 14ms/step - loss: 0.9156 -
 accuracy: 0.6678 - val_loss: 1.3959 - val_accuracy: 0.5300
 Epoch 19/20
 450/450 [=====] - 5s 12ms/step - loss: 0.8951 -
 accuracy: 0.6720 - val_loss: 1.1148 - val_accuracy: 0.6180
 Epoch 20/20
 450/450 [=====] - 6s 14ms/step - loss: 0.8747 -
 accuracy: 0.6816 - val_loss: 1.0702 - val_accuracy: 0.6260
 Epoch 1/20
 450/450 [=====] - 11s 14ms/step - loss: 2.2262 -
 accuracy: 0.2522 - val_loss: 1.9462 - val_accuracy: 0.2900
 Epoch 2/20
 450/450 [=====] - 5s 12ms/step - loss: 1.8246 -
 accuracy: 0.3378 - val_loss: 1.6210 - val_accuracy: 0.4000
 Epoch 3/20
 450/450 [=====] - 6s 14ms/step - loss: 1.6568 -
 accuracy: 0.3898 - val_loss: 1.5594 - val_accuracy: 0.4180
 Epoch 4/20
 450/450 [=====] - 5s 12ms/step - loss: 1.5673 -
 accuracy: 0.4262 - val_loss: 2.2974 - val_accuracy: 0.2980
 Epoch 5/20
 450/450 [=====] - 6s 14ms/step - loss: 1.4817 -
 accuracy: 0.4593 - val_loss: 1.3954 - val_accuracy: 0.4820
 Epoch 6/20
 450/450 [=====] - 5s 12ms/step - loss: 1.4050 -
 accuracy: 0.4771 - val_loss: 1.4313 - val_accuracy: 0.5060
 Epoch 7/20
 450/450 [=====] - 6s 14ms/step - loss: 1.3603 -
 accuracy: 0.4942 - val_loss: 1.5945 - val_accuracy: 0.4440
 Epoch 8/20
 450/450 [=====] - 5s 12ms/step - loss: 1.2967 -

accuracy: 0.5213 - val_loss: 1.2634 - val_accuracy: 0.5360
 Epoch 9/20
 450/450 [=====] - 5s 12ms/step - loss: 1.2393 -
 accuracy: 0.5373 - val_loss: 1.4708 - val_accuracy: 0.4800
 Epoch 10/20
 450/450 [=====] - 6s 14ms/step - loss: 1.1984 -
 accuracy: 0.5513 - val_loss: 1.3789 - val_accuracy: 0.5180
 Epoch 11/20
 450/450 [=====] - 5s 12ms/step - loss: 1.1469 -
 accuracy: 0.5780 - val_loss: 1.1787 - val_accuracy: 0.5680
 Epoch 12/20
 450/450 [=====] - 6s 13ms/step - loss: 1.1252 -
 accuracy: 0.5791 - val_loss: 1.4794 - val_accuracy: 0.5060
 Epoch 13/20
 450/450 [=====] - 5s 12ms/step - loss: 1.0845 -
 accuracy: 0.6033 - val_loss: 1.5160 - val_accuracy: 0.4860
 Epoch 14/20
 450/450 [=====] - 6s 14ms/step - loss: 1.0525 -
 accuracy: 0.6142 - val_loss: 1.3910 - val_accuracy: 0.5440
 Epoch 15/20
 450/450 [=====] - 5s 12ms/step - loss: 1.0242 -
 accuracy: 0.6160 - val_loss: 1.7600 - val_accuracy: 0.4880
 Epoch 16/20
 450/450 [=====] - 5s 12ms/step - loss: 0.9876 -
 accuracy: 0.6482 - val_loss: 1.2783 - val_accuracy: 0.5620
 Epoch 17/20
 450/450 [=====] - 6s 14ms/step - loss: 0.9623 -
 accuracy: 0.6491 - val_loss: 1.6105 - val_accuracy: 0.4980
 Epoch 18/20
 450/450 [=====] - 5s 12ms/step - loss: 0.9393 -
 accuracy: 0.6609 - val_loss: 1.4535 - val_accuracy: 0.5200
 Epoch 19/20
 450/450 [=====] - 6s 14ms/step - loss: 0.9060 -
 accuracy: 0.6744 - val_loss: 1.1326 - val_accuracy: 0.6240
 Epoch 20/20
 450/450 [=====] - 5s 12ms/step - loss: 0.8991 -
 accuracy: 0.6704 - val_loss: 1.1843 - val_accuracy: 0.6280
 Epoch 1/20
 450/450 [=====] - 10s 13ms/step - loss: 2.2104 -
 accuracy: 0.2636 - val_loss: 2.5101 - val_accuracy: 0.2780
 Epoch 2/20
 450/450 [=====] - 6s 14ms/step - loss: 1.7879 -
 accuracy: 0.3387 - val_loss: 2.6973 - val_accuracy: 0.2820
 Epoch 3/20
 450/450 [=====] - 5s 12ms/step - loss: 1.6782 -
 accuracy: 0.3733 - val_loss: 3.3359 - val_accuracy: 0.1840
 Epoch 4/20
 450/450 [=====] - 6s 14ms/step - loss: 1.5512 -

accuracy: 0.4251 - val_loss: 1.4515 - val_accuracy: 0.4600
 Epoch 5/20
 450/450 [=====] - 6s 13ms/step - loss: 1.4819 -
 accuracy: 0.4424 - val_loss: 1.8120 - val_accuracy: 0.3900
 Epoch 6/20
 450/450 [=====] - 6s 13ms/step - loss: 1.4198 -
 accuracy: 0.4776 - val_loss: 2.3389 - val_accuracy: 0.3000
 Epoch 7/20
 450/450 [=====] - 6s 13ms/step - loss: 1.3391 -
 accuracy: 0.5049 - val_loss: 1.3621 - val_accuracy: 0.5120
 Epoch 8/20
 450/450 [=====] - 6s 13ms/step - loss: 1.2726 -
 accuracy: 0.5269 - val_loss: 1.2940 - val_accuracy: 0.5440
 Epoch 9/20
 450/450 [=====] - 5s 12ms/step - loss: 1.2359 -
 accuracy: 0.5476 - val_loss: 1.8727 - val_accuracy: 0.4100
 Epoch 10/20
 450/450 [=====] - 6s 14ms/step - loss: 1.1881 -
 accuracy: 0.5513 - val_loss: 1.2728 - val_accuracy: 0.5100
 Epoch 11/20
 450/450 [=====] - 5s 12ms/step - loss: 1.1469 -
 accuracy: 0.5804 - val_loss: 1.2281 - val_accuracy: 0.5520
 Epoch 12/20
 450/450 [=====] - 6s 14ms/step - loss: 1.1189 -
 accuracy: 0.5813 - val_loss: 1.4507 - val_accuracy: 0.5080
 Epoch 13/20
 450/450 [=====] - 5s 12ms/step - loss: 1.0768 -
 accuracy: 0.6100 - val_loss: 1.0950 - val_accuracy: 0.6020
 Epoch 14/20
 450/450 [=====] - 5s 12ms/step - loss: 1.0618 -
 accuracy: 0.6129 - val_loss: 1.4317 - val_accuracy: 0.5660
 Epoch 15/20
 450/450 [=====] - 5s 12ms/step - loss: 1.0129 -
 accuracy: 0.6316 - val_loss: 1.0892 - val_accuracy: 0.6300
 Epoch 16/20
 450/450 [=====] - 6s 12ms/step - loss: 1.0033 -
 accuracy: 0.6300 - val_loss: 1.2042 - val_accuracy: 0.5900
 Epoch 17/20
 450/450 [=====] - 6s 12ms/step - loss: 0.9634 -
 accuracy: 0.6509 - val_loss: 1.8299 - val_accuracy: 0.4560
 Epoch 18/20
 450/450 [=====] - 6s 13ms/step - loss: 0.9452 -
 accuracy: 0.6607 - val_loss: 1.1356 - val_accuracy: 0.6120
 Epoch 19/20
 450/450 [=====] - 5s 12ms/step - loss: 0.9272 -
 accuracy: 0.6589 - val_loss: 1.3072 - val_accuracy: 0.5800
 Epoch 20/20
 450/450 [=====] - 6s 14ms/step - loss: 0.9116 -

accuracy: 0.6673 - val_loss: 1.0930 - val_accuracy: 0.6180
 Epoch 1/20
 450/450 [=====] - 9s 15ms/step - loss: 2.1962 -
 accuracy: 0.2596 - val_loss: 1.7909 - val_accuracy: 0.2980
 Epoch 2/20
 450/450 [=====] - 6s 13ms/step - loss: 1.8666 -
 accuracy: 0.3247 - val_loss: 2.1943 - val_accuracy: 0.2660
 Epoch 3/20
 450/450 [=====] - 6s 14ms/step - loss: 1.7486 -
 accuracy: 0.3536 - val_loss: 2.0690 - val_accuracy: 0.3300
 Epoch 4/20
 450/450 [=====] - 6s 14ms/step - loss: 1.6602 -
 accuracy: 0.3878 - val_loss: 1.9508 - val_accuracy: 0.3540
 Epoch 5/20
 450/450 [=====] - 6s 13ms/step - loss: 1.5763 -
 accuracy: 0.4138 - val_loss: 1.5519 - val_accuracy: 0.4360
 Epoch 6/20
 450/450 [=====] - 7s 15ms/step - loss: 1.5061 -
 accuracy: 0.4449 - val_loss: 2.3199 - val_accuracy: 0.3520
 Epoch 7/20
 450/450 [=====] - 6s 12ms/step - loss: 1.5075 -
 accuracy: 0.4413 - val_loss: 1.6965 - val_accuracy: 0.4120
 Epoch 8/20
 450/450 [=====] - 7s 14ms/step - loss: 1.4493 -
 accuracy: 0.4682 - val_loss: 1.7184 - val_accuracy: 0.4420
 Epoch 9/20
 450/450 [=====] - 6s 12ms/step - loss: 1.4069 -
 accuracy: 0.4771 - val_loss: 1.4615 - val_accuracy: 0.4820
 Epoch 10/20
 450/450 [=====] - 6s 14ms/step - loss: 1.3738 -
 accuracy: 0.4898 - val_loss: 1.3040 - val_accuracy: 0.5360
 Epoch 11/20
 450/450 [=====] - 6s 13ms/step - loss: 1.3305 -
 accuracy: 0.5013 - val_loss: 2.4997 - val_accuracy: 0.3680
 Epoch 12/20
 450/450 [=====] - 7s 15ms/step - loss: 1.3334 -
 accuracy: 0.5049 - val_loss: 1.4311 - val_accuracy: 0.4920
 Epoch 13/20
 450/450 [=====] - 6s 13ms/step - loss: 1.2870 -
 accuracy: 0.5231 - val_loss: 1.3582 - val_accuracy: 0.5260
 Epoch 14/20
 450/450 [=====] - 6s 14ms/step - loss: 1.2899 -
 accuracy: 0.5278 - val_loss: 1.2296 - val_accuracy: 0.5580
 Epoch 15/20
 450/450 [=====] - 6s 13ms/step - loss: 1.2610 -
 accuracy: 0.5402 - val_loss: 1.4789 - val_accuracy: 0.5040
 Epoch 16/20
 450/450 [=====] - 6s 14ms/step - loss: 1.2314 -

accuracy: 0.5507 - val_loss: 1.2562 - val_accuracy: 0.5580
 Epoch 17/20
 450/450 [=====] - 6s 12ms/step - loss: 1.1846 -
 accuracy: 0.5562 - val_loss: 1.2892 - val_accuracy: 0.5440
 Epoch 18/20
 450/450 [=====] - 6s 13ms/step - loss: 1.1665 -
 accuracy: 0.5731 - val_loss: 2.1407 - val_accuracy: 0.3940
 Epoch 19/20
 450/450 [=====] - 7s 15ms/step - loss: 1.1645 -
 accuracy: 0.5762 - val_loss: 1.2523 - val_accuracy: 0.5460
 Epoch 20/20
 450/450 [=====] - 6s 13ms/step - loss: 1.1485 -
 accuracy: 0.5840 - val_loss: 1.5985 - val_accuracy: 0.4920
 Epoch 1/20
 450/450 [=====] - 9s 14ms/step - loss: 2.1464 -
 accuracy: 0.2867 - val_loss: 2.1429 - val_accuracy: 0.2520
 Epoch 2/20
 450/450 [=====] - 7s 15ms/step - loss: 1.8704 -
 accuracy: 0.3380 - val_loss: 2.2154 - val_accuracy: 0.2780
 Epoch 3/20
 450/450 [=====] - 6s 12ms/step - loss: 1.6835 -
 accuracy: 0.3876 - val_loss: 1.5998 - val_accuracy: 0.4120
 Epoch 4/20
 450/450 [=====] - 7s 15ms/step - loss: 1.5907 -
 accuracy: 0.4142 - val_loss: 1.6238 - val_accuracy: 0.4380
 Epoch 5/20
 450/450 [=====] - 6s 12ms/step - loss: 1.5617 -
 accuracy: 0.4278 - val_loss: 1.6735 - val_accuracy: 0.4300
 Epoch 6/20
 450/450 [=====] - 7s 15ms/step - loss: 1.4707 -
 accuracy: 0.4636 - val_loss: 1.3083 - val_accuracy: 0.5240
 Epoch 7/20
 450/450 [=====] - 6s 12ms/step - loss: 1.4709 -
 accuracy: 0.4658 - val_loss: 1.2445 - val_accuracy: 0.5180
 Epoch 8/20
 450/450 [=====] - 7s 14ms/step - loss: 1.4039 -
 accuracy: 0.4838 - val_loss: 1.2951 - val_accuracy: 0.5440
 Epoch 9/20
 450/450 [=====] - 6s 12ms/step - loss: 1.3799 -
 accuracy: 0.4860 - val_loss: 1.5167 - val_accuracy: 0.4740
 Epoch 10/20
 450/450 [=====] - 6s 12ms/step - loss: 1.3333 -
 accuracy: 0.5098 - val_loss: 1.4785 - val_accuracy: 0.4700
 Epoch 11/20
 450/450 [=====] - 7s 14ms/step - loss: 1.2915 -
 accuracy: 0.5336 - val_loss: 1.4632 - val_accuracy: 0.4940
 Epoch 12/20
 450/450 [=====] - 6s 12ms/step - loss: 1.2998 -

accuracy: 0.5204 - val_loss: 1.3469 - val_accuracy: 0.5280
 Epoch 13/20
 450/450 [=====] - 6s 14ms/step - loss: 1.2550 -
 accuracy: 0.5376 - val_loss: 1.3886 - val_accuracy: 0.5200
 Epoch 14/20
 450/450 [=====] - 6s 13ms/step - loss: 1.2293 -
 accuracy: 0.5520 - val_loss: 1.4196 - val_accuracy: 0.5020
 Epoch 15/20
 450/450 [=====] - 7s 15ms/step - loss: 1.2298 -
 accuracy: 0.5513 - val_loss: 1.6946 - val_accuracy: 0.4480
 Epoch 16/20
 450/450 [=====] - 6s 13ms/step - loss: 1.2066 -
 accuracy: 0.5522 - val_loss: 1.4682 - val_accuracy: 0.5100
 Epoch 17/20
 450/450 [=====] - 6s 14ms/step - loss: 1.1609 -
 accuracy: 0.5667 - val_loss: 1.7355 - val_accuracy: 0.4380
 Epoch 18/20
 450/450 [=====] - 6s 12ms/step - loss: 1.1650 -
 accuracy: 0.5718 - val_loss: 1.2265 - val_accuracy: 0.5700
 Epoch 19/20
 450/450 [=====] - 6s 14ms/step - loss: 1.1173 -
 accuracy: 0.5989 - val_loss: 1.2678 - val_accuracy: 0.5460
 Epoch 20/20
 450/450 [=====] - 6s 12ms/step - loss: 1.1475 -
 accuracy: 0.5780 - val_loss: 1.3032 - val_accuracy: 0.5600
 Epoch 1/20
 450/450 [=====] - 10s 14ms/step - loss: 2.1304 -
 accuracy: 0.2751 - val_loss: 1.9710 - val_accuracy: 0.3080
 Epoch 2/20
 450/450 [=====] - 6s 14ms/step - loss: 1.8220 -
 accuracy: 0.3484 - val_loss: 2.5522 - val_accuracy: 0.2420
 Epoch 3/20
 450/450 [=====] - 6s 12ms/step - loss: 1.7116 -
 accuracy: 0.3884 - val_loss: 1.6028 - val_accuracy: 0.4020
 Epoch 4/20
 450/450 [=====] - 6s 14ms/step - loss: 1.6057 -
 accuracy: 0.4200 - val_loss: 1.5387 - val_accuracy: 0.4840
 Epoch 5/20
 450/450 [=====] - 6s 13ms/step - loss: 1.5308 -
 accuracy: 0.4329 - val_loss: 1.4979 - val_accuracy: 0.4260
 Epoch 6/20
 450/450 [=====] - 6s 14ms/step - loss: 1.4641 -
 accuracy: 0.4656 - val_loss: 1.4542 - val_accuracy: 0.4740
 Epoch 7/20
 450/450 [=====] - 6s 13ms/step - loss: 1.4211 -
 accuracy: 0.4816 - val_loss: 1.6783 - val_accuracy: 0.4500
 Epoch 8/20
 450/450 [=====] - 6s 14ms/step - loss: 1.3972 -

accuracy: 0.4920 - val_loss: 1.3543 - val_accuracy: 0.5020
 Epoch 9/20
 450/450 [=====] - 6s 13ms/step - loss: 1.3387 -
 accuracy: 0.5093 - val_loss: 1.2256 - val_accuracy: 0.5600
 Epoch 10/20
 450/450 [=====] - 7s 15ms/step - loss: 1.3237 -
 accuracy: 0.5142 - val_loss: 1.4115 - val_accuracy: 0.5120
 Epoch 11/20
 450/450 [=====] - 6s 12ms/step - loss: 1.2907 -
 accuracy: 0.5240 - val_loss: 1.3837 - val_accuracy: 0.5180
 Epoch 12/20
 450/450 [=====] - 7s 14ms/step - loss: 1.2733 -
 accuracy: 0.5358 - val_loss: 1.3216 - val_accuracy: 0.5440
 Epoch 13/20
 450/450 [=====] - 7s 15ms/step - loss: 1.2231 -
 accuracy: 0.5484 - val_loss: 1.1485 - val_accuracy: 0.5740
 Epoch 14/20
 450/450 [=====] - 6s 13ms/step - loss: 1.2128 -
 accuracy: 0.5538 - val_loss: 1.2171 - val_accuracy: 0.5580
 Epoch 15/20
 450/450 [=====] - 7s 15ms/step - loss: 1.1891 -
 accuracy: 0.5680 - val_loss: 1.2815 - val_accuracy: 0.5480
 Epoch 16/20
 450/450 [=====] - 6s 12ms/step - loss: 1.1667 -
 accuracy: 0.5658 - val_loss: 1.1443 - val_accuracy: 0.5980
 Epoch 17/20
 450/450 [=====] - 7s 15ms/step - loss: 1.1310 -
 accuracy: 0.5802 - val_loss: 1.2311 - val_accuracy: 0.5840
 Epoch 18/20
 450/450 [=====] - 6s 13ms/step - loss: 1.1186 -
 accuracy: 0.5900 - val_loss: 1.5499 - val_accuracy: 0.4860
 Epoch 19/20
 450/450 [=====] - 6s 14ms/step - loss: 1.1129 -
 accuracy: 0.5904 - val_loss: 1.4807 - val_accuracy: 0.5280
 Epoch 20/20
 450/450 [=====] - 6s 13ms/step - loss: 1.0809 -
 accuracy: 0.6058 - val_loss: 1.3984 - val_accuracy: 0.5300
 Epoch 1/20
 450/450 [=====] - 13s 14ms/step - loss: 2.2635 -
 accuracy: 0.2509 - val_loss: 1.8797 - val_accuracy: 0.2500
 Epoch 2/20
 450/450 [=====] - 7s 16ms/step - loss: 1.8533 -
 accuracy: 0.3260 - val_loss: 2.0486 - val_accuracy: 0.3200
 Epoch 3/20
 450/450 [=====] - 6s 14ms/step - loss: 1.6932 -
 accuracy: 0.3756 - val_loss: 2.0868 - val_accuracy: 0.2900
 Epoch 4/20
 450/450 [=====] - 7s 15ms/step - loss: 1.5910 -

accuracy: 0.4127 - val_loss: 1.4830 - val_accuracy: 0.4420
 Epoch 5/20
 450/450 [=====] - 6s 13ms/step - loss: 1.4967 -
 accuracy: 0.4440 - val_loss: 1.6783 - val_accuracy: 0.3780
 Epoch 6/20
 450/450 [=====] - 7s 15ms/step - loss: 1.4337 -
 accuracy: 0.4649 - val_loss: 1.7998 - val_accuracy: 0.3320
 Epoch 7/20
 450/450 [=====] - 6s 14ms/step - loss: 1.3942 -
 accuracy: 0.4798 - val_loss: 2.0606 - val_accuracy: 0.3800
 Epoch 8/20
 450/450 [=====] - 7s 15ms/step - loss: 1.3491 -
 accuracy: 0.4913 - val_loss: 4.2999 - val_accuracy: 0.1740
 Epoch 9/20
 450/450 [=====] - 7s 15ms/step - loss: 1.2911 -
 accuracy: 0.5178 - val_loss: 1.4714 - val_accuracy: 0.4640
 Epoch 10/20
 450/450 [=====] - 6s 13ms/step - loss: 1.2628 -
 accuracy: 0.5287 - val_loss: 1.1700 - val_accuracy: 0.5700
 Epoch 11/20
 450/450 [=====] - 7s 15ms/step - loss: 1.2219 -
 accuracy: 0.5469 - val_loss: 1.4337 - val_accuracy: 0.4740
 Epoch 12/20
 450/450 [=====] - 7s 15ms/step - loss: 1.2024 -
 accuracy: 0.5553 - val_loss: 1.5398 - val_accuracy: 0.4780
 Epoch 13/20
 450/450 [=====] - 6s 13ms/step - loss: 1.1542 -
 accuracy: 0.5718 - val_loss: 1.1351 - val_accuracy: 0.5760
 Epoch 14/20
 450/450 [=====] - 7s 16ms/step - loss: 1.1335 -
 accuracy: 0.5811 - val_loss: 1.2475 - val_accuracy: 0.5400
 Epoch 15/20
 450/450 [=====] - 6s 13ms/step - loss: 1.1182 -
 accuracy: 0.5824 - val_loss: 1.1590 - val_accuracy: 0.5940
 Epoch 16/20
 450/450 [=====] - 7s 15ms/step - loss: 1.0863 -
 accuracy: 0.6042 - val_loss: 1.1417 - val_accuracy: 0.5900
 Epoch 17/20
 450/450 [=====] - 6s 13ms/step - loss: 1.0533 -
 accuracy: 0.6127 - val_loss: 1.2680 - val_accuracy: 0.5820
 Epoch 18/20
 450/450 [=====] - 7s 15ms/step - loss: 1.0240 -
 accuracy: 0.6173 - val_loss: 1.3053 - val_accuracy: 0.5560
 Epoch 19/20
 450/450 [=====] - 6s 14ms/step - loss: 1.0034 -
 accuracy: 0.6351 - val_loss: 1.2813 - val_accuracy: 0.5240
 Epoch 20/20
 450/450 [=====] - 7s 16ms/step - loss: 0.9900 -

accuracy: 0.6449 - val_loss: 1.1672 - val_accuracy: 0.5900
 Epoch 1/20
 450/450 [=====] - 13s 16ms/step - loss: 2.2121 -
 accuracy: 0.2640 - val_loss: 1.7624 - val_accuracy: 0.3340
 Epoch 2/20
 450/450 [=====] - 6s 13ms/step - loss: 1.8658 -
 accuracy: 0.3269 - val_loss: 2.2417 - val_accuracy: 0.2860
 Epoch 3/20
 450/450 [=====] - 6s 14ms/step - loss: 1.7258 -
 accuracy: 0.3711 - val_loss: 1.4517 - val_accuracy: 0.4540
 Epoch 4/20
 450/450 [=====] - 7s 16ms/step - loss: 1.6088 -
 accuracy: 0.4007 - val_loss: 2.0467 - val_accuracy: 0.3020
 Epoch 5/20
 450/450 [=====] - 6s 14ms/step - loss: 1.5363 -
 accuracy: 0.4218 - val_loss: 1.5731 - val_accuracy: 0.4240
 Epoch 6/20
 450/450 [=====] - 7s 15ms/step - loss: 1.4784 -
 accuracy: 0.4540 - val_loss: 1.4232 - val_accuracy: 0.4700
 Epoch 7/20
 450/450 [=====] - 6s 14ms/step - loss: 1.4235 -
 accuracy: 0.4738 - val_loss: 1.6543 - val_accuracy: 0.4120
 Epoch 8/20
 450/450 [=====] - 7s 16ms/step - loss: 1.3700 -
 accuracy: 0.4902 - val_loss: 1.1865 - val_accuracy: 0.5740
 Epoch 9/20
 450/450 [=====] - 6s 14ms/step - loss: 1.3058 -
 accuracy: 0.5129 - val_loss: 1.7424 - val_accuracy: 0.3800
 Epoch 10/20
 450/450 [=====] - 7s 15ms/step - loss: 1.2750 -
 accuracy: 0.5300 - val_loss: 1.2818 - val_accuracy: 0.5120
 Epoch 11/20
 450/450 [=====] - 6s 14ms/step - loss: 1.2497 -
 accuracy: 0.5438 - val_loss: 1.2423 - val_accuracy: 0.5460
 Epoch 12/20
 450/450 [=====] - 7s 16ms/step - loss: 1.2243 -
 accuracy: 0.5464 - val_loss: 1.3115 - val_accuracy: 0.5020
 Epoch 13/20
 450/450 [=====] - 6s 13ms/step - loss: 1.1555 -
 accuracy: 0.5647 - val_loss: 1.3146 - val_accuracy: 0.5500
 Epoch 14/20
 450/450 [=====] - 7s 15ms/step - loss: 1.1490 -
 accuracy: 0.5751 - val_loss: 1.1561 - val_accuracy: 0.5840
 Epoch 15/20
 450/450 [=====] - 6s 14ms/step - loss: 1.1106 -
 accuracy: 0.5922 - val_loss: 1.1359 - val_accuracy: 0.5620
 Epoch 16/20
 450/450 [=====] - 7s 15ms/step - loss: 1.0876 -

accuracy: 0.6060 - val_loss: 1.1461 - val_accuracy: 0.5960
 Epoch 17/20
 450/450 [=====] - 7s 15ms/step - loss: 1.0595 -
 accuracy: 0.6067 - val_loss: 1.0617 - val_accuracy: 0.5940
 Epoch 18/20
 450/450 [=====] - 6s 14ms/step - loss: 1.0293 -
 accuracy: 0.6227 - val_loss: 1.3986 - val_accuracy: 0.5180
 Epoch 19/20
 450/450 [=====] - 7s 16ms/step - loss: 0.9960 -
 accuracy: 0.6313 - val_loss: 1.2651 - val_accuracy: 0.5340
 Epoch 20/20
 450/450 [=====] - 6s 14ms/step - loss: 0.9764 -
 accuracy: 0.6402 - val_loss: 1.5625 - val_accuracy: 0.4940
 Epoch 1/20
 450/450 [=====] - 13s 15ms/step - loss: 2.2939 -
 accuracy: 0.2484 - val_loss: 2.1683 - val_accuracy: 0.2580
 Epoch 2/20
 450/450 [=====] - 7s 16ms/step - loss: 1.8471 -
 accuracy: 0.3358 - val_loss: 3.1973 - val_accuracy: 0.2340
 Epoch 3/20
 450/450 [=====] - 6s 14ms/step - loss: 1.7073 -
 accuracy: 0.3767 - val_loss: 1.7782 - val_accuracy: 0.3340
 Epoch 4/20
 450/450 [=====] - 7s 16ms/step - loss: 1.6182 -
 accuracy: 0.3904 - val_loss: 1.7549 - val_accuracy: 0.3580
 Epoch 5/20
 450/450 [=====] - 6s 14ms/step - loss: 1.5464 -
 accuracy: 0.4302 - val_loss: 1.9235 - val_accuracy: 0.3160
 Epoch 6/20
 450/450 [=====] - 7s 16ms/step - loss: 1.4843 -
 accuracy: 0.4462 - val_loss: 1.6467 - val_accuracy: 0.4220
 Epoch 7/20
 450/450 [=====] - 7s 15ms/step - loss: 1.4270 -
 accuracy: 0.4691 - val_loss: 1.4983 - val_accuracy: 0.4460
 Epoch 8/20
 450/450 [=====] - 6s 14ms/step - loss: 1.3702 -
 accuracy: 0.4913 - val_loss: 1.5015 - val_accuracy: 0.4580
 Epoch 9/20
 450/450 [=====] - 7s 16ms/step - loss: 1.3195 -
 accuracy: 0.5120 - val_loss: 1.4252 - val_accuracy: 0.4980
 Epoch 10/20
 450/450 [=====] - 6s 14ms/step - loss: 1.2926 -
 accuracy: 0.5222 - val_loss: 1.2697 - val_accuracy: 0.5440
 Epoch 11/20
 450/450 [=====] - 7s 16ms/step - loss: 1.2392 -
 accuracy: 0.5451 - val_loss: 1.2201 - val_accuracy: 0.5460
 Epoch 12/20
 450/450 [=====] - 6s 14ms/step - loss: 1.2217 -

accuracy: 0.5478 - val_loss: 1.7812 - val_accuracy: 0.4320
 Epoch 13/20
 450/450 [=====] - 7s 15ms/step - loss: 1.1835 -
 accuracy: 0.5631 - val_loss: 1.2314 - val_accuracy: 0.6040
 Epoch 14/20
 450/450 [=====] - 6s 14ms/step - loss: 1.1592 -
 accuracy: 0.5776 - val_loss: 1.1544 - val_accuracy: 0.5900
 Epoch 15/20
 450/450 [=====] - 7s 16ms/step - loss: 1.1336 -
 accuracy: 0.5813 - val_loss: 1.1526 - val_accuracy: 0.5800
 Epoch 16/20
 450/450 [=====] - 6s 14ms/step - loss: 1.0931 -
 accuracy: 0.5938 - val_loss: 1.2468 - val_accuracy: 0.5540
 Epoch 17/20
 450/450 [=====] - 7s 16ms/step - loss: 1.0721 -
 accuracy: 0.6060 - val_loss: 1.3423 - val_accuracy: 0.5140
 Epoch 18/20
 450/450 [=====] - 6s 14ms/step - loss: 1.0506 -
 accuracy: 0.6147 - val_loss: 1.0784 - val_accuracy: 0.6120
 Epoch 19/20
 450/450 [=====] - 6s 14ms/step - loss: 1.0364 -
 accuracy: 0.6187 - val_loss: 1.1621 - val_accuracy: 0.5620
 Epoch 20/20
 450/450 [=====] - 7s 16ms/step - loss: 1.0056 -
 accuracy: 0.6353 - val_loss: 1.0974 - val_accuracy: 0.6100
 Epoch 1/20
 450/450 [=====] - 12s 15ms/step - loss: 2.3543 -
 accuracy: 0.2404 - val_loss: 1.7775 - val_accuracy: 0.3380
 Epoch 2/20
 450/450 [=====] - 7s 15ms/step - loss: 1.9088 -
 accuracy: 0.3011 - val_loss: 2.1206 - val_accuracy: 0.3220
 Epoch 3/20
 450/450 [=====] - 6s 13ms/step - loss: 1.7246 -
 accuracy: 0.3611 - val_loss: 2.2499 - val_accuracy: 0.3140
 Epoch 4/20
 450/450 [=====] - 7s 15ms/step - loss: 1.6331 -
 accuracy: 0.3867 - val_loss: 1.7057 - val_accuracy: 0.3740
 Epoch 5/20
 450/450 [=====] - 6s 13ms/step - loss: 1.5617 -
 accuracy: 0.4316 - val_loss: 1.3836 - val_accuracy: 0.4820
 Epoch 6/20
 450/450 [=====] - 7s 15ms/step - loss: 1.5051 -
 accuracy: 0.4416 - val_loss: 1.7223 - val_accuracy: 0.3540
 Epoch 7/20
 450/450 [=====] - 6s 13ms/step - loss: 1.4299 -
 accuracy: 0.4718 - val_loss: 1.6971 - val_accuracy: 0.4160
 Epoch 8/20
 450/450 [=====] - 7s 15ms/step - loss: 1.3887 -

accuracy: 0.4827 - val_loss: 1.8132 - val_accuracy: 0.4420
 Epoch 9/20
 450/450 [=====] - 6s 13ms/step - loss: 1.3512 -
 accuracy: 0.4962 - val_loss: 1.3524 - val_accuracy: 0.5100
 Epoch 10/20
 450/450 [=====] - 7s 15ms/step - loss: 1.2973 -
 accuracy: 0.5198 - val_loss: 1.2045 - val_accuracy: 0.5720
 Epoch 11/20
 450/450 [=====] - 6s 13ms/step - loss: 1.2749 -
 accuracy: 0.5360 - val_loss: 1.1673 - val_accuracy: 0.5820
 Epoch 12/20
 450/450 [=====] - 7s 15ms/step - loss: 1.2507 -
 accuracy: 0.5391 - val_loss: 1.6164 - val_accuracy: 0.4500
 Epoch 13/20
 450/450 [=====] - 6s 13ms/step - loss: 1.2225 -
 accuracy: 0.5398 - val_loss: 1.3208 - val_accuracy: 0.5320
 Epoch 14/20
 450/450 [=====] - 6s 13ms/step - loss: 1.1879 -
 accuracy: 0.5640 - val_loss: 1.4733 - val_accuracy: 0.4800
 Epoch 15/20
 450/450 [=====] - 6s 14ms/step - loss: 1.1834 -
 accuracy: 0.5671 - val_loss: 1.2811 - val_accuracy: 0.5500
 Epoch 16/20
 450/450 [=====] - 6s 14ms/step - loss: 1.1496 -
 accuracy: 0.5816 - val_loss: 1.4910 - val_accuracy: 0.5220
 Epoch 17/20
 450/450 [=====] - 6s 13ms/step - loss: 1.1355 -
 accuracy: 0.5838 - val_loss: 1.2061 - val_accuracy: 0.5680
 Epoch 18/20
 450/450 [=====] - 6s 13ms/step - loss: 1.1049 -
 accuracy: 0.6051 - val_loss: 2.0394 - val_accuracy: 0.4100
 Epoch 19/20
 450/450 [=====] - 7s 15ms/step - loss: 1.0898 -
 accuracy: 0.5980 - val_loss: 1.4095 - val_accuracy: 0.5660
 Epoch 20/20
 450/450 [=====] - 6s 13ms/step - loss: 1.0848 -
 accuracy: 0.6013 - val_loss: 1.5566 - val_accuracy: 0.5040
 Epoch 1/20
 450/450 [=====] - 11s 16ms/step - loss: 2.3601 -
 accuracy: 0.2380 - val_loss: 2.0943 - val_accuracy: 0.2640
 Epoch 2/20
 450/450 [=====] - 6s 13ms/step - loss: 1.8501 -
 accuracy: 0.3280 - val_loss: 1.8725 - val_accuracy: 0.3560
 Epoch 3/20
 450/450 [=====] - 6s 13ms/step - loss: 1.6969 -
 accuracy: 0.3738 - val_loss: 1.6307 - val_accuracy: 0.3880
 Epoch 4/20
 450/450 [=====] - 7s 15ms/step - loss: 1.5782 -

accuracy: 0.4171 - val_loss: 1.7486 - val_accuracy: 0.3700
Epoch 5/20
450/450 [=====] - 6s 13ms/step - loss: 1.4969 -
accuracy: 0.4387 - val_loss: 1.7330 - val_accuracy: 0.4400
Epoch 6/20
450/450 [=====] - 7s 15ms/step - loss: 1.4361 -
accuracy: 0.4642 - val_loss: 1.5843 - val_accuracy: 0.4240
Epoch 7/20
450/450 [=====] - 6s 13ms/step - loss: 1.3584 -
accuracy: 0.4927 - val_loss: 1.6401 - val_accuracy: 0.4000
Epoch 8/20
450/450 [=====] - 7s 15ms/step - loss: 1.3349 -
accuracy: 0.5042 - val_loss: 1.5856 - val_accuracy: 0.4600
Epoch 9/20
450/450 [=====] - 6s 13ms/step - loss: 1.2697 -
accuracy: 0.5267 - val_loss: 1.3777 - val_accuracy: 0.5300
Epoch 10/20
450/450 [=====] - 6s 13ms/step - loss: 1.2632 -
accuracy: 0.5320 - val_loss: 1.5808 - val_accuracy: 0.4660
Epoch 11/20
450/450 [=====] - 7s 15ms/step - loss: 1.2074 -
accuracy: 0.5533 - val_loss: 1.2768 - val_accuracy: 0.5340
Epoch 12/20
450/450 [=====] - 6s 13ms/step - loss: 1.1871 -
accuracy: 0.5562 - val_loss: 2.7392 - val_accuracy: 0.3100
Epoch 13/20
450/450 [=====] - 7s 15ms/step - loss: 1.1634 -
accuracy: 0.5771 - val_loss: 1.1452 - val_accuracy: 0.5860
Epoch 14/20
450/450 [=====] - 6s 13ms/step - loss: 1.1338 -
accuracy: 0.5884 - val_loss: 1.3755 - val_accuracy: 0.4960
Epoch 15/20
450/450 [=====] - 7s 15ms/step - loss: 1.1107 -
accuracy: 0.5918 - val_loss: 1.3395 - val_accuracy: 0.5620
Epoch 16/20
450/450 [=====] - 6s 13ms/step - loss: 1.0927 -
accuracy: 0.5991 - val_loss: 1.1603 - val_accuracy: 0.5940
Epoch 17/20
450/450 [=====] - 6s 14ms/step - loss: 1.0623 -
accuracy: 0.6071 - val_loss: 1.6576 - val_accuracy: 0.4800
Epoch 18/20
450/450 [=====] - 7s 15ms/step - loss: 1.0547 -
accuracy: 0.6171 - val_loss: 1.1621 - val_accuracy: 0.5960
Epoch 19/20
450/450 [=====] - 7s 15ms/step - loss: 1.0483 -
accuracy: 0.6251 - val_loss: 2.0975 - val_accuracy: 0.4400
Epoch 20/20
450/450 [=====] - 6s 13ms/step - loss: 1.0329 -

accuracy: 0.6180 - val_loss: 1.2791 - val_accuracy: 0.5660
 Epoch 1/20
 450/450 [=====] - 13s 16ms/step - loss: 2.3083 -
 accuracy: 0.2327 - val_loss: 1.8645 - val_accuracy: 0.2600
 Epoch 2/20
 450/450 [=====] - 6s 14ms/step - loss: 1.9012 -
 accuracy: 0.3202 - val_loss: 2.0664 - val_accuracy: 0.3040
 Epoch 3/20
 450/450 [=====] - 6s 13ms/step - loss: 1.7378 -
 accuracy: 0.3571 - val_loss: 2.1954 - val_accuracy: 0.2720
 Epoch 4/20
 450/450 [=====] - 7s 15ms/step - loss: 1.6359 -
 accuracy: 0.3947 - val_loss: 1.7279 - val_accuracy: 0.3760
 Epoch 5/20
 450/450 [=====] - 6s 13ms/step - loss: 1.5506 -
 accuracy: 0.4178 - val_loss: 1.7001 - val_accuracy: 0.4560
 Epoch 6/20
 450/450 [=====] - 7s 15ms/step - loss: 1.5031 -
 accuracy: 0.4451 - val_loss: 1.4683 - val_accuracy: 0.4460
 Epoch 7/20
 450/450 [=====] - 6s 13ms/step - loss: 1.4313 -
 accuracy: 0.4680 - val_loss: 2.4538 - val_accuracy: 0.3160
 Epoch 8/20
 450/450 [=====] - 7s 16ms/step - loss: 1.3960 -
 accuracy: 0.4860 - val_loss: 2.1260 - val_accuracy: 0.3120
 Epoch 9/20
 450/450 [=====] - 6s 13ms/step - loss: 1.3404 -
 accuracy: 0.5022 - val_loss: 1.2735 - val_accuracy: 0.5420
 Epoch 10/20
 450/450 [=====] - 7s 15ms/step - loss: 1.2904 -
 accuracy: 0.5251 - val_loss: 1.5240 - val_accuracy: 0.4600
 Epoch 11/20
 450/450 [=====] - 6s 13ms/step - loss: 1.2704 -
 accuracy: 0.5240 - val_loss: 1.4069 - val_accuracy: 0.5180
 Epoch 12/20
 450/450 [=====] - 7s 15ms/step - loss: 1.2546 -
 accuracy: 0.5364 - val_loss: 1.2891 - val_accuracy: 0.5320
 Epoch 13/20
 450/450 [=====] - 6s 13ms/step - loss: 1.2035 -
 accuracy: 0.5569 - val_loss: 1.1302 - val_accuracy: 0.5820
 Epoch 14/20
 450/450 [=====] - 7s 15ms/step - loss: 1.1805 -
 accuracy: 0.5624 - val_loss: 1.1863 - val_accuracy: 0.5980
 Epoch 15/20
 450/450 [=====] - 6s 14ms/step - loss: 1.1629 -
 accuracy: 0.5682 - val_loss: 1.2411 - val_accuracy: 0.5780
 Epoch 16/20
 450/450 [=====] - 6s 13ms/step - loss: 1.1413 -

accuracy: 0.5811 - val_loss: 1.1660 - val_accuracy: 0.5820
 Epoch 17/20
 450/450 [=====] - 7s 15ms/step - loss: 1.1285 -
 accuracy: 0.5882 - val_loss: 1.1493 - val_accuracy: 0.5960
 Epoch 18/20
 450/450 [=====] - 6s 13ms/step - loss: 1.1087 -
 accuracy: 0.5942 - val_loss: 1.5144 - val_accuracy: 0.4960
 Epoch 19/20
 450/450 [=====] - 7s 15ms/step - loss: 1.0755 -
 accuracy: 0.6062 - val_loss: 1.2106 - val_accuracy: 0.5820
 Epoch 20/20
 450/450 [=====] - 6s 14ms/step - loss: 1.0645 -
 accuracy: 0.6087 - val_loss: 1.1192 - val_accuracy: 0.6040
 Epoch 1/20
 450/450 [=====] - 10s 17ms/step - loss: 2.5665 -
 accuracy: 0.2004 - val_loss: 2.7995 - val_accuracy: 0.1320
 Epoch 2/20
 450/450 [=====] - 6s 13ms/step - loss: 2.1599 -
 accuracy: 0.2693 - val_loss: 1.6630 - val_accuracy: 0.4000
 Epoch 3/20
 450/450 [=====] - 7s 15ms/step - loss: 2.0276 -
 accuracy: 0.2713 - val_loss: 1.7291 - val_accuracy: 0.3760
 Epoch 4/20
 450/450 [=====] - 6s 13ms/step - loss: 1.8972 -
 accuracy: 0.3136 - val_loss: 1.7352 - val_accuracy: 0.3560
 Epoch 5/20
 450/450 [=====] - 7s 15ms/step - loss: 1.8038 -
 accuracy: 0.3362 - val_loss: 1.6340 - val_accuracy: 0.4040
 Epoch 6/20
 450/450 [=====] - 6s 14ms/step - loss: 1.7673 -
 accuracy: 0.3569 - val_loss: 2.5412 - val_accuracy: 0.2460
 Epoch 7/20
 450/450 [=====] - 7s 15ms/step - loss: 1.7641 -
 accuracy: 0.3571 - val_loss: 1.5293 - val_accuracy: 0.4240
 Epoch 8/20
 450/450 [=====] - 6s 13ms/step - loss: 1.7103 -
 accuracy: 0.3671 - val_loss: 1.6540 - val_accuracy: 0.3840
 Epoch 9/20
 450/450 [=====] - 7s 15ms/step - loss: 1.6678 -
 accuracy: 0.3769 - val_loss: 1.4944 - val_accuracy: 0.4340
 Epoch 10/20
 450/450 [=====] - 6s 14ms/step - loss: 1.6567 -
 accuracy: 0.3891 - val_loss: 1.6068 - val_accuracy: 0.3820
 Epoch 11/20
 450/450 [=====] - 6s 14ms/step - loss: 1.6441 -
 accuracy: 0.3880 - val_loss: 1.6601 - val_accuracy: 0.3780
 Epoch 12/20
 450/450 [=====] - 7s 15ms/step - loss: 1.6332 -

accuracy: 0.3947 - val_loss: 1.5207 - val_accuracy: 0.3900
 Epoch 13/20
 450/450 [=====] - 6s 13ms/step - loss: 1.6144 -
 accuracy: 0.3969 - val_loss: 1.4486 - val_accuracy: 0.4720
 Epoch 14/20
 450/450 [=====] - 7s 16ms/step - loss: 1.5744 -
 accuracy: 0.4033 - val_loss: 1.7728 - val_accuracy: 0.3720
 Epoch 15/20
 450/450 [=====] - 6s 13ms/step - loss: 1.5797 -
 accuracy: 0.4164 - val_loss: 1.5218 - val_accuracy: 0.4240
 Epoch 16/20
 450/450 [=====] - 7s 16ms/step - loss: 1.5662 -
 accuracy: 0.4129 - val_loss: 1.5139 - val_accuracy: 0.4200
 Epoch 17/20
 450/450 [=====] - 6s 13ms/step - loss: 1.5360 -
 accuracy: 0.4260 - val_loss: 1.4518 - val_accuracy: 0.4660
 Epoch 18/20
 450/450 [=====] - 7s 16ms/step - loss: 1.5035 -
 accuracy: 0.4391 - val_loss: 1.4066 - val_accuracy: 0.4580
 Epoch 19/20
 450/450 [=====] - 7s 15ms/step - loss: 1.5192 -
 accuracy: 0.4420 - val_loss: 1.6251 - val_accuracy: 0.4160
 Epoch 20/20
 450/450 [=====] - 6s 13ms/step - loss: 1.4908 -
 accuracy: 0.4498 - val_loss: 1.5508 - val_accuracy: 0.4220
 Epoch 1/20
 450/450 [=====] - 9s 14ms/step - loss: 2.5868 -
 accuracy: 0.2100 - val_loss: 2.1528 - val_accuracy: 0.1960
 Epoch 2/20
 450/450 [=====] - 6s 13ms/step - loss: 2.2179 -
 accuracy: 0.2629 - val_loss: 1.6317 - val_accuracy: 0.3780
 Epoch 3/20
 450/450 [=====] - 7s 15ms/step - loss: 1.9905 -
 accuracy: 0.2991 - val_loss: 1.7236 - val_accuracy: 0.3360
 Epoch 4/20
 450/450 [=====] - 6s 13ms/step - loss: 1.8963 -
 accuracy: 0.3136 - val_loss: 1.7423 - val_accuracy: 0.3460
 Epoch 5/20
 450/450 [=====] - 7s 15ms/step - loss: 1.8472 -
 accuracy: 0.3282 - val_loss: 1.5677 - val_accuracy: 0.4000
 Epoch 6/20
 450/450 [=====] - 6s 13ms/step - loss: 1.7843 -
 accuracy: 0.3484 - val_loss: 1.7485 - val_accuracy: 0.3920
 Epoch 7/20
 450/450 [=====] - 7s 15ms/step - loss: 1.7513 -
 accuracy: 0.3647 - val_loss: 1.7023 - val_accuracy: 0.3800
 Epoch 8/20
 450/450 [=====] - 6s 13ms/step - loss: 1.6989 -

accuracy: 0.3642 - val_loss: 1.7353 - val_accuracy: 0.3620
 Epoch 9/20
 450/450 [=====] - 7s 15ms/step - loss: 1.6759 -
 accuracy: 0.3716 - val_loss: 1.7685 - val_accuracy: 0.3800
 Epoch 10/20
 450/450 [=====] - 6s 13ms/step - loss: 1.6589 -
 accuracy: 0.3851 - val_loss: 1.3760 - val_accuracy: 0.4920
 Epoch 11/20
 450/450 [=====] - 7s 16ms/step - loss: 1.6415 -
 accuracy: 0.3884 - val_loss: 1.4699 - val_accuracy: 0.4480
 Epoch 12/20
 450/450 [=====] - 6s 14ms/step - loss: 1.6010 -
 accuracy: 0.4104 - val_loss: 1.4315 - val_accuracy: 0.4620
 Epoch 13/20
 450/450 [=====] - 6s 14ms/step - loss: 1.5951 -
 accuracy: 0.4011 - val_loss: 1.4784 - val_accuracy: 0.4180
 Epoch 14/20
 450/450 [=====] - 7s 14ms/step - loss: 1.5801 -
 accuracy: 0.4069 - val_loss: 1.4735 - val_accuracy: 0.4640
 Epoch 15/20
 450/450 [=====] - 6s 13ms/step - loss: 1.5677 -
 accuracy: 0.4193 - val_loss: 1.3822 - val_accuracy: 0.4620
 Epoch 16/20
 450/450 [=====] - 7s 15ms/step - loss: 1.5554 -
 accuracy: 0.4267 - val_loss: 1.3892 - val_accuracy: 0.4860
 Epoch 17/20
 450/450 [=====] - 7s 16ms/step - loss: 1.5514 -
 accuracy: 0.4191 - val_loss: 1.3401 - val_accuracy: 0.5080
 Epoch 18/20
 450/450 [=====] - 6s 13ms/step - loss: 1.5021 -
 accuracy: 0.4349 - val_loss: 1.3720 - val_accuracy: 0.5160
 Epoch 19/20
 450/450 [=====] - 6s 14ms/step - loss: 1.4916 -
 accuracy: 0.4416 - val_loss: 1.3114 - val_accuracy: 0.5140
 Epoch 20/20
 450/450 [=====] - 7s 16ms/step - loss: 1.4967 -
 accuracy: 0.4387 - val_loss: 1.2931 - val_accuracy: 0.5140
 Epoch 1/20
 450/450 [=====] - 10s 17ms/step - loss: 2.5633 -
 accuracy: 0.2144 - val_loss: 2.0449 - val_accuracy: 0.3020
 Epoch 2/20
 450/450 [=====] - 7s 15ms/step - loss: 2.1736 -
 accuracy: 0.2727 - val_loss: 1.6892 - val_accuracy: 0.3460
 Epoch 3/20
 450/450 [=====] - 6s 13ms/step - loss: 2.0044 -
 accuracy: 0.3058 - val_loss: 2.1766 - val_accuracy: 0.2500
 Epoch 4/20
 450/450 [=====] - 6s 13ms/step - loss: 1.8685 -

accuracy: 0.3278 - val_loss: 1.5869 - val_accuracy: 0.3880
 Epoch 5/20
 450/450 [=====] - 7s 15ms/step - loss: 1.8238 -
 accuracy: 0.3358 - val_loss: 1.7463 - val_accuracy: 0.3340
 Epoch 6/20
 450/450 [=====] - 6s 13ms/step - loss: 1.7783 -
 accuracy: 0.3460 - val_loss: 1.7712 - val_accuracy: 0.3720
 Epoch 7/20
 450/450 [=====] - 7s 15ms/step - loss: 1.7539 -
 accuracy: 0.3636 - val_loss: 1.7977 - val_accuracy: 0.3500
 Epoch 8/20
 450/450 [=====] - 6s 13ms/step - loss: 1.7213 -
 accuracy: 0.3693 - val_loss: 2.0569 - val_accuracy: 0.3560
 Epoch 9/20
 450/450 [=====] - 7s 16ms/step - loss: 1.6760 -
 accuracy: 0.3740 - val_loss: 1.4328 - val_accuracy: 0.4520
 Epoch 10/20
 450/450 [=====] - 6s 13ms/step - loss: 1.6529 -
 accuracy: 0.3871 - val_loss: 1.8422 - val_accuracy: 0.3800
 Epoch 11/20
 450/450 [=====] - 7s 15ms/step - loss: 1.6437 -
 accuracy: 0.3967 - val_loss: 1.6570 - val_accuracy: 0.4060
 Epoch 12/20
 450/450 [=====] - 6s 13ms/step - loss: 1.6272 -
 accuracy: 0.3929 - val_loss: 1.5762 - val_accuracy: 0.4140
 Epoch 13/20
 450/450 [=====] - 7s 15ms/step - loss: 1.5855 -
 accuracy: 0.4127 - val_loss: 1.5994 - val_accuracy: 0.3940
 Epoch 14/20
 450/450 [=====] - 7s 15ms/step - loss: 1.5985 -
 accuracy: 0.4113 - val_loss: 1.6559 - val_accuracy: 0.4180
 Epoch 15/20
 450/450 [=====] - 6s 13ms/step - loss: 1.5858 -
 accuracy: 0.4049 - val_loss: 1.4932 - val_accuracy: 0.4440
 Epoch 16/20
 450/450 [=====] - 7s 15ms/step - loss: 1.5470 -
 accuracy: 0.4116 - val_loss: 1.5915 - val_accuracy: 0.4500
 Epoch 17/20
 450/450 [=====] - 6s 14ms/step - loss: 1.5313 -
 accuracy: 0.4236 - val_loss: 1.5504 - val_accuracy: 0.4400
 Epoch 18/20
 450/450 [=====] - 6s 13ms/step - loss: 1.4994 -
 accuracy: 0.4453 - val_loss: 1.4782 - val_accuracy: 0.4360
 Epoch 19/20
 450/450 [=====] - 6s 13ms/step - loss: 1.4963 -
 accuracy: 0.4382 - val_loss: 1.6498 - val_accuracy: 0.4000
 Epoch 20/20
 450/450 [=====] - 7s 15ms/step - loss: 1.5019 -


```

accuracy: 0.4404 - val_loss: 1.3240 - val_accuracy: 0.5000
Epoch 1/20
450/450 [=====] - 13s 15ms/step - loss: 2.5475 -
accuracy: 0.2322 - val_loss: 1.9661 - val_accuracy: 0.2860
Epoch 2/20
450/450 [=====] - 7s 16ms/step - loss: 2.0037 -
accuracy: 0.2927 - val_loss: 1.7624 - val_accuracy: 0.3040
Epoch 3/20
450/450 [=====] - 6s 14ms/step - loss: 1.8329 -
accuracy: 0.3260 - val_loss: 1.6290 - val_accuracy: 0.3680
Epoch 4/20
450/450 [=====] - 7s 16ms/step - loss: 1.7437 -
accuracy: 0.3436 - val_loss: 1.7639 - val_accuracy: 0.3800
Epoch 5/20
450/450 [=====] - 6s 14ms/step - loss: 1.6700 -
accuracy: 0.3731 - val_loss: 1.6377 - val_accuracy: 0.3880
Epoch 6/20
450/450 [=====] - 7s 15ms/step - loss: 1.6099 -
accuracy: 0.3949 - val_loss: 1.7196 - val_accuracy: 0.3780
Epoch 7/20
450/450 [=====] - 7s 16ms/step - loss: 1.6016 -
accuracy: 0.3953 - val_loss: 1.5057 - val_accuracy: 0.4420
Epoch 8/20
450/450 [=====] - 7s 16ms/step - loss: 1.5423 -
accuracy: 0.4182 - val_loss: 1.4811 - val_accuracy: 0.4620
Epoch 9/20
450/450 [=====] - 6s 14ms/step - loss: 1.5071 -
accuracy: 0.4367 - val_loss: 1.5056 - val_accuracy: 0.4460
Epoch 10/20
450/450 [=====] - 7s 16ms/step - loss: 1.4809 -
accuracy: 0.4460 - val_loss: 1.4777 - val_accuracy: 0.4440
Epoch 11/20
450/450 [=====] - 6s 14ms/step - loss: 1.4764 -
accuracy: 0.4491 - val_loss: 1.5137 - val_accuracy: 0.4460
Epoch 12/20
450/450 [=====] - 7s 14ms/step - loss: 1.4597 -
accuracy: 0.4536 - val_loss: 1.3566 - val_accuracy: 0.4760
Epoch 13/20
450/450 [=====] - 7s 16ms/step - loss: 1.4153 -
accuracy: 0.4813 - val_loss: 1.5391 - val_accuracy: 0.4540
Epoch 14/20
450/450 [=====] - 7s 15ms/step - loss: 1.3857 -
accuracy: 0.4780 - val_loss: 1.6236 - val_accuracy: 0.3980
Epoch 15/20
450/450 [=====] - 7s 15ms/step - loss: 1.3833 -
accuracy: 0.4831 - val_loss: 1.5213 - val_accuracy: 0.4600
Epoch 16/20
450/450 [=====] - 7s 16ms/step - loss: 1.3632 -

```

accuracy: 0.4909 - val_loss: 1.2884 - val_accuracy: 0.5280
 Epoch 17/20
 450/450 [=====] - 6s 14ms/step - loss: 1.3263 -
 accuracy: 0.5042 - val_loss: 1.3422 - val_accuracy: 0.5140
 Epoch 18/20
 450/450 [=====] - 7s 16ms/step - loss: 1.2948 -
 accuracy: 0.5133 - val_loss: 1.5975 - val_accuracy: 0.4600
 Epoch 19/20
 450/450 [=====] - 7s 15ms/step - loss: 1.3082 -
 accuracy: 0.5164 - val_loss: 1.3008 - val_accuracy: 0.5180
 Epoch 20/20
 450/450 [=====] - 6s 14ms/step - loss: 1.2771 -
 accuracy: 0.5191 - val_loss: 1.0809 - val_accuracy: 0.6100
 Epoch 1/20
 450/450 [=====] - 13s 17ms/step - loss: 2.5657 -
 accuracy: 0.2100 - val_loss: 2.5462 - val_accuracy: 0.1740
 Epoch 2/20
 450/450 [=====] - 7s 15ms/step - loss: 2.0054 -
 accuracy: 0.2873 - val_loss: 2.4650 - val_accuracy: 0.2320
 Epoch 3/20
 450/450 [=====] - 7s 16ms/step - loss: 1.8632 -
 accuracy: 0.3136 - val_loss: 1.7847 - val_accuracy: 0.3580
 Epoch 4/20
 450/450 [=====] - 6s 14ms/step - loss: 1.7553 -
 accuracy: 0.3382 - val_loss: 1.5284 - val_accuracy: 0.4060
 Epoch 5/20
 450/450 [=====] - 7s 15ms/step - loss: 1.7323 -
 accuracy: 0.3460 - val_loss: 1.4928 - val_accuracy: 0.4180
 Epoch 6/20
 450/450 [=====] - 7s 16ms/step - loss: 1.6570 -
 accuracy: 0.3769 - val_loss: 1.8045 - val_accuracy: 0.3000
 Epoch 7/20
 450/450 [=====] - 7s 16ms/step - loss: 1.6327 -
 accuracy: 0.3898 - val_loss: 1.3850 - val_accuracy: 0.4760
 Epoch 8/20
 450/450 [=====] - 7s 14ms/step - loss: 1.5809 -
 accuracy: 0.4078 - val_loss: 1.7877 - val_accuracy: 0.3380
 Epoch 9/20
 450/450 [=====] - 7s 14ms/step - loss: 1.5444 -
 accuracy: 0.4078 - val_loss: 1.4520 - val_accuracy: 0.4520
 Epoch 10/20
 450/450 [=====] - 7s 16ms/step - loss: 1.5166 -
 accuracy: 0.4187 - val_loss: 1.3696 - val_accuracy: 0.5100
 Epoch 11/20
 450/450 [=====] - 7s 15ms/step - loss: 1.4916 -
 accuracy: 0.4422 - val_loss: 1.5882 - val_accuracy: 0.4400
 Epoch 12/20
 450/450 [=====] - 7s 16ms/step - loss: 1.4805 -

accuracy: 0.4471 - val_loss: 1.3564 - val_accuracy: 0.4980
 Epoch 13/20
 450/450 [=====] - 7s 16ms/step - loss: 1.4377 -
 accuracy: 0.4627 - val_loss: 1.5282 - val_accuracy: 0.4320
 Epoch 14/20
 450/450 [=====] - 7s 15ms/step - loss: 1.4299 -
 accuracy: 0.4609 - val_loss: 1.6250 - val_accuracy: 0.4520
 Epoch 15/20
 450/450 [=====] - 7s 17ms/step - loss: 1.4097 -
 accuracy: 0.4822 - val_loss: 1.3405 - val_accuracy: 0.5060
 Epoch 16/20
 450/450 [=====] - 6s 14ms/step - loss: 1.3898 -
 accuracy: 0.4829 - val_loss: 1.6742 - val_accuracy: 0.4020
 Epoch 17/20
 450/450 [=====] - 6s 14ms/step - loss: 1.3480 -
 accuracy: 0.4962 - val_loss: 1.2651 - val_accuracy: 0.5220
 Epoch 18/20
 450/450 [=====] - 7s 16ms/step - loss: 1.3302 -
 accuracy: 0.5049 - val_loss: 1.4665 - val_accuracy: 0.4440
 Epoch 19/20
 450/450 [=====] - 7s 16ms/step - loss: 1.3113 -
 accuracy: 0.5200 - val_loss: 1.1807 - val_accuracy: 0.5600
 Epoch 20/20
 450/450 [=====] - 7s 15ms/step - loss: 1.3057 -
 accuracy: 0.5149 - val_loss: 1.4363 - val_accuracy: 0.5060
 Epoch 1/20
 450/450 [=====] - 14s 15ms/step - loss: 2.5873 -
 accuracy: 0.1973 - val_loss: 2.6974 - val_accuracy: 0.2340
 Epoch 2/20
 450/450 [=====] - 7s 16ms/step - loss: 2.0393 -
 accuracy: 0.2764 - val_loss: 1.9451 - val_accuracy: 0.2580
 Epoch 3/20
 450/450 [=====] - 7s 16ms/step - loss: 1.8640 -
 accuracy: 0.3107 - val_loss: 1.6483 - val_accuracy: 0.3540
 Epoch 4/20
 450/450 [=====] - 7s 15ms/step - loss: 1.7776 -
 accuracy: 0.3322 - val_loss: 1.5553 - val_accuracy: 0.3860
 Epoch 5/20
 450/450 [=====] - 8s 17ms/step - loss: 1.7089 -
 accuracy: 0.3567 - val_loss: 2.2110 - val_accuracy: 0.2840
 Epoch 6/20
 450/450 [=====] - 7s 17ms/step - loss: 1.6678 -
 accuracy: 0.3760 - val_loss: 1.6545 - val_accuracy: 0.3700
 Epoch 7/20
 450/450 [=====] - 7s 16ms/step - loss: 1.6316 -
 accuracy: 0.3891 - val_loss: 1.6232 - val_accuracy: 0.4020
 Epoch 8/20
 450/450 [=====] - 7s 15ms/step - loss: 1.5892 -

accuracy: 0.4027 - val_loss: 1.4825 - val_accuracy: 0.4700
Epoch 9/20
450/450 [=====] - 7s 17ms/step - loss: 1.5640 -
accuracy: 0.4118 - val_loss: 1.9248 - val_accuracy: 0.3960
Epoch 10/20
450/450 [=====] - 7s 15ms/step - loss: 1.5277 -
accuracy: 0.4244 - val_loss: 1.5967 - val_accuracy: 0.3860
Epoch 11/20
450/450 [=====] - 7s 15ms/step - loss: 1.4979 -
accuracy: 0.4462 - val_loss: 1.3474 - val_accuracy: 0.5100
Epoch 12/20
450/450 [=====] - 7s 15ms/step - loss: 1.4672 -
accuracy: 0.4422 - val_loss: 1.3908 - val_accuracy: 0.4540
Epoch 13/20
450/450 [=====] - 7s 17ms/step - loss: 1.4392 -
accuracy: 0.4698 - val_loss: 1.3484 - val_accuracy: 0.5020
Epoch 14/20
450/450 [=====] - 7s 15ms/step - loss: 1.4286 -
accuracy: 0.4560 - val_loss: 1.3638 - val_accuracy: 0.4860
Epoch 15/20
450/450 [=====] - 7s 15ms/step - loss: 1.4078 -
accuracy: 0.4756 - val_loss: 1.3330 - val_accuracy: 0.4920
Epoch 16/20
450/450 [=====] - 7s 17ms/step - loss: 1.3776 -
accuracy: 0.4820 - val_loss: 1.4396 - val_accuracy: 0.4880
Epoch 17/20
450/450 [=====] - 7s 15ms/step - loss: 1.3573 -
accuracy: 0.4913 - val_loss: 1.1705 - val_accuracy: 0.5580
Epoch 18/20
450/450 [=====] - 8s 17ms/step - loss: 1.3345 -
accuracy: 0.4984 - val_loss: 1.2380 - val_accuracy: 0.5140
Epoch 19/20
450/450 [=====] - 7s 16ms/step - loss: 1.3246 -
accuracy: 0.5036 - val_loss: 1.3344 - val_accuracy: 0.5100
Epoch 20/20
450/450 [=====] - 7s 15ms/step - loss: 1.3069 -
accuracy: 0.5136 - val_loss: 1.4068 - val_accuracy: 0.4900
Epoch 1/20
450/450 [=====] - 12s 17ms/step - loss: 2.5791 -
accuracy: 0.2207 - val_loss: 1.8342 - val_accuracy: 0.3040
Epoch 2/20
450/450 [=====] - 7s 15ms/step - loss: 2.0722 -
accuracy: 0.2864 - val_loss: 2.2624 - val_accuracy: 0.2480
Epoch 3/20
450/450 [=====] - 7s 16ms/step - loss: 1.8987 -
accuracy: 0.3171 - val_loss: 2.6109 - val_accuracy: 0.2660
Epoch 4/20
450/450 [=====] - 7s 15ms/step - loss: 1.7938 -

accuracy: 0.3362 - val_loss: 1.6085 - val_accuracy: 0.3980
Epoch 5/20
450/450 [=====] - 7s 16ms/step - loss: 1.7234 -
accuracy: 0.3649 - val_loss: 1.5861 - val_accuracy: 0.4420
Epoch 6/20
450/450 [=====] - 7s 16ms/step - loss: 1.6852 -
accuracy: 0.3780 - val_loss: 1.5338 - val_accuracy: 0.4000
Epoch 7/20
450/450 [=====] - 6s 14ms/step - loss: 1.6515 -
accuracy: 0.3909 - val_loss: 1.5358 - val_accuracy: 0.4180
Epoch 8/20
450/450 [=====] - 7s 16ms/step - loss: 1.6000 -
accuracy: 0.4051 - val_loss: 1.4676 - val_accuracy: 0.4420
Epoch 9/20
450/450 [=====] - 7s 16ms/step - loss: 1.5683 -
accuracy: 0.4142 - val_loss: 1.4970 - val_accuracy: 0.4600
Epoch 10/20
450/450 [=====] - 6s 14ms/step - loss: 1.5530 -
accuracy: 0.4240 - val_loss: 1.5933 - val_accuracy: 0.3920
Epoch 11/20
450/450 [=====] - 6s 14ms/step - loss: 1.5369 -
accuracy: 0.4238 - val_loss: 1.4796 - val_accuracy: 0.4700
Epoch 12/20
450/450 [=====] - 7s 16ms/step - loss: 1.5044 -
accuracy: 0.4396 - val_loss: 1.5923 - val_accuracy: 0.4040
Epoch 13/20
450/450 [=====] - 6s 14ms/step - loss: 1.4762 -
accuracy: 0.4464 - val_loss: 1.3921 - val_accuracy: 0.4740
Epoch 14/20
450/450 [=====] - 7s 16ms/step - loss: 1.4877 -
accuracy: 0.4509 - val_loss: 1.2703 - val_accuracy: 0.5260
Epoch 15/20
450/450 [=====] - 6s 14ms/step - loss: 1.4654 -
accuracy: 0.4602 - val_loss: 1.2739 - val_accuracy: 0.5120
Epoch 16/20
450/450 [=====] - 7s 15ms/step - loss: 1.4482 -
accuracy: 0.4549 - val_loss: 1.4724 - val_accuracy: 0.4960
Epoch 17/20
450/450 [=====] - 7s 16ms/step - loss: 1.4216 -
accuracy: 0.4776 - val_loss: 1.2853 - val_accuracy: 0.5240
Epoch 18/20
450/450 [=====] - 6s 14ms/step - loss: 1.4066 -
accuracy: 0.4816 - val_loss: 2.5824 - val_accuracy: 0.3540
Epoch 19/20
450/450 [=====] - 7s 16ms/step - loss: 1.3969 -
accuracy: 0.4782 - val_loss: 1.3408 - val_accuracy: 0.5160
Epoch 20/20
450/450 [=====] - 6s 14ms/step - loss: 1.3933 -

accuracy: 0.4902 - val_loss: 1.2774 - val_accuracy: 0.5500
 Epoch 1/20
 450/450 [=====] - 12s 16ms/step - loss: 2.5922 -
 accuracy: 0.2224 - val_loss: 1.9549 - val_accuracy: 0.2660
 Epoch 2/20
 450/450 [=====] - 7s 16ms/step - loss: 2.0475 -
 accuracy: 0.2849 - val_loss: 1.9337 - val_accuracy: 0.3220
 Epoch 3/20
 450/450 [=====] - 6s 14ms/step - loss: 1.8982 -
 accuracy: 0.3087 - val_loss: 1.6656 - val_accuracy: 0.3800
 Epoch 4/20
 450/450 [=====] - 7s 16ms/step - loss: 1.8048 -
 accuracy: 0.3402 - val_loss: 2.5178 - val_accuracy: 0.2320
 Epoch 5/20
 450/450 [=====] - 6s 14ms/step - loss: 1.7365 -
 accuracy: 0.3658 - val_loss: 3.3059 - val_accuracy: 0.1440
 Epoch 6/20
 450/450 [=====] - 6s 14ms/step - loss: 1.6624 -
 accuracy: 0.3840 - val_loss: 1.9980 - val_accuracy: 0.3060
 Epoch 7/20
 450/450 [=====] - 7s 16ms/step - loss: 1.6347 -
 accuracy: 0.3853 - val_loss: 2.1807 - val_accuracy: 0.2860
 Epoch 8/20
 450/450 [=====] - 7s 15ms/step - loss: 1.5959 -
 accuracy: 0.4122 - val_loss: 1.5730 - val_accuracy: 0.4040
 Epoch 9/20
 450/450 [=====] - 7s 15ms/step - loss: 1.5646 -
 accuracy: 0.4184 - val_loss: 1.4593 - val_accuracy: 0.4920
 Epoch 10/20
 450/450 [=====] - 6s 14ms/step - loss: 1.5366 -
 accuracy: 0.4271 - val_loss: 2.6299 - val_accuracy: 0.3360
 Epoch 11/20
 450/450 [=====] - 7s 16ms/step - loss: 1.5113 -
 accuracy: 0.4444 - val_loss: 2.8390 - val_accuracy: 0.2900
 Epoch 12/20
 450/450 [=====] - 6s 14ms/step - loss: 1.4816 -
 accuracy: 0.4518 - val_loss: 1.3703 - val_accuracy: 0.4900
 Epoch 13/20
 450/450 [=====] - 7s 16ms/step - loss: 1.4534 -
 accuracy: 0.4611 - val_loss: 2.0454 - val_accuracy: 0.3460
 Epoch 14/20
 450/450 [=====] - 6s 14ms/step - loss: 1.4404 -
 accuracy: 0.4598 - val_loss: 1.3010 - val_accuracy: 0.5260
 Epoch 15/20
 450/450 [=====] - 6s 14ms/step - loss: 1.4310 -
 accuracy: 0.4793 - val_loss: 1.3318 - val_accuracy: 0.5240
 Epoch 16/20
 450/450 [=====] - 7s 16ms/step - loss: 1.4079 -

accuracy: 0.4702 - val_loss: 1.3448 - val_accuracy: 0.5280
 Epoch 17/20
 450/450 [=====] - 6s 14ms/step - loss: 1.3968 -
 accuracy: 0.4804 - val_loss: 1.3468 - val_accuracy: 0.5020
 Epoch 18/20
 450/450 [=====] - 7s 16ms/step - loss: 1.3760 -
 accuracy: 0.4844 - val_loss: 1.1821 - val_accuracy: 0.5540
 Epoch 19/20
 450/450 [=====] - 6s 14ms/step - loss: 1.3759 -
 accuracy: 0.4938 - val_loss: 1.5980 - val_accuracy: 0.4240
 Epoch 20/20
 450/450 [=====] - 7s 16ms/step - loss: 1.3675 -
 accuracy: 0.4978 - val_loss: 1.8367 - val_accuracy: 0.3760
 Epoch 1/20
 450/450 [=====] - 12s 15ms/step - loss: 2.5707 -
 accuracy: 0.2224 - val_loss: 4.9271 - val_accuracy: 0.1920
 Epoch 2/20
 450/450 [=====] - 7s 15ms/step - loss: 2.0397 -
 accuracy: 0.2838 - val_loss: 2.7546 - val_accuracy: 0.2740
 Epoch 3/20
 450/450 [=====] - 6s 14ms/step - loss: 1.8596 -
 accuracy: 0.3176 - val_loss: 1.7542 - val_accuracy: 0.3240
 Epoch 4/20
 450/450 [=====] - 7s 16ms/step - loss: 1.7611 -
 accuracy: 0.3373 - val_loss: 1.7150 - val_accuracy: 0.3300
 Epoch 5/20
 450/450 [=====] - 7s 15ms/step - loss: 1.7109 -
 accuracy: 0.3618 - val_loss: 1.8955 - val_accuracy: 0.3340
 Epoch 6/20
 450/450 [=====] - 6s 14ms/step - loss: 1.6648 -
 accuracy: 0.3858 - val_loss: 1.4522 - val_accuracy: 0.4320
 Epoch 7/20
 450/450 [=====] - 6s 13ms/step - loss: 1.6135 -
 accuracy: 0.3918 - val_loss: 1.5633 - val_accuracy: 0.3880
 Epoch 8/20
 450/450 [=====] - 7s 16ms/step - loss: 1.5989 -
 accuracy: 0.4089 - val_loss: 1.4988 - val_accuracy: 0.4460
 Epoch 9/20
 450/450 [=====] - 7s 15ms/step - loss: 1.5539 -
 accuracy: 0.4207 - val_loss: 1.5793 - val_accuracy: 0.4220
 Epoch 10/20
 450/450 [=====] - 6s 14ms/step - loss: 1.5289 -
 accuracy: 0.4300 - val_loss: 1.8832 - val_accuracy: 0.3600
 Epoch 11/20
 450/450 [=====] - 7s 16ms/step - loss: 1.5053 -
 accuracy: 0.4380 - val_loss: 2.0407 - val_accuracy: 0.3360
 Epoch 12/20
 450/450 [=====] - 7s 16ms/step - loss: 1.4815 -

```

accuracy: 0.4473 - val_loss: 2.8356 - val_accuracy: 0.3020
Epoch 13/20
450/450 [=====] - 7s 15ms/step - loss: 1.4644 -
accuracy: 0.4593 - val_loss: 1.5464 - val_accuracy: 0.4240
Epoch 14/20
450/450 [=====] - 6s 14ms/step - loss: 1.4515 -
accuracy: 0.4553 - val_loss: 1.2489 - val_accuracy: 0.5260
Epoch 15/20
450/450 [=====] - 7s 16ms/step - loss: 1.4290 -
accuracy: 0.4724 - val_loss: 1.3826 - val_accuracy: 0.5120
Epoch 16/20
450/450 [=====] - 6s 14ms/step - loss: 1.3940 -
accuracy: 0.4822 - val_loss: 1.4307 - val_accuracy: 0.4900
Epoch 17/20
450/450 [=====] - 6s 14ms/step - loss: 1.3936 -
accuracy: 0.4793 - val_loss: 1.3970 - val_accuracy: 0.4600
Epoch 18/20
450/450 [=====] - 7s 16ms/step - loss: 1.3632 -
accuracy: 0.4960 - val_loss: 1.3327 - val_accuracy: 0.5140
Epoch 19/20
450/450 [=====] - 6s 14ms/step - loss: 1.3593 -
accuracy: 0.5033 - val_loss: 1.2468 - val_accuracy: 0.5280
Epoch 20/20
450/450 [=====] - 6s 14ms/step - loss: 1.3635 -
accuracy: 0.5027 - val_loss: 1.3052 - val_accuracy: 0.5040
Best Parameters: {'drop_rate': 0.5, 'optimizer': 'rmsprop', 'alpha': 0.5}
Best Validation Performance: 0

```

```
[ ]: print(best_val_performance)
```

```
0.527999997138977
```

```
parameters: {'drop_rate': 0.5, 'optimizer': 'rmsprop', 'alpha': 0.5} validation accuracy:0.528
```

```
[ ]: model_3_7_3 = YourModel(name='q3_7_3',
                                feature_maps=32,
                                num_classes=data_manager.n_classes,
                                num_blocks=4,
                                drop_rate=0.5,
                                batch_norm=True,
                                is_augmentation=True,
                                use_skip=True,
                                optimizer='rmsprop',
                                num_epochs=100,
                                learning_rate=0.001)

                                # Build and train the model
model_3_7_3.build_cnn()
```



```
model_3_7_3.fit(data_manager, use_mixup=True, alpha=0.5)
```

Epoch 1/100

450/450 [=====] - 10s 13ms/step - loss: 2.5998 - accuracy: 0.2149 - val_loss: 2.2869 - val_accuracy: 0.2760

Epoch 2/100

450/450 [=====] - 6s 14ms/step - loss: 2.0635 - accuracy: 0.2860 - val_loss: 4.0022 - val_accuracy: 0.2400

Epoch 3/100

450/450 [=====] - 6s 12ms/step - loss: 1.9129 - accuracy: 0.3156 - val_loss: 2.3024 - val_accuracy: 0.3080

Epoch 4/100

450/450 [=====] - 6s 14ms/step - loss: 1.7950 - accuracy: 0.3300 - val_loss: 1.6124 - val_accuracy: 0.3740

Epoch 5/100

450/450 [=====] - 6s 12ms/step - loss: 1.7564 - accuracy: 0.3529 - val_loss: 1.6863 - val_accuracy: 0.3440

Epoch 6/100

450/450 [=====] - 6s 13ms/step - loss: 1.6869 - accuracy: 0.3771 - val_loss: 1.5898 - val_accuracy: 0.3840

Epoch 7/100

450/450 [=====] - 6s 14ms/step - loss: 1.6699 - accuracy: 0.3789 - val_loss: 1.5121 - val_accuracy: 0.4100

Epoch 8/100

450/450 [=====] - 6s 13ms/step - loss: 1.6155 - accuracy: 0.4049 - val_loss: 1.9935 - val_accuracy: 0.3580

Epoch 9/100

450/450 [=====] - 6s 12ms/step - loss: 1.5722 - accuracy: 0.4047 - val_loss: 2.1283 - val_accuracy: 0.2840

Epoch 10/100

450/450 [=====] - 6s 14ms/step - loss: 1.5400 - accuracy: 0.4204 - val_loss: 1.4777 - val_accuracy: 0.4460

Epoch 11/100

450/450 [=====] - 6s 12ms/step - loss: 1.5519 - accuracy: 0.4187 - val_loss: 1.4593 - val_accuracy: 0.4300

Epoch 12/100

450/450 [=====] - 7s 15ms/step - loss: 1.5240 - accuracy: 0.4273 - val_loss: 2.0205 - val_accuracy: 0.3300

Epoch 13/100

450/450 [=====] - 6s 13ms/step - loss: 1.4983 - accuracy: 0.4409 - val_loss: 1.7936 - val_accuracy: 0.3660

Epoch 14/100

450/450 [=====] - 7s 15ms/step - loss: 1.4722 - accuracy: 0.4547 - val_loss: 3.5521 - val_accuracy: 0.2660

Epoch 15/100

450/450 [=====] - 6s 12ms/step - loss: 1.4589 - accuracy: 0.4660 - val_loss: 1.4798 - val_accuracy: 0.4680

Epoch 16/100
450/450 [=====] - 6s 14ms/step - loss: 1.4256 - accuracy: 0.4616 - val_loss: 1.3579 - val_accuracy: 0.4800
Epoch 17/100
450/450 [=====] - 6s 13ms/step - loss: 1.4090 - accuracy: 0.4682 - val_loss: 1.2221 - val_accuracy: 0.5460
Epoch 18/100
450/450 [=====] - 7s 15ms/step - loss: 1.4117 - accuracy: 0.4831 - val_loss: 1.7445 - val_accuracy: 0.3980
Epoch 19/100
450/450 [=====] - 6s 12ms/step - loss: 1.3881 - accuracy: 0.4824 - val_loss: 1.3782 - val_accuracy: 0.4720
Epoch 20/100
450/450 [=====] - 6s 14ms/step - loss: 1.3933 - accuracy: 0.4800 - val_loss: 1.4160 - val_accuracy: 0.4960
Epoch 21/100
450/450 [=====] - 7s 15ms/step - loss: 1.3732 - accuracy: 0.4982 - val_loss: 1.3235 - val_accuracy: 0.5220
Epoch 22/100
450/450 [=====] - 6s 12ms/step - loss: 1.3691 - accuracy: 0.5000 - val_loss: 1.4279 - val_accuracy: 0.4780
Epoch 23/100
450/450 [=====] - 6s 14ms/step - loss: 1.3445 - accuracy: 0.5056 - val_loss: 1.1892 - val_accuracy: 0.5640
Epoch 24/100
450/450 [=====] - 6s 12ms/step - loss: 1.3320 - accuracy: 0.5089 - val_loss: 1.4513 - val_accuracy: 0.4880
Epoch 25/100
450/450 [=====] - 7s 15ms/step - loss: 1.3317 - accuracy: 0.5104 - val_loss: 1.3240 - val_accuracy: 0.5220
Epoch 26/100
450/450 [=====] - 6s 12ms/step - loss: 1.3116 - accuracy: 0.5207 - val_loss: 1.2255 - val_accuracy: 0.5340
Epoch 27/100
450/450 [=====] - 7s 15ms/step - loss: 1.3161 - accuracy: 0.5162 - val_loss: 1.2525 - val_accuracy: 0.5180
Epoch 28/100
450/450 [=====] - 6s 13ms/step - loss: 1.3053 - accuracy: 0.5180 - val_loss: 1.1635 - val_accuracy: 0.5980
Epoch 29/100
450/450 [=====] - 7s 15ms/step - loss: 1.2903 - accuracy: 0.5229 - val_loss: 1.1538 - val_accuracy: 0.6000
Epoch 30/100
450/450 [=====] - 6s 12ms/step - loss: 1.2791 - accuracy: 0.5298 - val_loss: 1.7265 - val_accuracy: 0.4480
Epoch 31/100
450/450 [=====] - 6s 14ms/step - loss: 1.2914 - accuracy: 0.5269 - val_loss: 1.2148 - val_accuracy: 0.5680

Epoch 32/100
450/450 [=====] - 6s 13ms/step - loss: 1.2689 - accuracy: 0.5300 - val_loss: 1.1514 - val_accuracy: 0.5920
Epoch 33/100
450/450 [=====] - 6s 14ms/step - loss: 1.2539 - accuracy: 0.5376 - val_loss: 1.2365 - val_accuracy: 0.5460
Epoch 34/100
450/450 [=====] - 6s 12ms/step - loss: 1.2408 - accuracy: 0.5418 - val_loss: 1.4547 - val_accuracy: 0.5100
Epoch 35/100
450/450 [=====] - 7s 14ms/step - loss: 1.2354 - accuracy: 0.5482 - val_loss: 1.2517 - val_accuracy: 0.5640
Epoch 36/100
450/450 [=====] - 6s 12ms/step - loss: 1.2478 - accuracy: 0.5431 - val_loss: 1.1473 - val_accuracy: 0.5820
Epoch 37/100
450/450 [=====] - 7s 14ms/step - loss: 1.2325 - accuracy: 0.5604 - val_loss: 1.2154 - val_accuracy: 0.5660
Epoch 38/100
450/450 [=====] - 5s 12ms/step - loss: 1.2225 - accuracy: 0.5558 - val_loss: 1.2052 - val_accuracy: 0.5960
Epoch 39/100
450/450 [=====] - 6s 14ms/step - loss: 1.2157 - accuracy: 0.5533 - val_loss: 1.1928 - val_accuracy: 0.5640
Epoch 40/100
450/450 [=====] - 6s 13ms/step - loss: 1.2312 - accuracy: 0.5500 - val_loss: 1.1695 - val_accuracy: 0.5980
Epoch 41/100
450/450 [=====] - 7s 15ms/step - loss: 1.1942 - accuracy: 0.5631 - val_loss: 1.1628 - val_accuracy: 0.5740
Epoch 42/100
450/450 [=====] - 6s 12ms/step - loss: 1.1753 - accuracy: 0.5796 - val_loss: 1.2074 - val_accuracy: 0.6160
Epoch 43/100
450/450 [=====] - 6s 14ms/step - loss: 1.1893 - accuracy: 0.5649 - val_loss: 1.3253 - val_accuracy: 0.5360
Epoch 44/100
450/450 [=====] - 6s 12ms/step - loss: 1.2337 - accuracy: 0.5627 - val_loss: 1.2316 - val_accuracy: 0.5600
Epoch 45/100
450/450 [=====] - 7s 15ms/step - loss: 1.1773 - accuracy: 0.5716 - val_loss: 1.1840 - val_accuracy: 0.5620
Epoch 46/100
450/450 [=====] - 6s 12ms/step - loss: 1.1712 - accuracy: 0.5744 - val_loss: 1.1397 - val_accuracy: 0.6060
Epoch 47/100
450/450 [=====] - 7s 15ms/step - loss: 1.1687 - accuracy: 0.5778 - val_loss: 1.1424 - val_accuracy: 0.6000

Epoch 48/100
450/450 [=====] - 6s 13ms/step - loss: 1.1828 -
accuracy: 0.5693 - val_loss: 1.1638 - val_accuracy: 0.5740
Epoch 49/100
450/450 [=====] - 7s 14ms/step - loss: 1.1581 -
accuracy: 0.5771 - val_loss: 1.3094 - val_accuracy: 0.5580
Epoch 50/100
450/450 [=====] - 6s 12ms/step - loss: 1.1559 -
accuracy: 0.5796 - val_loss: 1.2321 - val_accuracy: 0.5600
Epoch 51/100
450/450 [=====] - 7s 15ms/step - loss: 1.1859 -
accuracy: 0.5742 - val_loss: 1.2658 - val_accuracy: 0.5740
Epoch 52/100
450/450 [=====] - 6s 13ms/step - loss: 1.1670 -
accuracy: 0.5853 - val_loss: 1.4662 - val_accuracy: 0.5040
Epoch 53/100
450/450 [=====] - 6s 14ms/step - loss: 1.1625 -
accuracy: 0.5833 - val_loss: 1.3327 - val_accuracy: 0.5740
Epoch 54/100
450/450 [=====] - 6s 12ms/step - loss: 1.1430 -
accuracy: 0.5791 - val_loss: 1.2842 - val_accuracy: 0.5820
Epoch 55/100
450/450 [=====] - 6s 14ms/step - loss: 1.1210 -
accuracy: 0.5909 - val_loss: 1.6690 - val_accuracy: 0.5140
Epoch 56/100
450/450 [=====] - 6s 12ms/step - loss: 1.1395 -
accuracy: 0.5891 - val_loss: 1.2477 - val_accuracy: 0.5900
Epoch 57/100
450/450 [=====] - 7s 15ms/step - loss: 1.1416 -
accuracy: 0.5884 - val_loss: 1.0695 - val_accuracy: 0.6340
Epoch 58/100
450/450 [=====] - 6s 13ms/step - loss: 1.1429 -
accuracy: 0.5842 - val_loss: 1.2281 - val_accuracy: 0.5820
Epoch 59/100
450/450 [=====] - 6s 14ms/step - loss: 1.1460 -
accuracy: 0.5898 - val_loss: 1.0874 - val_accuracy: 0.6200
Epoch 60/100
450/450 [=====] - 7s 15ms/step - loss: 1.1250 -
accuracy: 0.5911 - val_loss: 1.1943 - val_accuracy: 0.5940
Epoch 61/100
450/450 [=====] - 6s 12ms/step - loss: 1.1216 -
accuracy: 0.5862 - val_loss: 1.2558 - val_accuracy: 0.5500
Epoch 62/100
450/450 [=====] - 7s 15ms/step - loss: 1.1145 -
accuracy: 0.5944 - val_loss: 1.1502 - val_accuracy: 0.5840
Epoch 63/100
450/450 [=====] - 6s 12ms/step - loss: 1.1128 -
accuracy: 0.6020 - val_loss: 1.3086 - val_accuracy: 0.5960

Epoch 64/100
450/450 [=====] - 7s 14ms/step - loss: 1.1131 -
accuracy: 0.5976 - val_loss: 1.2359 - val_accuracy: 0.5860
Epoch 65/100
450/450 [=====] - 6s 12ms/step - loss: 1.1120 -
accuracy: 0.5993 - val_loss: 1.1795 - val_accuracy: 0.5780
Epoch 66/100
450/450 [=====] - 6s 13ms/step - loss: 1.0987 -
accuracy: 0.6058 - val_loss: 1.1759 - val_accuracy: 0.6120
Epoch 67/100
450/450 [=====] - 6s 13ms/step - loss: 1.1085 -
accuracy: 0.6016 - val_loss: 1.0956 - val_accuracy: 0.6280
Epoch 68/100
450/450 [=====] - 6s 13ms/step - loss: 1.1028 -
accuracy: 0.5987 - val_loss: 1.1208 - val_accuracy: 0.6080
Epoch 69/100
450/450 [=====] - 6s 13ms/step - loss: 1.1054 -
accuracy: 0.5993 - val_loss: 1.1493 - val_accuracy: 0.6220
Epoch 70/100
450/450 [=====] - 6s 13ms/step - loss: 1.0805 -
accuracy: 0.6076 - val_loss: 1.3525 - val_accuracy: 0.5880
Epoch 71/100
450/450 [=====] - 6s 13ms/step - loss: 1.1125 -
accuracy: 0.5902 - val_loss: 1.2326 - val_accuracy: 0.6320
Epoch 72/100
450/450 [=====] - 6s 13ms/step - loss: 1.1127 -
accuracy: 0.5913 - val_loss: 1.0367 - val_accuracy: 0.6380
Epoch 73/100
450/450 [=====] - 7s 15ms/step - loss: 1.0679 -
accuracy: 0.6162 - val_loss: 1.2289 - val_accuracy: 0.6060
Epoch 74/100
450/450 [=====] - 6s 12ms/step - loss: 1.0815 -
accuracy: 0.6127 - val_loss: 1.1822 - val_accuracy: 0.5960
Epoch 75/100
450/450 [=====] - 7s 14ms/step - loss: 1.0811 -
accuracy: 0.6062 - val_loss: 1.0663 - val_accuracy: 0.6120
Epoch 76/100
450/450 [=====] - 7s 15ms/step - loss: 1.0974 -
accuracy: 0.6027 - val_loss: 1.0462 - val_accuracy: 0.6380
Epoch 77/100
450/450 [=====] - 6s 12ms/step - loss: 1.0553 -
accuracy: 0.6149 - val_loss: 1.2851 - val_accuracy: 0.5880
Epoch 78/100
450/450 [=====] - 6s 14ms/step - loss: 1.0917 -
accuracy: 0.6016 - val_loss: 1.0616 - val_accuracy: 0.6220
Epoch 79/100
450/450 [=====] - 6s 12ms/step - loss: 1.0823 -
accuracy: 0.6118 - val_loss: 1.1079 - val_accuracy: 0.6180

Epoch 80/100
450/450 [=====] - 7s 15ms/step - loss: 1.0740 - accuracy: 0.6120 - val_loss: 1.2189 - val_accuracy: 0.5940
Epoch 81/100
450/450 [=====] - 6s 13ms/step - loss: 1.0808 - accuracy: 0.6113 - val_loss: 1.1286 - val_accuracy: 0.6260
Epoch 82/100
450/450 [=====] - 7s 15ms/step - loss: 1.0804 - accuracy: 0.6113 - val_loss: 1.0766 - val_accuracy: 0.6340
Epoch 83/100
450/450 [=====] - 6s 13ms/step - loss: 1.0702 - accuracy: 0.6171 - val_loss: 1.4214 - val_accuracy: 0.5340
Epoch 84/100
450/450 [=====] - 7s 15ms/step - loss: 1.0908 - accuracy: 0.6096 - val_loss: 1.4015 - val_accuracy: 0.5840
Epoch 85/100
450/450 [=====] - 6s 12ms/step - loss: 1.0674 - accuracy: 0.6156 - val_loss: 1.1400 - val_accuracy: 0.6180
Epoch 86/100
450/450 [=====] - 6s 14ms/step - loss: 1.0827 - accuracy: 0.6076 - val_loss: 1.0858 - val_accuracy: 0.6160
Epoch 87/100
450/450 [=====] - 6s 13ms/step - loss: 1.0504 - accuracy: 0.6213 - val_loss: 1.1134 - val_accuracy: 0.6120
Epoch 88/100
450/450 [=====] - 6s 14ms/step - loss: 1.0759 - accuracy: 0.6116 - val_loss: 1.1505 - val_accuracy: 0.6260
Epoch 89/100
450/450 [=====] - 6s 12ms/step - loss: 1.0474 - accuracy: 0.6278 - val_loss: 1.0894 - val_accuracy: 0.6600
Epoch 90/100
450/450 [=====] - 6s 14ms/step - loss: 1.0892 - accuracy: 0.6160 - val_loss: 1.1571 - val_accuracy: 0.6100
Epoch 91/100
450/450 [=====] - 6s 13ms/step - loss: 1.0517 - accuracy: 0.6313 - val_loss: 1.0817 - val_accuracy: 0.6440
Epoch 92/100
450/450 [=====] - 6s 14ms/step - loss: 1.0614 - accuracy: 0.6156 - val_loss: 1.4039 - val_accuracy: 0.6140
Epoch 93/100
450/450 [=====] - 6s 12ms/step - loss: 1.0534 - accuracy: 0.6307 - val_loss: 1.2693 - val_accuracy: 0.5960
Epoch 94/100
450/450 [=====] - 7s 15ms/step - loss: 1.0379 - accuracy: 0.6282 - val_loss: 1.1655 - val_accuracy: 0.6300
Epoch 95/100
450/450 [=====] - 6s 13ms/step - loss: 1.0436 - accuracy: 0.6209 - val_loss: 1.2218 - val_accuracy: 0.6020

```
Epoch 96/100
450/450 [=====] - 6s 14ms/step - loss: 1.0629 -
accuracy: 0.6291 - val_loss: 1.1238 - val_accuracy: 0.6240
Epoch 97/100
450/450 [=====] - 6s 12ms/step - loss: 1.0338 -
accuracy: 0.6287 - val_loss: 1.2165 - val_accuracy: 0.6160
Epoch 98/100
450/450 [=====] - 6s 14ms/step - loss: 1.0401 -
accuracy: 0.6320 - val_loss: 1.2140 - val_accuracy: 0.6200
Epoch 99/100
450/450 [=====] - 6s 13ms/step - loss: 1.0427 -
accuracy: 0.6313 - val_loss: 1.2121 - val_accuracy: 0.5840
Epoch 100/100
450/450 [=====] - 7s 14ms/step - loss: 1.0303 -
accuracy: 0.6429 - val_loss: 1.2009 - val_accuracy: 0.6360
```

The best model so far parameter is: num_block=4, dropout=0.5, alpha=0.5, num_epoch=100, optimizer='rmsprop', learning_rate=0.001, use_skip=True, is_augmentation=True, mixup=0.1, Validation accuracy=0.636

```
[ ]: model_3_7_3.model.save("model_3_7_3.h5")#save the model, and this is the best
      ↪model found so far
```

```
[ ]: from tensorflow.keras.models import load_model

loaded_model = load_model("model_3_7_3.h5")
```

```
[ ]: num_samples = 25
sample_dataset = data_manager.ds_test.take(num_samples)
model_3_7_3.predict(sample_dataset.batch(num_samples), data_manager.ds_info)#
      ↪test the model on test dataset
```

```
1/1 [=====] - 1s 641ms/step
Sample 1: Predicted label - car
Sample 2: Predicted label - cat
Sample 3: Predicted label - cat
Sample 4: Predicted label - bird
Sample 5: Predicted label - monkey
Sample 6: Predicted label - cat
Sample 7: Predicted label - cat
Sample 8: Predicted label - dog
Sample 9: Predicted label - horse
Sample 10: Predicted label - deer
Sample 11: Predicted label - deer
Sample 12: Predicted label - cat
Sample 13: Predicted label - dog
Sample 14: Predicted label - ship
Sample 15: Predicted label - deer
```

Sample 16: Predicted label - deer
 Sample 17: Predicted label - dog
 Sample 18: Predicted label - airplane
 Sample 19: Predicted label - airplane
 Sample 20: Predicted label - monkey
 Sample 21: Predicted label - horse
 Sample 22: Predicted label - dog
 Sample 23: Predicted label - ship
 Sample 24: Predicted label - airplane
 Sample 25: Predicted label - airplane

```
[ ]: num_samples = 20
sample_dataset = data_manager.ds_test.take(num_samples)
model_3_7_3.plot_predictions(sample_dataset, data_manager.ds_info,
    ↳ num_samples=num_samples, grid_shape=(4, 5))#visualize the model on test
    ↳ dataset
```

1/1 [=====] - 0s 78ms/step



In conclusion: the best model found so far is: number of block:4 learning rate:0.001 use skip_connection use data_augmentation dropout=0.5 alpha for mixup=0.5 optimizer:rmsprop number of epoches:100

And best model found has been saved to “model_3_7_3.h5”

For questions **3.8 and 3.9**, you can reuse code in lectures or labs. You should not use third-party libraries such as ClevenHans.

0.1.8 Question 3.8: Attack your model

Attack your best obtained model with PGD, MIM, and FGSM attacks with $\epsilon = 0.0313$, $k = 20$, $\eta = 0.002$ on the testing set. Write the code for the attacks and report the robust accuracies. Also choose a random set of 20 clean images in the testing set and visualize the original and attacked images.

[5 points]

```
[ ]: # YOU ARE REQUIRED TO INSERT YOUR CODES IN THIS CELL
     # You can add more cells if necessary
```

0.1.9 Question 3.9: Train a robust model

Train a robust model using adversarial training with PGD $\epsilon = 0.0313$, $k = 10$, $\eta = 0.002$. Write the code for the adversarial training and report the robust accuracies. After finishing the training, you need to store your best robust model in the folder `./models` and load the model to evaluate the robust accuracies for PGD, MIM, and FGSM attacks with $\epsilon = 0.0313$, $k = 20$, $\eta = 0.002$ on the testing set.

[5 points]

```
[ ]: # YOU ARE REQUIRED TO INSERT YOUR CODES IN THIS CELL
     # You can add more cells if necessary
```

The following is an exploring question with bonus points. It is great if you try to do this question, but it is **totally optional**. In this question, we will investigate a recent SOTA technique to improve the generalization ability of deep nets named *Sharpness-Aware Minimization (SAM)* ([link to the main paper](#)). Furthermore, SAM is simple and efficient technique, but roughly doubles the training time due to its required computation. If you have an idea to improve SAM, it would be a great paper to top-tier venues in machine learning and computer vision. Highly recommend to give it a try.

0.1.10 Question 3.10 (bonus question)

Read the SAM paper ([link to the main paper](#)). Try to apply this technique to the best obtained model and report the results. For the purpose of implementing SAM, we can flexibly add more cells and extensions to the `model.py` file.

[5 points]

```
[ ]: # YOU ARE REQUIRED TO INSERT YOUR CODES IN THIS CELL  
# You can add more cells if necessary
```

**

END OF ASSIGNMENT

GOOD LUCK WITH YOUR ASSIGNMENT 1!

**